

University of Zurich[UZH]

ETH

Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Multi-Currency Portfolio Optimization with General Risk Measures

Master Thesis

Anej Rozman

Supervisor: Prof. Patrick Cheridito
Co-Supervisor: Dr. Urban Ulrych

Department of Mathematics, ETH Zürich

Acknowledgments

Abstract

Multi-currency portfolio construction requires managing asset and currency exposures together, yet it remains unclear whether these decisions are best made simultaneously or sequentially as the investable universe grows. In this thesis, we build on the work of [30, 31]. We investigate whether joint optimization of asset and currency weights provides a benefit over the standard sequential currency overlay, across five risk measures and three universe sizes. The empirical backtests show that joint optimization does not deliver a meaningful improvement in realized excess returns or risk-adjusted returns over the overlay, in any configuration tested, with the possible exception of the smallest universe. This motivates a closer look at where joint optimization differs from the overlay by construction, which is risk minimization. The thesis identifies two mechanisms behind the joint program's in-sample risk advantage for the tail measures, a tail-fitting mechanism and an implicit risk-return tradeoff, and shows that neither generalizes out of sample. For the deviation measures, the in-sample advantage is small and most of it survives out of sample. The out-of-sample advantage of joint optimization is positive only in the small universe, where the currency weights are a larger share of the decision variables.

Contents

Contents iii

1 Introduction 1
- 1.1 Motivation: International Diversification, Currency Risk, and Tail Risk 1
- 1.2 Contributions 1
- 1.3 Thesis Structure 1

2 Literature Review 2

3 Optimization Framework 3
- 3.1 Model-Free Multi-Currency Portfolio Returns 3
  - 3.1.1 Unhedged Portfolio Returns 3
  - 3.1.2 Hedged Portfolio Returns 4
  - 3.1.3 Matrix Representation and Decomposition 5
- 3.2 Mean–Risk Portfolio Optimization 6
- 3.3 Variance and Semi-Variance 6
- 3.4 Mean Absolute Deviation 6
- 3.5 Expected Shortfall and Second Order Expected Shortfall 6

4 Data and Implementation 7
- 4.1 Data 7
- 4.2 Backtest Implementation 11
  - 4.2.1 Backtesting Framework 11
  - 4.2.2 Backtest Configuration 12
  - 4.2.3 Performance Evaluation Metrics 13

5 Backtest Analysis 16
- 5.1 Baseline Backtest Results 16
- 5.2 Regularization 21
- 5.3 Mean-Risk Optimization 26
- 5.4 Conclusion 30

6 Joint versus Sequential Optimization 32
- 6.1 Advantage of Joint Optimization for Risk Minimization 32
- 6.2 Tail Measures 36
  - 6.2.1 Level Analysis 36
  - 6.2.2 De-meaning the Tail Measures 37
- 6.3 Deviation Measures 38

6.4 Conclusion 39

7 Conclusion **40**
7.1 Findings 40
7.2 Limitations and Future Research 42

Appendices **43**

A Optimization Derivation: Expected Shortfall and Second-Order Expected Shortfall **44**
A.1 Auxiliary Formulations of Expected Shortfall and Second-Order Expected Shortfall 44
A.1.1 Expected Shortfall 44
A.1.2 Second-Order Expected Shortfall 46
A.2 Representation for Portfolio Optimization 47
A.2.1 Expected Shortfall 47
A.2.2 Second-Order Expected Shortfall 48
A.3 Regularization 49
A.3.1 L^1 Regularization 49
A.3.2 L^2 Regularization 50
A.3.3 Leverage Constraint 50
A.3.4 Expected Shortfall 50
A.3.5 Second-Order Expected Shortfall 51
A.4 Sequential Optimization 51
A.4.1 Currency Overlay under Expected Shortfall 52
A.4.2 Currency Overlay under Second-Order Expected Shortfall 53
A.5 Joint Optimization 53
A.5.1 Joint Optimization under Expected Shortfall 54
A.5.2 Joint Optimization under Second-Order Expected Shortfall 55

B Optimization Derivation: Variance and Semi-Variance **56**
B.1 Quadratic Program Representation of the Mean-Variance Portfolio Optimization 56
B.2 Regularization 57
B.2.1 L_1, L_2 Regularization and Leverage Constraint 57
B.2.2 Estimators 58
B.2.3 Quadratic Program Representation of the Regularized Portfolio Optimization 59
B.3 Separate (Currency Overlay) Optimization 59
B.3.1 Asset Optimization 60
B.3.2 Currency Overlay Optimization 60
B.4 Joint Optimization 61

C Optimization Derivation: Mean Absolute Deviation **64**
C.1 Linear Program Represenation of the Mean Absolute Deviation Portfolio Optimization Problem 64
C.2 Regularization 64
C.2.1 L_1, L_2 Regularization and Leverage Constraint 64
C.2.2 Quadratic Program Representaion of the Regularized Portfolio Optimization Problem 64
C.3 Sequential Overlay Optimization 64
C.3.1 Asset Optimizaiton 64
C.3.2 Currency Overlay Optimization 64

C.4 Joint Optimization 64
C.5 Adding the Mean Term 64

D Coherent Risk Measures and Deviation Measures **65**

E Numerical Implementation of the Optimization Problems **68**
E.1 The CVXPY Library 68
E.2 Numerical Implementation 68
E.3 Computational Complexity and Performance 69

F Transaction Costs **70**

G Additional Results **72**
G.1 Hyperparameter Sensitivity Analysis 72
G.2 Portfolio Sparsity Analysis 76
G.3 Risk Aversion Instability under Mean Variance Optimization 77
G.4 Additional Backtest Performance Results 77

Bibliography **81**

Chapter 1

Introduction

1.1 Motivation: International Diversification, Currency Risk, and Tail Risk

1.2 Contributions

a systematic comaprison across different risk measures, universe sizes, which the exsititng literature covers only partially.

the mechanism analysis of the tail measures

the universe size results. The joint advantage depends on the currency share of the decision variables.. dimensionality effect.

1.3 Thesis Structure

Chapter 2

Literature Review

Cite the following: International allocation and currency risk management, Currency overlay vs. joint optimization, Downside risk measures and Expected Shortfall, Instability of ES optimization and the role of regularization, Generalized objectives and decision-focused / end-to-end learning.

In this chapter, we review the relevant literature on multi-currency portfolio optimization. We discuss the contributions and methodologies in the field, highlighting how our work builds upon and extends existing research.

The core of our analysis builds on two recent studies that apply regularization techniques to international asset allocation. Burkhardt and Ulrych (2023) introduce a joint optimization approach within a mean-variance framework. By using L1 and L2 regularization, they show that jointly optimizing assets and currencies leads to better out-of-sample performance than the standard industry practice of using a separate currency overlay.

Extending this framework, Lucescu and Ulrych (2026) consider an Expected Shortfall optimization framework to better manage downside tail risk. To fix the numerical instability common in Expected Shortfall minimization, they apply similar regularization methods. Their findings reveal a dimensional trade-off. Joint optimization works best in small asset universes, but the separate currency overlay approach becomes more effective for larger universes due to its built-in dimensionality reduction.

Reproducing and comparing these two baseline frameworks serves as the starting point for the extensions proposed in this thesis.

Chapter 3

Optimization Framework

This chapter develops the optimization framework used throughout the thesis, building up from a model of multi-currency portfolio returns to the risk objectives, the regularization that stabilizes them, and the two formulations through which asset and currency decisions are combined. The derivations follow the methodology presented in [31] and [30], use simple returns, and assume that rebalancing occurs at the end of each investment period. We begin in Section 3.1 with a model-free representation of one-period portfolio returns for an investor who holds assets denominated in multiple currencies and can take currency forward positions for hedging, decomposing the unhedged return of a foreign asset into its local, currency, and cross-currency components, introducing the currency-mapping matrix that links assets to their implicit currency exposures, and expressing the hedged portfolio return in terms of net currency exposures.

With the return model in place, we formulate the optimization problem. Section **??** states the classical mean-variance criterion, and Section **??** introduces Expected Shortfall as a coherent downside-risk measure, together with the Rockafellar–Uryasev representation that casts its minimization as a tractable linear program. Both are presented as instances of a common mean–risk objective, so that the framework treats the risk measure as a swappable component rather than a fixed choice. Section **??** then introduces the L_1 and L_2 penalties and the leverage constraint used to stabilize the optimization and improve out-of-sample performance.

Finally, Sections **??** and **??** assemble these elements into the two approaches to multi-currency allocation studied throughout the thesis: the separate (currency overlay) formulation, in which asset and currency decisions are made sequentially, and the joint formulation, in which they are optimized simultaneously. The detailed derivations of the resulting quadratic programs are found in Appendix A and Appendix B.

3.1 Model-Free Multi-Currency Portfolio Returns

3.1.1 Unhedged Portfolio Returns

We consider an international investor who allocates their wealth across N assets denominated in up to $M+1$ currencies. We fix a discrete time index $t \in \{0, 1, 2, \dots\}$ and one investment period from t to $t{+}1$. Let the investor's *domestic currency* be indexed by 0 and let the remaining M foreign currencies be indexed by $\mathcal{C} \coloneqq \{1, 2, \dots, M\}$.

Let $P_{i,t} > 0$ be the price of asset i at time t in its local currency $c_i \in \mathcal{C} \cup \{0\}$. Define the

corresponding *simple local-currency asset return* over $(t, t{+}1]$ as

$$R_{i,t+1} \coloneqq \frac{P_{i,t+1} - P_{i,t}}{P_{i,t}}.$$

Thus $P_{i,t+1} = P_{i,t}(1 + R_{i,t+1})$.

For each foreign currency $c \in \mathcal{C}$, let $S_{c,t} > 0$ denote the exchange rate at time t, quoted as units of domestic currency per one unit of foreign currency c. The associated *simple currency return* over $(t, t{+}1]$ is

$$e_{c,t+1} \coloneqq \frac{S_{c,t+1} - S_{c,t}}{S_{c,t}}.$$

Equivalently, $S_{c,t+1} = S_{c,t}(1+e_{c,t+1})$. For the domestic currency, we set $S_{0,t} \equiv 1$ and $e_{0,t+1} \equiv 0$.

The domestic-currency value of holding one unit of asset i at time t is $P_{i,t}\, S_{c_i,t}$. Therefore, the *unhedged domestic-currency simple return* of asset i is

$$\begin{aligned} R^u_{i,t+1} &= \frac{P_{i,t+1}S_{c_i,t+1} - P_{i,t}S_{c_i,t}}{P_{i,t}S_{c_i,t}} \\ &= \frac{P_{i,t}(1 + R_{i,t+1})\, S_{c_i,t}(1 + e_{c_i,t+1}) - P_{i,t}S_{c_i,t}}{P_{i,t}S_{c_i,t}} \\ &= R_{i,t+1} + e_{c_i,t+1} + R_{i,t+1}\, e_{c_i,t+1}. \end{aligned} \tag{3.1}$$

Let $\boldsymbol{x}_t \in \mathbb{R}^N$ be the vector of asset portfolio weights chosen at time t and held until $t{+}1$, where $x_{i,t}$ is the fraction of total portfolio value invested in asset i at time t. We impose the standard budget constraint $\mathbf{1}^\mathsf{T}\boldsymbol{x}_t = 1$.

Assuming no rebalancing between times t and $t{+}1$, the domestic-currency portfolio's unhedged return is given by

$$R^u_{\mathcal{P},t+1} = \sum_{i=1}^{N} x_{i,t}\, R^u_{i,t+1}.$$

3.1.2 Hedged Portfolio Returns

To manage currency risk, the investor can take positions in one-period currency forward contracts. For each foreign currency $c \in \mathcal{C}$, let $F_{c,t}$ be the forward rate quoted at time t for settlement at $t{+}1$, in units of domestic currency per unit of foreign currency c. We define the *forward premium* by

$$f_{c,t} \coloneqq \frac{F_{c,t} - S_{c,t}}{S_{c,t}}.$$

The investor hedges by selling foreign currency forward. Let $\phi_{c,t}$ denote the relative notional of the *short* forward position in currency c at time t, expressed as a fraction of total portfolio value, with the convention that $\phi_{c,t} > 0$ corresponds to hedging. Per unit of notional, the short position pays $f_{c,t} - e_{c,t+1}$ in domestic currency over the period. Adding these forward payoffs to the unhedged asset returns, the *hedged portfolio return* is

$$R^h_{\mathcal{P},t+1} = \underbrace{\sum_{i=1}^{N} x_{i,t}\, R^u_{i,t+1}}_{R^u_{\mathcal{P},t+1}} + \sum_{c=1}^{M} \phi_{c,t}\,(f_{c,t} - e_{c,t+1}). \tag{3.2}$$

3.1.3 Matrix Representation and Decomposition

To express the returns in a compact form suitable for optimization, we define the currency incidence matrix $\boldsymbol{C} \in \{0,1\}^{N\times M}$ by

$$C_{i,c} \coloneqq \begin{cases} 1, & \text{if asset } i \text{ is denominated in foreign currency } c, \\ 0, & \text{otherwise.} \end{cases}$$

Each row of $\boldsymbol{C}$ has at most one nonzero entry, and the rows corresponding to domestic assets are identically zero. This allows us to define the portfolio's *implicit foreign-currency exposure* as

$$\boldsymbol{w}_t \coloneqq \boldsymbol{C}^\mathsf{T}\boldsymbol{x}_t \in \mathbb{R}^M, \qquad w_{c,t} = \sum_{i:\, c_i=c} x_{i,t},$$

such that $w_{c,t}$ is the total portfolio weight invested in assets denominated in currency c. The implicit domestic exposure follows residually as $w_{0,t} = 1 - \mathbf{1}_M^\mathsf{T}\boldsymbol{w}_t$, since the asset weights sum to one.

We express returns in vector form $\boldsymbol{R}_{t+1} \in \mathbb{R}^N$, $\boldsymbol{e}_{t+1} \in \mathbb{R}^M$, and the forward premiums as $\boldsymbol{f}_t \in \mathbb{R}^M$. The product $\boldsymbol{C}\boldsymbol{e}_{t+1} \in \mathbb{R}^N$ aligns currency returns to assets, with ith component $e_{c_i,t+1}$ for foreign assets and 0 for domestic ones. Using (3.1), the vector of *unhedged domestic returns* across assets can be written as

$$\boldsymbol{R}^u_{t+1} = \boldsymbol{R}_{t+1} + \boldsymbol{C}\boldsymbol{e}_{t+1} + (\boldsymbol{R}_{t+1} \odot \boldsymbol{C}\boldsymbol{e}_{t+1}),$$

where $\odot$ denotes the Hadamard product.

Expressing the forward positions in vector form as $\boldsymbol{\phi}_t \in \mathbb{R}^M$, the hedged portfolio return (3.2) can be written as

$$R^h_{\mathcal{P},t+1} = \boldsymbol{x}_t^\mathsf{T}\boldsymbol{R}^u_{t+1} + \boldsymbol{\phi}_t^\mathsf{T}(\boldsymbol{f}_t - \boldsymbol{e}_{t+1}). \tag{3.3}$$

Setting $\boldsymbol{\phi}_t = \boldsymbol{w}_t = \boldsymbol{C}^\mathsf{T}\boldsymbol{x}_t$ neutralizes the implicit currency exposure. Substituting this together with the unhedged return vector into (3.3), the *fully hedged portfolio return* is

$$\begin{aligned} R^{fh}_{\mathcal{P},t+1} &= \boldsymbol{x}_t^\mathsf{T}\boldsymbol{R}^u_{t+1} + (\boldsymbol{C}^\mathsf{T}\boldsymbol{x}_t)^\mathsf{T}(\boldsymbol{f}_t - \boldsymbol{e}_{t+1}) \\ &= \boldsymbol{x}_t^\mathsf{T}\Big[\boldsymbol{R}_{t+1} + \boldsymbol{C}\boldsymbol{e}_{t+1} + (\boldsymbol{R}_{t+1} \odot \boldsymbol{C}\boldsymbol{e}_{t+1})\Big] + \boldsymbol{x}_t^\mathsf{T}\boldsymbol{C}(\boldsymbol{f}_t - \boldsymbol{e}_{t+1}) \\ &= \boldsymbol{x}_t^\mathsf{T}\Big[\boldsymbol{R}_{t+1} + \boldsymbol{C}\boldsymbol{e}_{t+1} + (\boldsymbol{R}_{t+1} \odot \boldsymbol{C}\boldsymbol{e}_{t+1}) + \boldsymbol{C}\boldsymbol{f}_t - \boldsymbol{C}\boldsymbol{e}_{t+1}\Big] \\ &= \boldsymbol{x}_t^\mathsf{T}\Big[\boldsymbol{R}_{t+1} + \boldsymbol{C}\boldsymbol{f}_t + (\boldsymbol{R}_{t+1} \odot \boldsymbol{C}\boldsymbol{e}_{t+1})\Big] \\ &= \boldsymbol{x}_t^\mathsf{T}\boldsymbol{R}^{fh}_{t+1}. \end{aligned}$$

The first-order currency terms $\pm\boldsymbol{x}_t^\mathsf{T}\boldsymbol{C}\boldsymbol{e}_{t+1}$ cancel, leaving the *fully hedged asset return vector*

$$\boldsymbol{R}^{fh}_{t+1} = \boldsymbol{R}_{t+1} + \boldsymbol{C}\boldsymbol{f}_t + (\boldsymbol{R}_{t+1} \odot \boldsymbol{C}\boldsymbol{e}_{t+1}).$$

More generally, define the *net currency exposure* $\boldsymbol{\psi}_t \coloneqq \boldsymbol{w}_t - \boldsymbol{\phi}_t$. Imposing $\sum_{c=0}^M \phi_{c,t} = 1$ gives $\sum_{c=0}^M \psi_{c,t} = 0$, so the currency overlay is a zero-investment portfolio. Substituting $\boldsymbol{\phi}_t = \boldsymbol{w}_t - \boldsymbol{\psi}_t$ into (3.3) expresses the hedged return in terms of net exposures,

$$\begin{aligned} R^h_{\mathcal{P},t+1} &= \boldsymbol{x}_t^\mathsf{T}\boldsymbol{R}^u_{t+1} + (\boldsymbol{w}_t - \boldsymbol{\psi}_t)^\mathsf{T}(\boldsymbol{f}_t - \boldsymbol{e}_{t+1}) \\ &= \boldsymbol{x}_t^\mathsf{T}\boldsymbol{R}^u_{t+1} + \boldsymbol{w}_t^\mathsf{T}(\boldsymbol{f}_t - \boldsymbol{e}_{t+1}) + \boldsymbol{\psi}_t^\mathsf{T}(\boldsymbol{e}_{t+1} - \boldsymbol{f}_t) \\ &= \boldsymbol{x}_t^\mathsf{T}\boldsymbol{R}^{fh}_{t+1} + \boldsymbol{\psi}_t^\mathsf{T}(\boldsymbol{e}_{t+1} - \boldsymbol{f}_t), \end{aligned} \tag{3.4}$$

which we adopt throughout the following portfolio optimization formulations, as it makes the asset weights $\boldsymbol{x}_t$ and the net currency exposures $\boldsymbol{\psi}_t$ the explicit variables of the optimization.

3.2 Mean–Risk Portfolio Optimization

3.3 Variance and Semi-Variance

3.4 Mean Absolute Deviation

3.5 Expected Shortfall and Second Order Expected Shortfall

Chapter 4

Data and Implementation

In this chapter, we describe the dataset and the backtest implementation used for the empirical evaluation of the portfolio optimization strategies in Chapters 5 and 6. In Section 4.1 we detail the construction of three investable universes of different sizes from the largest stocks of 14 global economies. We analyze the fully hedged asset returns and the excess currency returns, and show that as the universe size increases, the average asset returns tend to decrease and their volatilities rise. We further show that the correlations between the asset and excess currency returns are mostly positive, with the commodity and risk-sensitive currencies being the most positively correlated and the safe-haven currencies negatively correlated with most economies. This dependence structure stays stable across universe sizes. In Section 4.2 we describe the implementation of the backtest. We first present in Section 4.2.1 the modular structure of the backtesting framework. In Section 4.2.2 we present the backtested strategies and their benchmarks, and specify the rebalancing frequency, the lookback window, the selection of the hyperparameters, and the modelling of transaction cost. Finally, in Section 4.2.3 we define the performance metrics on which the strategies are evaluated in Chapter 5.

4.1 Data

The empirical evaluation of all portfolio optimization strategies in this thesis relies on a dataset of international equities, foreign exchange spot rates, and forward exchange rates sourced from Refinitiv Datastream. The one-month U.S. Treasury bill rate used in the return and performance metric calculations is obtained from Kenneth R. French's online data library.[1] The datasets span from June 1990 to June 2023, covering several different market regimes, including the dot-com crash, the global financial crisis, the European sovereign debt crisis, and the COVID-19 pandemic.

We construct the investable universe from 14 economies: Australia, Canada, Denmark, Norway, Sweden, Hong Kong, Japan, New Zealand, Singapore, South Africa, Switzerland, the Eurozone, the United Kingdom, and the United States. The universe is rebuilt at the beginning of each year from the largest stocks by market capitalization at the time (a rolling-window approach that mitigates survivorship bias). At the start of each year we rank stocks within each economy and restrict the following year's investments to that set. To study how the universe size affects the optimization approaches considered, we form three cross-sections, comprising of the top 10, top 50, and top 100 stocks per economy. We refer to these as the small, medium and large universe. Because the constituents change each year, each universe spans more distinct stocks

[1] https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/data_library.html

over the full sample than are held on any single date, 407 for the top 10, 1705 for the top 50, and 3189 for the top 100, respectively.

The US dollar serves as the base currency throughout our analysis. For each stock we compute fully hedged returns, and for each currency the excess currency return, following the model-free return representation in Chapter 3.

Economies	Assets (%)						Currencies (%)	
	Top 10		Top 50		Top 100			
	Return	Vol.	Return	Vol.	Return	Vol.	Excess Return	Vol.
Australia	4.57	29.94	3.10	37.72	-1.88	58.00	1.31	11.59
Canada	9.10	32.89	6.58	40.24	1.86	59.24	-0.03	7.63
Denmark	8.35	34.05	3.04	39.00	-3.47	48.52	-0.08	9.68
Eurozone	5.09	32.24	5.05	34.01	3.25	35.46	-0.40	9.58
Hong Kong	4.42	37.59	-0.15	45.77	-3.48	52.59	-0.16	0.56
Japan	5.37	33.69	4.41	34.01	4.09	35.38	-2.07	10.47
New Zealand	2.88	29.93	-1.13	37.16	-4.56	41.49	2.48	11.36
Norway	3.17	43.30	-1.59	50.29	-2.61	53.26	-0.36	11.73
Singapore	3.32	35.43	-1.99	54.57	-4.84	62.12	0.18	5.35
South Africa	3.53	37.72	1.01	43.84	-4.39	53.94	1.22	15.18
Sweden	9.54	34.50	6.92	42.50	5.37	47.65	-1.16	11.50
Switzerland	6.38	29.04	5.67	30.32	3.97	33.65	-0.12	10.99
United Kingdom	2.76	31.50	2.63	34.71	1.29	36.22	0.12	9.42
United States	12.59	33.58	7.99	39.07	6.86	42.84	0.00	0.00

Table 4.1: Summary statistics of fully hedged equity returns and excess currency returns by economy. Asset statistics are reported for three investable universes (the top 10, 50, and 100 stocks per economy). Excess currency returns are defined as the currency return net of the forward premium. All numbers are reported as annualized percentages, from the perspective of a USD-based investor.

Table 4.1 reports annualized average returns and volatilities for the fully hedged equity returns and excess currency returns, by economy. For each economy, the asset figures are the cross-sectional average of the daily returns of its individual stocks.

We can see that as the investable universe becomes larger, the average returns tend to fall. In the small universe every economy has a positive average return, ranging from 2.76% (the United Kingdom) to 12.59% (the United States). In the medium universe, the range drops from -1.99% (Singapore) to 7.99% (the United States), and in the large universe the range goes from -4.84% (Singapore) to 6.86% (the United States), with many economies now having negative average returns. The trend reflects the firm size, the largest stocks post the strongest returns, and the smaller firms added pull the average down. Volatilities exhibit the opposite pattern, rising with universe size, as smaller, more volatile stocks are added to the investable universe. Switzerland exhibits the lowest volatility in every case, its volatility rises from 29.04% to 33.65%, and the upper end rises from 43.30% (Norway) in the small universe to 62.12% (Singapore) in the large universe.

Excess currency returns range between -2.07% for the Japanese yen and 2.48% for the New Zealand dollar, and excess currency volatilities range from 0.56% for the Hong Kong dollar, whose peg to the US dollar keeps it stable, up to 15.18% for the South African rand.

To quantify the dependence between asset and currency returns, we compute Pearson correlations between economy-level fully hedged equity returns and excess currency returns. For each economy we form an equally weighted daily portfolio of the fully hedged returns of its constituent stocks, resetting the constituents at each annual universe reconstitution.

Correlating each economy's series with each currency's excess return over the full sample gives a 14×13 correlation matrix. The US dollar is dropped from the currency axis, since its excess return is zero by construction. We repeat this for the three universe sizes. The resulting heatmaps are shown in Figures 4.1 to 4.3.

Across all three universes the correlations range from -0.2 to 0.4, and remain mostly positive. The results are intuitive and agree with known empirical findings. The strongest positive correlations belong to commodity and risk-sensitive currencies such as the AUD, CAD, NZD and ZAR, whose equity markets and currencies rise together in good times and fall together in bad [4, 17, 20]. The JPY and CHF are the only negatively correlated currencies with most other economies (ranging from -0.2 to 0.03), in line with their role as safe havens that strengthen when equities fall [25]. The EUR, DKK, and HKD are minimally correlated, reflecting the HKD's peg to the US dollar, and the DKK's peg to the EUR. This split between risky and safe-haven currencies is well documented [3, 25].

This is the dependence structure the overlay program discards when it optimizes the asset weights. It sets the asset weights first, using only equity information, and adds the currency hedge afterwards, whereas the joint program can act on it directly. The structure barely changes across universe sizes, and this matters for the results that follow. A shift in the relative performance of the joint and overlay programs across universe sizes in Chapters 5 and 6 is therefore not due to a change in the underlying dependence structure, and we analyze its causes in Chapter 6.

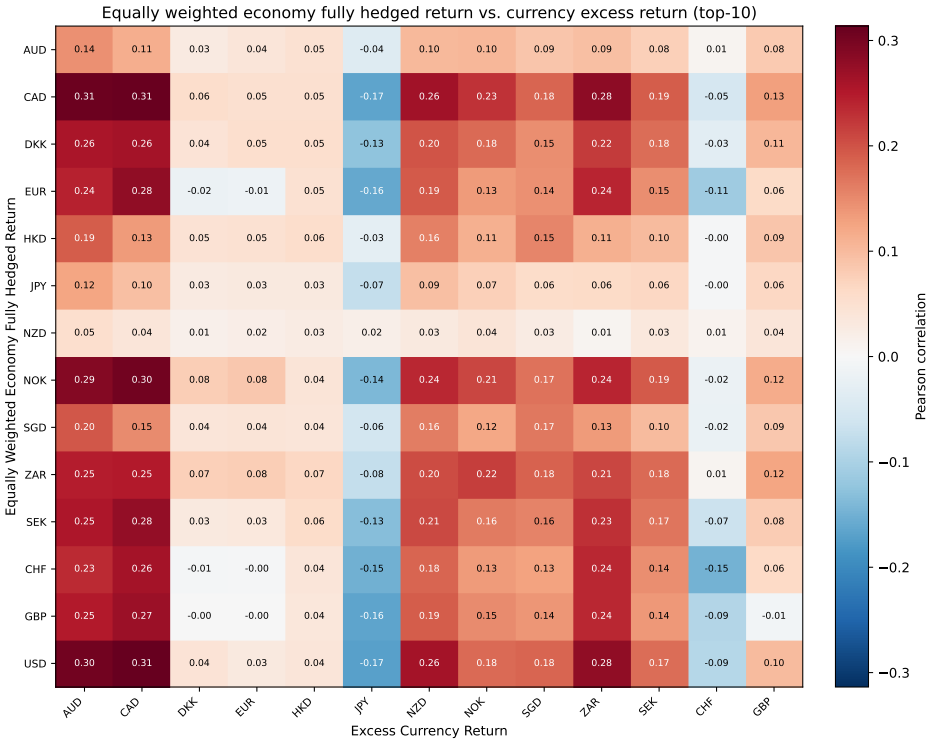


Figure 4.1: Pearson correlation between the fully hedged returns of the equally weighted economy portfolios and the excess currency returns in the small universe (top 10 stocks per economy). The US dollar is excluded from the currency axis, since it serves as the base currency and therefore its excess return is zero.

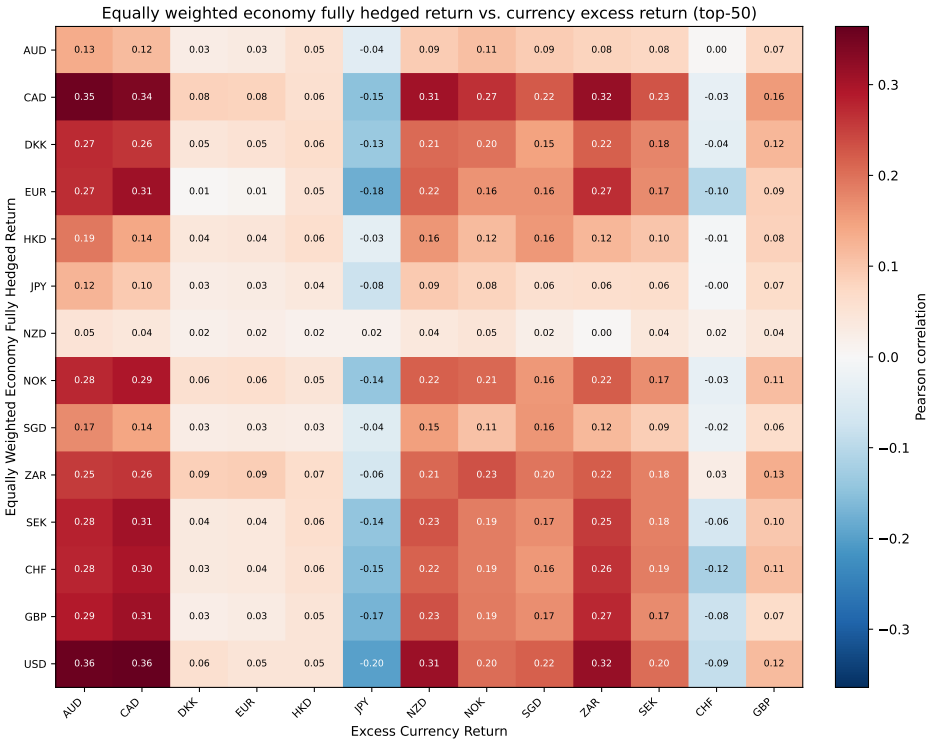


Figure 4.2: Pearson correlation between the fully hedged returns of the equally weighted economy portfolios and the excess currency returns in the medium universe (top 50 stocks per economy). The US dollar is excluded from the currency axis, since it serves as the base currency and therefore its excess return is zero.

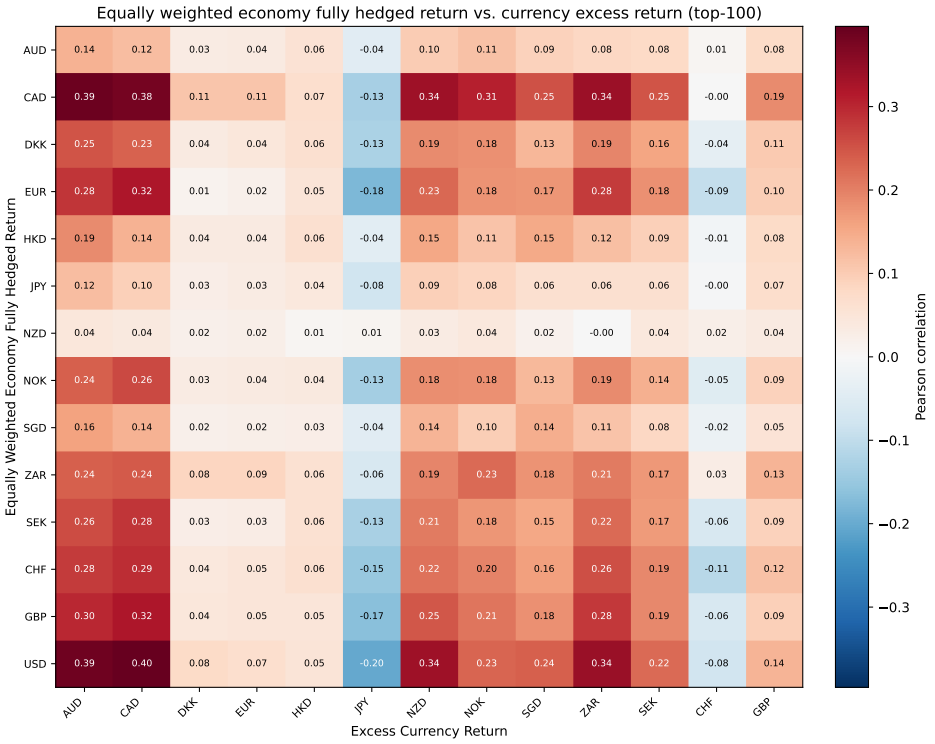


Figure 4.3: Pearson correlation between the fully hedged returns of the equally weighted economy portfolios and the excess currency returns in the large universe (top 100 stocks per economy). The US dollar is excluded from the currency axis, since it serves as the base currency and therefore its excess return is zero.

4.2 Backtest Implementation

4.2.1 Backtesting Framework

In Chapter 5 we evaluate a range of optimization strategies under different risk measures (variance, semi-variance, mean absolute deviation, Expected Shortfall and second-order Expected Shortfall), their treatment of currency exposure (overlay or joint), and the size of the investable universe. To keep these comparisons fair, the backtesting framework separates between the orchestration of the backtest and the optimization logic it evaluates.[2] The main backtesting loop that ingests data, rolls the estimation window forward, selects hyperparameters, and records performance is fixed, while the optimizer that produces portfolio weights is a swappable component. Any optimizer that conforms to a common interface is therefore evaluated under the same conditions. The framework is composed of three main components, a data handler, the backtester, and a family of optimizers under a common abstract base class, whose interaction is summarized in Figure 4.4.

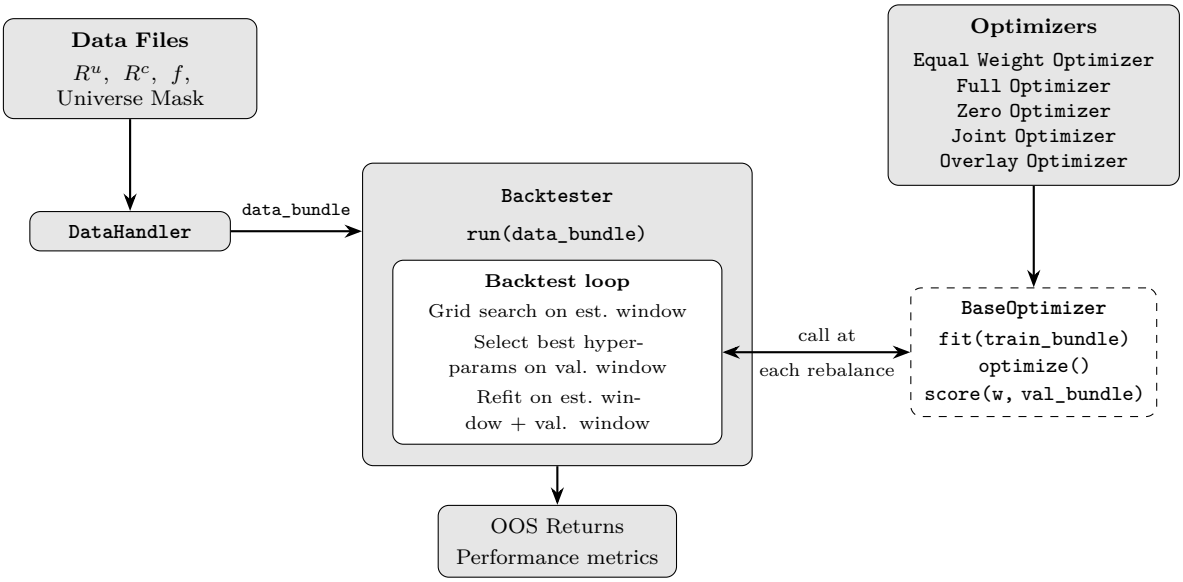


Figure 4.4: Schematic of the backtesting framework.

Data ingestion and alignment are implemented in the `DataHandler`. It loads the input data, the local unhedged asset returns R^u, the currency returns R^c, the forward premia f, and the universe mask that identifies which assets constitute the investable universe at each date. It then aligns them into a `data_bundle` which is passed to the backtester. The optimizers operate only on the bundled series, and the construction of returns from these inputs follows Section 3.1.

All implemented optimizers inherit from an abstract base class, `BaseOptimizer`, that defines the three-method interface through which the backtester runs them. The `fit` method receives a training data bundle and estimates and stores the quantities the optimizer requires for its objective (e.g. the covariance matrix for the variance optimizers), so that estimation is performed entirely within each optimizer. The `optimize` method solves the specific optimization problem. The `score` method takes a set of weights and a validation bundle and returns the validation performance used for hyperparameter selection. More details on the estimators and

[2]The full codebase with the backtest implementation is publicly available on GitHub: https://github.com/anejrozman/Multi-Currency-Portfolio-Optimization-with-General-Risk-Measures.

solver internals invoked inside the methods `fit` and `optimize` can be found in Appendix E.2.

The backtest itself is implemented in the `Backtester.run` method, which advances through the data sample on a set rebalancing frequency (daily, weekly, monthly, quarterly, or annual). At each rebalancing date, the optimizer is calibrated and deployed using a fixed rolling lookback window, partitioned into an estimation window and a validation window. Hyperparameters are selected by a grid search over this partition. Each configuration is fitted on the estimation window via `fit`, used to produce weights via `optimize`, and evaluated on the validation window via `score`, and the configuration with the best validation score is retrained on both the estimation and validation windows and finally used to produce the portfolio weights for the next out-of-sample period. These weights are held until the next rebalancing date, the window then rolls forward and the procedure repeats.

Each invocation of the backtester for a given optimizer produces a self-contained run directory. A new `run_<n>` folder is created in `results/runs/<optimizer_name>/` and collects the full record of that run: the out-of-sample return series `oos_returns.csv`, the performance metrics `metrics.json`, a cumulative-return plot `cumulative_returns_plot.png`, the hyperparameters selected at each rebalancing date `selected_params.csv`, the asset and currency weights `weights_assets.csv` and `weights_currencies.csv`, and a configuration file `config.json` recording the run specification. Storing the configuration alongside its outputs makes each result reproducible and attributable to the settings that generated it.

4.2.2 Backtest Configuration

We now specify the configuration under which the considered strategies in the thesis are backtested. We first list the backtested strategies, then describe the chosen backtest specification, and finally the execution of the currency positions and the transaction costs.

We backtest the joint and the sequential currency overlay optimizers under the five risk measures considered in the thesis, variance, semi-variance, mean absolute deviation, Expected Shortfall and second-order Expected Shortfall, each in an unregularized, a regularized and a mean-risk formulation, as introduced in Chapter 3. As benchmarks we use the equally weighted portfolio and two regularized asset-only optimizers, Full and Zero, under the same five risk measures. EW and Full are evaluated on fully hedged returns, and the joint and overlay optimizers on fully hedged returns plus the excess currency returns of their currency positions. Zero optimizes the same objective as Full, but is evaluated on unhedged returns, so it keeps the full currency exposure of its asset positions. The following portfolio strategies are considered:

- **EW**: The fully hedged equally weighted portfolio.
- **Full**: The fully hedged regularized asset portfolio with both L_1 and L_2 penalties.
- **Zero**: The unhedged regularized asset portfolio with both L_1 and L_2 penalties.
- **Overlay**: The unregularized sequential currency overlay portfolio.
- **Regularized Overlay**: The overlay portfolio with both L_1 and L_2 penalties on the asset and currency weights.
- **Mean-risk Overlay**: The unregularized overlay portfolio with an added mean term.
- **Joint**: The unregularized joint optimization portfolio.
- **Regularized Joint**: The joint optimization portfolio with both L_1 and L_2 penalties on the asset and currency weights.

- **Mean-risk Joint**: The unregularized joint optimization portfolio with an added mean term.

Portfolios are rebalanced on the last trading day of each quarter. The first portfolio is formed on 30 June 2000, and the last on 30 December 2022, which gives 91 rebalancing dates. The lookback window covers the ten years before each rebalancing date, approximately 2,500 daily observations. All strategies invest in the constituents of the current year's universe that have a complete return history over the lookback window.

The penalty weights of the regularized optimizers are selected at every rebalancing date by a grid search, for which the lookback window is partitioned into a nine-year estimation window and a one-year validation window. Each configuration is fitted on the estimation window, and the resulting weights are scored by the risk measure they realize on the validation window. The configuration with the lowest validation risk is then re-fitted on both the estimation and validation windows, so that the final weights also use the most recent year of data. More details on the grid search and the hyperparameter range analysis are given in Appendix G.

The mean-risk optimizers also require expected asset and currency returns, which are estimated at each rebalancing date on the lookback window. Since individual stock means are hard to estimate accurately, we combine them by economy, as derived in Chapter 3. That is, the means of the equally weighted economy returns are shrunk toward zero with a James–Stein estimator, and every asset inherits the shrunk mean of its economy. The expected excess currency returns are not estimated from returns, but set to the negative forward premium at the end of the lookback window. Due to computational complexity considerations, the risk aversion coefficient γ is not selected on the validation window. Instead, we run each mean-risk optimizer over a grid of γ values on the first lookback window between June 1990 and June 2000 and fix one value for each risk measure. More details on the risk aversion parameter selection can be found in Appendix G.

The currency positions of the hedged strategies are implemented with three-month forward contracts. Since each contract expires at the next rebalancing date, the full forward position is traded at every rebalancing date rather than only its change, which is why the currency turnover (4.1) in Section 4.2.3 equals the notional of the forward positions. All returns are reported net of transaction costs. We assume constant half-spreads of 10 basis points for equities and 1 basis point for currencies, which gives transaction costs of $\tau^{asset} \approx 0.002$ and $\tau^{curr} \approx 0.0002$ per unit traded. We derive the transaction costs in detail in Appendix F.

4.2.3 Performance Evaluation Metrics

In Chapter 5 we evaluate the performance of all considered portfolio optimization strategies on a common set of performance metrics computed from their out-of-sample return series, each capturing a different view of the portfolio's performance. How well the joint and overlay programs reduce the risk measure they minimize is analyzed separately in Chapter 6. Throughout this section, let T be the number of trading days in the out-of-sample period, and let $R_{\mathcal{P},t}$ denote the realized daily out-of-sample portfolio return, $\bar{R}_{\mathcal{P}}$ its sample mean, $\sigma_{\mathcal{P}}$ its annualized volatility, $\sigma_{\mathcal{P},\text{down}}$ its annualized downside volatility, $R_{f,t}$ the daily risk-free rate and $\bar{R}_f$ the average daily risk-free rate. Daily figures are annualized using the standard 252 trading days per year assumption. The metrics we report are the annualized excess return, annualized volatility, maximum drawdown, the Sharpe ratio, the Sortino ratio, and average asset and currency turnover.

The annualized excess return is the compound annual growth rate of the portfolio in excess of

that of the risk-free rate

$$\mathrm{ER}_{\mathcal{P}} = \left(\prod_{t=1}^{T}(1+R_{\mathcal{P},t})\right)^{252/T} - \left(\prod_{t=1}^{T}(1+R_{f,t})\right)^{252/T}.$$

It is the annual return an investor holding the portfolio over the full sample would have earned on top of the risk-free rate. Since the risk-minimizing optimizers do not target returns, inspecting the excess return shows how much return they give up in exchange for lower risk.

Annualized volatility is the standard deviation of portfolio returns

$$\sigma_{\mathcal{P}} = \sqrt{252}\sqrt{\frac{1}{T-1}\sum_{t=1}^{T}\left(R_{\mathcal{P},t}-\bar{R}_{\mathcal{P}}\right)^2}.$$

It is the most common measure of risk used to evaluate portfolio performance, and is also used in the calculation of the Sharpe ratio to measure risk-adjusted performance. Because a portfolio can achieve low volatility while still suffering severe cumulative losses over time, it is most informative when interpreted together with the maximum drawdown.

Maximum drawdown is the largest decline of the cumulative return path of the portfolio. Unlike volatility, which summarizes a one-period return distribution, drawdown is a property of the realized wealth path, capturing the cumulative and sequential experience of losses that an investor goes through. Such path-dependent losses matter in practice. For example, deep drawdowns can breach margin requirements. Importantly, none of the optimizers considered in the thesis targets drawdown directly, so it serves as an out-of-sample robustness check. Let $V_{\mathcal{P},t} = \prod_{s\le t}(1+R_{\mathcal{P},s})$ denote the cumulative wealth of the portfolio, with $V_{\mathcal{P},0} = 1$, and $P_{\mathcal{P},t} = \max_{s\le t} V_{\mathcal{P},s}$ its running maximum, the maximum drawdown is defined as

$$\mathrm{MD}_{\mathcal{P}} = \max_{t}\frac{P_{\mathcal{P},t}-V_{\mathcal{P},t}}{P_{\mathcal{P},t}}.$$

The Sharpe ratio is the annualized ratio of the average excess portfolio return relative to its total volatility. It is the standard measure of return per unit of total risk. The Sharpe ratio is given by

$$\mathrm{SR}_{\mathcal{P}} = 252\,\frac{\bar{R}_{\mathcal{P}}-\bar{R}_f}{\sigma_{\mathcal{P}}}.$$

Since the risk-minimizing optimizers run at a much lower volatility than the equally weighted portfolio, their returns alone cannot be compared fairly. The Sharpe ratio corrects for this, and a strategy with a higher Sharpe ratio earns a higher return at any level of risk an investor chooses. Like volatility, however, it penalizes favorable and adverse deviations equally.

The Sortino ratio replaces volatility in the Sharpe ratio with downside volatility, so that a strategy is not penalized for upside dispersion. Read against the Sharpe ratio, it indicates whether a strategy's risk-adjusted performance stems from downside protection or from lower overall variability. The Sortino ratio is given by

$$\mathrm{SOR}_{\mathcal{P}} = 252\,\frac{\bar{R}_{\mathcal{P}}-\bar{R}_f}{\sigma_{\mathcal{P},\mathrm{down}}},$$

where the annualized downside volatility

$$\sigma_{\mathcal{P},\mathrm{down}} = \sqrt{252}\sqrt{\frac{1}{T}\sum_{t=1}^{T}\min\left(R_{\mathcal{P},t},\,0\right)^2}$$

only considers negative returns.

Finally, turnover proxies the trading activity and the implementation costs of each strategy. Although the reported returns are already net of transaction costs, turnover remains informative because it measures how much the portfolio must trade to maintain its positions, its capacity to scale, and its exposure to frictions such as slippage and market impact that our analysis does not capture. A strategy whose performance depends on heavy rebalancing can be hard to implement, so low and stable turnover is a desirable property. Turnover is reported separately for the asset and currency legs, since the two are traded through different instruments with different cost levels. Asset turnover is the average rebalancing volume over the out-of-sample period

$$\text{Turnover}^a = \frac{1}{D}\sum_{d=1}^{D}\sum_{i=1}^{N}|x_{i,d} - x_{i,d}^{-}|,$$

where D is the number of out-of-sample rebalancing dates, $x_{i,d}^{-}$ the weight of asset i just before the rebalance at date d, and $x_{i,d}$ the weight set at that rebalance. Currency turnover is the average absolute notional of the forward positions

$$\text{Turnover}^c = \frac{1}{D}\sum_{d=1}^{D}\sum_{c=1}^{M}|\phi_{c,d}|. \tag{4.1}$$

Chapter 5

Backtest Analysis

In this chapter, we analyze the out-of-sample performance of the joint and sequential overlay programs across the five risk measures and three universe sizes introduced in Chapters 3 and 4. In Section 5.1 we report the baseline backtests of the unregularized joint and overlay programs, together with the equally weighted benchmark. We show that the strategies trade excess return for lower volatility and drawdowns, that their risk-adjusted returns exceed those of the equally weighted portfolio in the medium and large universes but not consistently in the small universe, and that joint and sequential optimization produce similar realized performance across measures and universes. In Section 5.2 we repeat the analysis for the regularized joint and overlay programs and compare them against the regularized asset-only benchmarks, the fully hedged and unhedged portfolios. We show that regularization improves the Sharpe ratios, the most in the large universe, and that actively trading the currencies performs on par with full hedging. In Section 5.3 we consider the mean-risk formulation of the optimizers and show that adding a mean term substantially improves the risk-adjusted returns. We further show that the deviation measures benefit from the addition of a mean term more than the tail measures. Across the three considered backtests, the only metric where joint optimization consistently outperforms sequential optimization is volatility, i.e. risk reduction, where the joint program achieves lower realized volatility across all risk measures and universe sizes. This motivates the analysis of Chapter 6, where we focus on comparing the optimization approaches in terms of their risk reduction.

5.1 Baseline Backtest Results

We first report the baseline backtests of the unregularized Joint and Overlay optimizers across the five risk measures and three universe sizes considered in the thesis. We compare them against the equally weighted portfolio which serves as the no-optimization benchmark.

For the small, medium and large universes, we report in Table 5.1, the out-of-sample annualized excess return, annualized volatility, maximum drawdown, Sharpe ratio, Sortino ratio, average asset turnover and average currency turnover of the Overlay and Joint optimizers and of the equally weighted portfolio. All quantities are computed on the out-of-sample excess return series produced by the backtest configuration presented in Section 4.2.2. We also report in Figures 5.1 to 5.3 the cumulative out-of-sample returns of the Joint optimizers together with the equally weighted benchmark, and in the lower panel of each figure the cumulative difference between the Joint and Overlay optimizers for each risk measure. The latter helps us visualize the realized return difference between the two optimization approaches over time.

Baseline Backtest Performance across Universe Sizes

Panel A: Small Investable Universe (top 10 stocks per economy)											
		Variance		**Semi-Variance**		**MAD**		**ES**		**SOES**	
Metric	**EW**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**
ER (%)	1.77	1.12	1.25	0.67	0.61	1.26	1.39	0.61	0.62	0.69	0.19
Vol (%)	13.96	7.97	7.73	8.15	7.84	7.87	7.71	8.18	7.97	9.17	8.72
MD (%)	54.77	28.23	27.92	32.17	32.63	27.67	27.31	29.31	28.31	34.43	35.38
SR	0.13	0.14	0.16	0.08	0.08	0.16	0.18	0.08	0.08	0.07	0.02
SOR	0.17	0.19	0.22	0.11	0.11	0.22	0.25	0.10	0.11	0.10	0.03
Turnover^a	0.30	0.87	0.82	0.99	0.97	1.10	1.06	1.30	1.27	1.38	1.47
Turnover^c	0.92	1.70	1.88	2.13	2.38	1.51	1.84	2.39	2.98	2.66	3.46
Panel B: Medium Investable Universe (top 50 stocks per economy)											
		Variance		**Semi-Variance**		**MAD**		**ES**		**SOES**	
Metric	**EW**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**
ER (%)	4.09	2.41	2.53	2.37	2.54	2.75	2.65	3.16	2.95	2.39	2.70
Vol (%)	12.50	5.17	5.09	5.27	5.16	5.73	5.56	5.68	5.53	5.86	5.64
MD (%)	55.75	35.85	34.90	37.88	36.71	35.93	36.94	38.59	38.84	35.05	33.09
SR	0.33	0.47	0.50	0.45	0.49	0.48	0.48	0.56	0.53	0.41	0.48
SOR	0.44	0.62	0.67	0.60	0.66	0.69	0.68	0.75	0.73	0.55	0.66
Turnover^a	0.27	1.05	1.05	1.25	1.26	1.66	1.68	1.78	1.79	2.13	2.23
Turnover^c	0.92	1.16	1.37	1.20	1.36	0.98	1.25	0.97	1.67	1.02	1.85
Panel C: Large Investable Universe (top 100 stocks per economy)											
		Variance		**Semi-Variance**		**MAD**		**ES**		**SOES**	
Metric	**EW**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**
ER (%)	4.26	2.74	2.79	2.73	2.76	2.44	2.40	3.28	3.44	2.39	2.46
Vol (%)	11.66	4.24	4.19	4.26	4.20	4.63	4.57	4.57	4.55	4.67	4.66
MD (%)	56.26	26.72	26.37	28.74	28.29	31.25	31.42	31.32	30.46	31.29	30.31
SR	0.37	0.65	0.66	0.64	0.66	0.53	0.52	0.72	0.76	0.51	0.53
SOR	0.49	0.85	0.88	0.84	0.87	0.70	0.70	0.95	1.01	0.68	0.70
Turnover^a	0.28	1.09	1.09	1.34	1.36	1.76	1.77	1.89	1.91	2.15	2.18
Turnover^c	0.91	1.06	1.24	1.07	1.21	0.97	1.16	0.97	1.43	0.99	1.62

Table 5.1: Performance metrics of the baseline backtests in the small, medium and large investable universes

Table 5.1 reports out-of-sample performance metrics for the unregularized optimizers in the three investable universes, with portfolios rebalanced quarterly. Panel A covers the small, Panel B the medium and Panel C the large investable universe. Each panel reports the unregularized sequential currency overlay (Overlay) and joint (Joint) optimizers, together with the equally weighted portfolio (EW). Columns are grouped by the risk measure that is minimized, that is, variance, semi-variance, mean absolute deviation (MAD), Expected Shortfall (ES) and second-order Expected Shortfall (SOES). The realized out-of-sample performance metrics shown are the annualized excess return (ER), volatility (Vol), maximum drawdown (MD), the Sharpe ratio (SR), the Sortino ratio (SOR), and average asset and currency turnover, respectively. The sample period runs from June 1990 to June 2023, and the first out-of-sample portfolio is formed in June 2000. More details on the optimization approaches, backtest configuration, and performance metrics are provided in Chapters 3 and 4.

Across all five risk measures and universe sizes, the Overlay and Joint optimizers achieve substantially lower volatility than the equally weighted portfolio. In the small universe, their volatility ranges from 7.71% to 9.17%, against 13.96% for the equally weighted portfolio, a reduction of 34% to 45%. In the medium and large universes, the reduction is even larger, by up to 59% and 64%, respectively. The equally weighted portfolio ignores the dependence between asset returns. Since it already holds a broad cross-section of stocks, most idiosyncratic risk is diversified away, and further adding stocks does little to reduce the remaining market risk,

which can only be addressed by taking the correlation structure of the assets and currencies into account. The optimizers instead tilt the portfolio toward assets with low risk, and a larger universe offers more such assets to choose from. As a result, maximum drawdowns are also substantially reduced, by up to 27 percentage points in the small, 23 in the medium and 30 in the large universe, relative to drawdowns of 54.77% to 56.26% for the equally weighted portfolio. Large drawdowns occur in market crashes, when the correlations between assets rise and diversification offers less protection. The optimizers reduce the portfolio's exposure to these crashes, but cannot completely mitigate them, since in a crash even low-risk equities tend to fall together, which is why the drawdowns remain sizeable for all strategies.

For excess returns, we observe the opposite effect. The Overlay and Joint optimizers deliver lower excess returns than the equally weighted portfolio for every risk measure and universe size. This is partially due to the fact that the optimizers are minimizing risk without any explicit return term in the objective function. We address this in more detail in Section 5.3. Secondly, since much of the volatility reduction achieved by the risk minimizing portfolios comes from a lower exposure to the common movement of the equity markets (i.e. lower market beta), the optimizers also earn a smaller share of the equity risk premium. But excess returns alone say little about performance when the portfolios carry different levels of risk, so we compare them on a risk-adjusted basis.

In the small universe, the Sharpe ratios of the optimizers range from 0.02 to 0.18, against 0.13 for the equally weighted portfolio, and only the variance and mean absolute deviation optimizers improve on the benchmark. In the medium and large universes, every optimizer outperforms the equally weighted portfolio, with Sharpe ratios between 0.41 and 0.56 against 0.33, and between 0.51 and 0.76 against 0.37, respectively. In the small universe, the optimizers lose more excess return than they reduce risk, while in larger universes they can diversify across more low-risk stocks, and the reduction in volatility outweighs the lost return. Sortino ratios follow the same pattern, so accounting only for downside risk does not change the comparison.

Both the Overlay and Joint optimizers have much higher asset and currency turnover than the equally weighted portfolio. This is expected since the equally weighted portfolio has fixed target weights and only trades to undo the price drift. Comparing the asset turnover of the Joint and Overlay optimizers, we see that it is very similar, differing at most by 6%. This tells us that the asset weights are determined primarily by the risk measure being optimized and not by the optimization approach itself. The difference between the two optimizers is in currency turnover, where Joint is higher than Overlay in every measure-universe combination, by between 11% and 82%. Since returns are reported net of transaction costs, the comparisons above already incorporate these differences. The extra currency turnover does not reduce the risk-adjusted returns of Joint relative to Overlay, and the higher turnover of both optimizers does not change their advantage over the equally weighted portfolio.

Finally, comparing the performance of the unregularized Joint and Overlay optimizers under the same risk measures, across all reported metrics, we see that they perform very similarly. The Joint optimizer does consistently achieve a lower volatility, but the differences shrink as we increase the universe size, from up to a 7% decrease in the small universe, to essentially no difference in the large universe. This is in line with the joint program being able to exploit the asset and currency dependence structure when the number of assets compared to the number of currencies is smaller. We study the risk reduction performance and mechanism in more detail in Chapter 6. The other metrics differ by small amounts, with excess returns differing by at most 0.5 percentage points, maximum drawdowns by at most 1.96 percentage points and Sharpe ratios by at most 0.07. However, the sign of these differences changes across risk measures within every universe, so one optimization approach does not consistently

outperform the other.

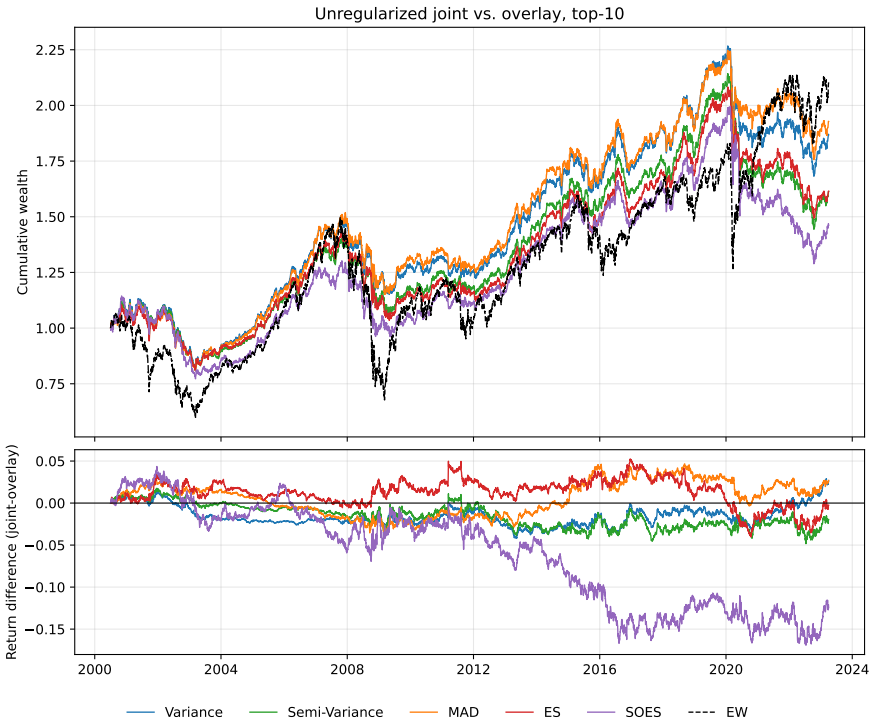


Figure 5.1: Cumulative performance of the unregularized joint optimizers and their return difference to the overlay optimizers in the small investable universe (top 10 stocks per economy).

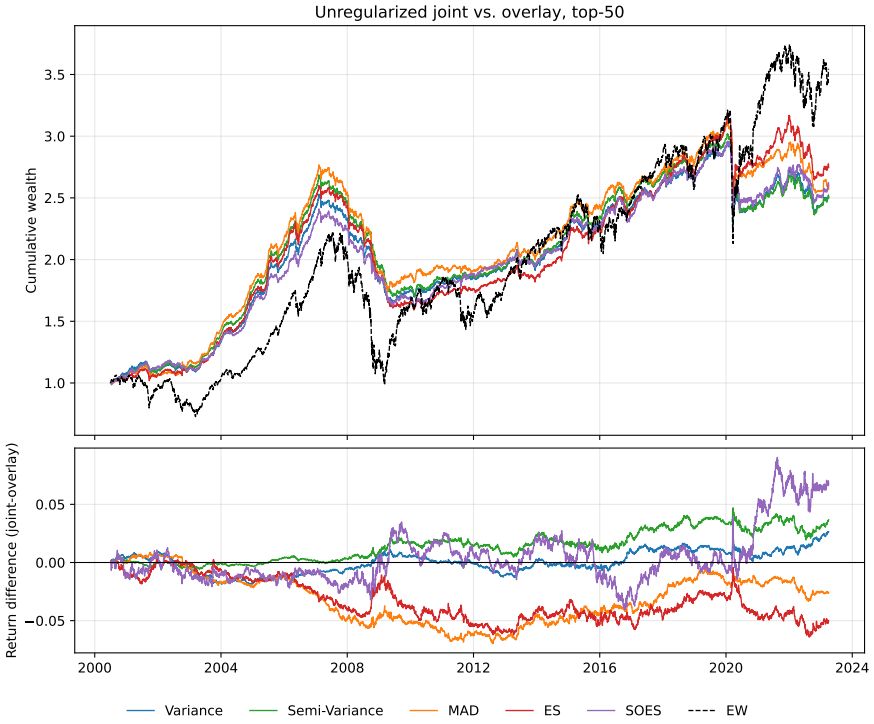


Figure 5.2: Cumulative performance of the unregularized joint optimizers and their return difference to the overlay optimizers in the medium investable universe (top 50 stocks per economy).

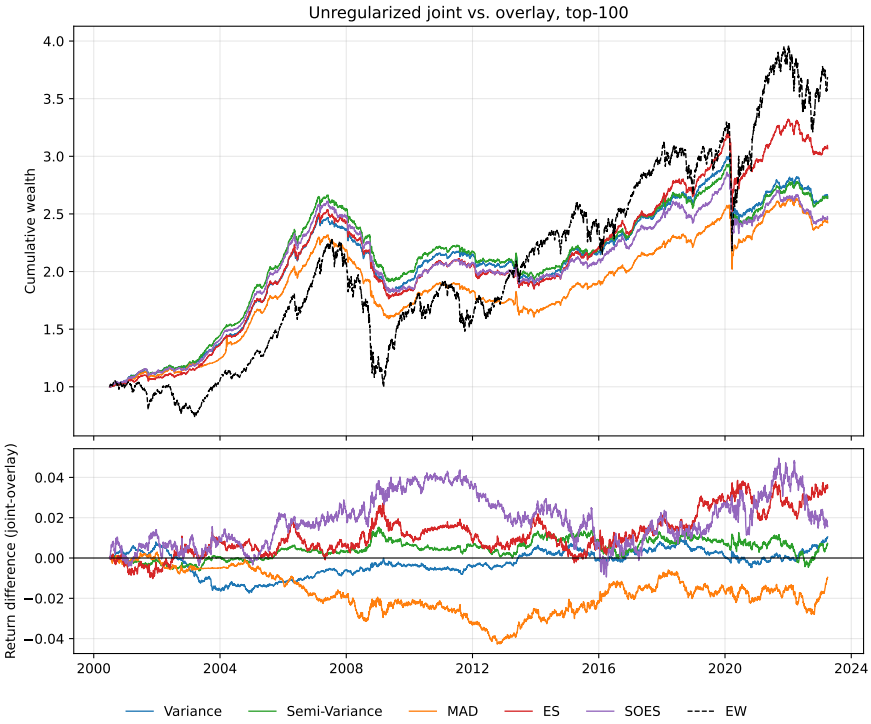


Figure 5.3: Cumulative performance of the unregularized joint optimizers and their return difference to the overlay optimizers in the large investable universe (top 100 stocks per economy).

Figures 5.1 to 5.3 show the out-of-sample performance of the unregularized optimizers in the small, medium and large investable universes, respectively. In each figure, the top panel shows the cumulative wealth of the joint optimizers and the equally weighted portfolio, and the bottom panel shows the cumulative sum of the daily return differences between the joint and the overlay optimizers. The portfolios are rebalanced quarterly and reported net of transaction costs. The corresponding performance metrics are reported in Panels A, B and C of Table 5.1, respectively.

Figures 5.1 to 5.3 present the cumulative performance of the Overlay and Joint optimizers alongside the equally weighted benchmark. In each universe, the top panel shows that the Joint optimizers perform better than the equally weighted portfolio during the first major market crashes in the sample, the dot-com crash and the global financial crisis. This advantage is particularly evident in the medium and large universes during the global financial crisis, where the optimizer strategies exhibit a much smaller drawdown compared to the equally weighted portfolio. The equally weighted portfolio recovers quickly after the COVID-19 market shock and ultimately attains the highest terminal wealth across all three universes, albeit at considerably higher volatility and with deeper drawdowns throughout the sample. In contrast, the Overlay and Joint optimizers follow more stable growth paths, consistent with their lower realized volatility reported in Table 5.1.

The bottom panels show the cumulative daily return difference between the Joint and Overlay optimizers for each risk measure, and confirm that the two optimization approaches track each other closely throughout the sample in terms of returns. The cumulative differences oscilate around zero in every universe size and for every measure, with the exception of second-order Expected Shortfall in the small universe, where the Joint optimizer drifts persistently below the Overlay optimizer from the early 2010s onward.

Overall, we see a consistent pattern across all the considered risk measures. The Overlay and Joint optimizers reduce volatility and drawdowns substantially relative to the equally weighted portfolio. The Sharpe comparison depends on the universe size. In the small

universe, the Joint and Overlay optimizers do not consistently outperform the equally weighted portfolio, while in the medium and large universes they do. The Joint and Overlay optimizers produce very similar out-of-sample metrics across all risk measures and universe sizes, with Joint consistently delivering lower volatility and higher currency turnover but no consistent advantage on risk-adjusted returns.

5.2 Regularization

We now repeat the analysis of Section 5.1 for the regularized versions of the Joint and Overlay optimizers. We add L_1 and L_2 penalties on the asset weights and, separately, on the currency weights to the objective, following the regularization scheme of Chapter 3. The penalty weights are selected at every rebalancing date by a grid search over the validation window described in Section 4.2.2. All other aspects of the backtests remain unchanged. We also compare the Joint and Overlay optimizers against the regularized Assets-Full (Full) and Assets-Zero (Zero) optimizers which serve as benchmarks for optimizing only the asset weights under the same risk measures, with the currency weights either fully hedged or set to zero, respectively. Table 5.2 reports the change in the Sharpe ratio between the unregularized and regularized optimizers, and Tables 5.3 to 5.5 report performance metrics of the regularized optimizers and the three benchmarks.

Change in Sharpe Ratio from Regularization

	Overlay			**Joint**		
Risk Measure	**Small**	**Medium**	**Large**	**Small**	**Medium**	**Large**
Variance	+0.02	−0.02	+0.12	0.00	−0.05	+0.10
Semi-Variance	+0.05	+0.02	+0.19	+0.03	−0.03	+0.18
MAD	+0.04	−0.02	+0.02	+0.03	+0.01	+0.04
ES	+0.05	−0.01	+0.15	+0.09	+0.02	+0.13
SOES	+0.04	+0.12	+0.38	−0.02	+0.08	+0.38
Average	+0.04	+0.02	+0.17	+0.03	+0.01	+0.17

Table 5.2: Change in the out-of-sample Sharpe ratio of the Overlay and Joint optimizers from regularization

Table 5.2 reports the change in the out-of-sample Sharpe ratio from regularization, computed as the Sharpe ratio of each regularized optimizer in Tables 5.3 to 5.5 minus that of its unregularized version in Table 5.1, for the Overlay and Joint optimizers in the small, medium and large investable universes. Rows correspond to the risk measure that is minimized, and the last row reports the average change across the five risk measures.

We see in Table 5.2 that regularization improves the Sharpe ratios of the Overlay and Joint optimizers for most risk measures and universe sizes. The improvement is the largest in the large universe, where the Sharpe ratio rises by 0.17 on average, while the effect is modest in the small universe and mixed in the medium universe. The penalties stabilize the weights by limiting the extreme positions that the unregularized programs fit to noise in the lookback window, which becomes more important when the number of assets increases relative to the length of the lookback window. The unregularized second-order Expected Shortfall optimizers are the least stable, since the measure is driven by the squared losses of the few most extreme observations in the lookback window. Accordingly, regularization benefits them the most, raising their Sharpe ratios by 0.38 in the large universe.

Regularized Backtest Performance: Small Universe (top 10 stocks per economy)

Panel A1: Regularized Overlay and Joint Optimizers

	Variance		Semi-Variance		MAD		ES		SOES	
Metric	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**
ER (%)	1.29	1.24	1.05	0.84	1.60	1.64	1.07	1.31	0.98	0.03
Vol (%)	7.82	7.76	7.91	7.72	7.81	7.75	7.99	7.85	8.52	8.64
MD (%)	27.81	27.80	30.63	30.43	27.83	26.01	29.96	28.89	36.02	36.22
SR	0.16	0.16	0.13	0.11	0.20	0.21	0.13	0.17	0.11	0.00
SOR	0.22	0.22	0.18	0.15	0.28	0.29	0.18	0.23	0.16	0.00
Turnover^a	0.93	0.96	1.09	1.01	1.04	1.06	1.07	1.17	1.40	1.59
Turnover^c	1.43	1.47	1.72	1.77	1.36	1.44	1.70	1.77	2.04	2.12

Panel B1: Benchmarks

		Variance		Semi-Variance		MAD		ES		SOES	
Metric	**EW**	**Full**	**Zero**	**Full**	**Zero**	**Full**	**Zero**	**Full**	**Zero**	**Full**	**Zero**
ER (%)	1.77	1.05	2.64	1.08	2.65	1.66	3.13	0.85	2.44	0.55	1.93
Vol (%)	13.96	8.39	11.80	8.41	11.72	8.36	11.70	8.51	11.84	9.12	12.35
MD (%)	54.77	28.00	42.18	29.21	41.12	28.59	40.88	30.81	44.28	33.06	44.84
SR	0.13	0.12	0.22	0.13	0.23	0.20	0.27	0.10	0.21	0.06	0.16
SOR	0.17	0.17	0.30	0.17	0.31	0.27	0.36	0.13	0.28	0.08	0.21
Turnover^a	0.30	1.02	1.02	1.16	1.17	1.08	1.08	1.27	1.23	1.66	1.67
Turnover^c	0.92	0.93	0	0.93	0	0.92	0	0.93	0	0.94	0

Table 5.3: Performance metrics of the regularized optimizers in the small investable universe (top 10 stocks per economy)

Regularized Backtest Performance: Medium Universe (top 50 stocks per economy)

Panel A2: Regularized Overlay and Joint Optimizers

	Variance		Semi-Variance		MAD		ES		SOES	
Metric	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**
ER (%)	2.38	2.35	2.41	2.33	2.83	2.95	3.16	3.08	2.91	3.10
Vol (%)	5.26	5.19	5.16	5.08	6.16	6.08	5.74	5.58	5.52	5.52
MD (%)	35.03	34.25	36.39	34.57	35.93	35.52	38.57	36.38	30.41	30.20
SR	0.45	0.45	0.47	0.46	0.46	0.49	0.55	0.55	0.53	0.56
SOR	0.59	0.59	0.62	0.61	0.75	0.79	0.77	0.76	0.71	0.76
Turnover^a	1.14	1.13	1.24	1.27	1.31	1.32	1.53	1.57	1.56	1.97
Turnover^c	1.11	1.11	1.18	1.17	1.09	1.06	1.26	1.30	1.39	1.32

Panel B2: Benchmarks

		Variance		Semi-Variance		MAD		ES		SOES	
Metric	**EW**	**Full**	**Zero**	**Full**	**Zero**	**Full**	**Zero**	**Full**	**Zero**	**Full**	**Zero**
ER (%)	4.09	2.52	3.86	2.55	3.92	2.94	3.84	3.26	4.54	3.02	4.48
Vol (%)	12.50	5.31	9.58	5.22	9.43	6.19	9.97	5.84	9.87	5.60	9.78
MD (%)	55.75	35.05	42.05	36.53	41.31	35.90	43.81	38.54	42.37	29.76	36.74
SR	0.33	0.48	0.40	0.49	0.42	0.48	0.38	0.56	0.46	0.54	0.46
SOR	0.44	0.62	0.55	0.65	0.57	0.77	0.57	0.78	0.64	0.72	0.63
Turnover^a	0.27	1.07	1.07	1.34	1.35	1.30	1.30	1.60	1.60	1.81	1.81
Turnover^c	0.92	0.95	0	0.95	0	0.96	0	0.95	0	0.95	0

Table 5.4: Performance metrics of the regularized optimizers in the medium investable universe (top 50 stocks per economy)

Regularized Backtest Performance: Large Universe (top 100 stocks per economy)

Panel A3: Regularized Overlay and Joint Optimizers										
	Variance		Semi-Variance		MAD		ES		SOES	
Metric	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**
ER (%)	3.33	3.25	3.56	3.52	2.72	2.73	3.88	3.91	4.05	4.09
Vol (%)	4.35	4.29	4.27	4.19	4.97	4.85	4.45	4.40	4.57	4.52
MD (%)	26.27	26.40	27.33	27.29	31.28	31.63	30.61	30.11	25.84	24.70
SR	0.77	0.76	0.83	0.84	0.55	0.56	0.87	0.89	0.89	0.91
SOR	0.99	0.98	1.09	1.10	0.73	0.75	1.14	1.17	1.17	1.20
Turnover[a]	1.14	1.10	1.40	1.33	1.60	1.56	1.42	1.64	1.65	1.87
Turnover[c]	1.06	1.03	1.06	1.06	1.01	1.03	1.15	1.23	1.23	1.26

Panel B3: Benchmarks											
		Variance		Semi-Variance		MAD		ES		SOES	
Metric	**EW**	**Full**	**Zero**	**Full**	**Zero**	**Full**	**Zero**	**Full**	**Zero**	**Full**	**Zero**
ER (%)	4.26	3.16	4.40	3.56	4.94	2.83	4.13	4.00	5.28	3.75	4.92
Vol (%)	11.66	4.37	8.99	4.28	8.92	4.94	9.36	4.46	9.21	4.57	9.28
MD (%)	56.26	26.60	33.06	27.84	34.27	31.28	36.99	31.60	35.12	30.41	35.50
SR	0.37	0.72	0.49	0.83	0.55	0.57	0.44	0.90	0.57	0.82	0.53
SOR	0.49	0.93	0.68	1.09	0.77	0.77	0.62	1.18	0.80	1.09	0.74
Turnover[a]	0.28	1.10	1.10	1.35	1.35	1.59	1.59	1.61	1.62	1.84	1.84
Turnover[c]	0.91	0.96	0	0.97	0	0.97	0	0.96	0	0.97	0

Table 5.5: Performance metrics of the regularized optimizers in the large investable universe (top 100 stocks per economy)

Tables 5.3 to 5.5 report out-of-sample performance metrics for the regularized optimizers in the three investable universes, with portfolios rebalanced quarterly. Panels A1, A2 and A3 report the regularized sequential currency overlay (Overlay) and joint (Joint) optimizers, and Panels B1, B2 and B3 the benchmarks, the equally weighted portfolio (EW) and the asset-only optimizers Assets-Full (Full) and Assets-Zero (Zero). Columns are grouped by the risk measure that is minimized, that is, variance, semi-variance, mean absolute deviation (MAD), Expected Shortfall (ES) and second-order Expected Shortfall (SOES). The realized out-of-sample performance metrics shown are the annualized excess return (ER), volatility (Vol), maximum drawdown (MD), the Sharpe ratio (SR), the Sortino ratio (SOR), and average asset and currency turnover, respectively. The sample period runs from June 1990 to June 2023, and the first out-of-sample portfolio is formed in June 2000. More details on the optimization approaches, backtest configuration, and performance metrics are provided in Chapters 3 and 4.

Comparing the performance of the Full optimizer with Overlay and Joint optimizers in Tables 5.3 to 5.5, we see that fully hedging the currency exposure performs very similarly to actively trading the currencies. Full solves the same asset-only problem as the first step of the overlay program, and the regularized Overlay and Joint optimizers achieve Sharpe ratios on par with Full in every universe. A likely reason is that the objective contains no mean term, so the currency weights are chosen only to reduce risk and not to earn the excess currency returns. Volatility in the small universe is where the Joint and Overlay outperform the Full optimizer, by up to 10%. Looking at the currency turnover, we see that active currency trading is the mechanism through which this is achieved. In the medium and large universes, the currency turnover of the Overlay and Joint optimizers is only 4% to 46% higher than that of Full, so their currency positions deviate only moderately from a full hedge, whereas in the small universe it is 1.5 to 2.3 times as high. On the other hand, the Zero optimizer keeps the full currency exposure and with it the excess currency returns that the risk-minimizing optimizers give up, and earns a higher excess return at a much higher volatility. In the small universe this gives Zero the higher Sharpe ratio for every risk measure, while in the medium and large universes the Overlay and Joint optimizers outperform it for every risk measure, by

up to 70%. This performance advantage of the risk-minimizing optimizers over Zero shows the benefits of active currency management in larger universes, where equity risk is reduced further and currency risk makes up a larger share of the portfolio's total risk.

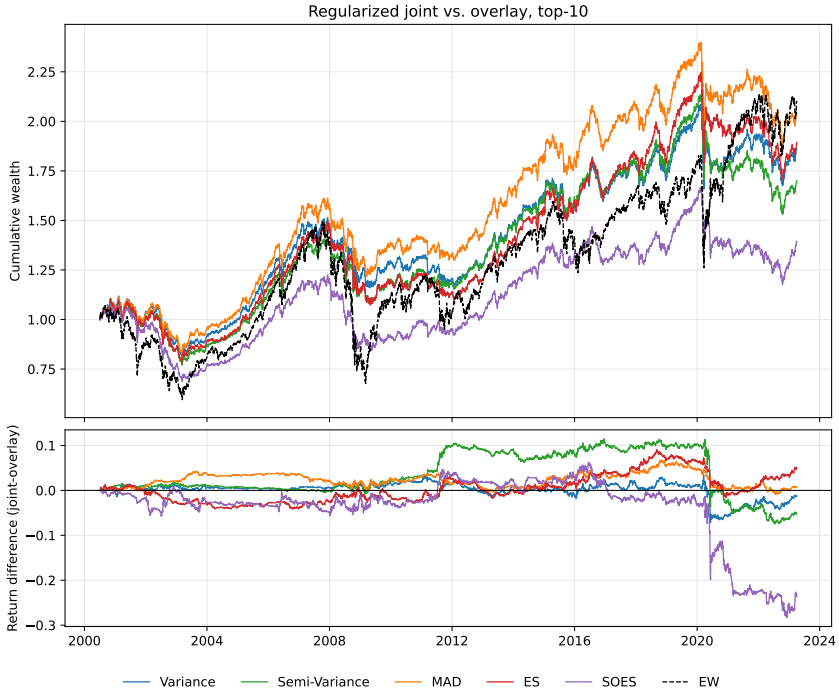


Figure 5.4: Cumulative performance of the regularized joint optimizers and their return difference to the overlay optimizers in the small investable universe (top 10 stocks per economy).

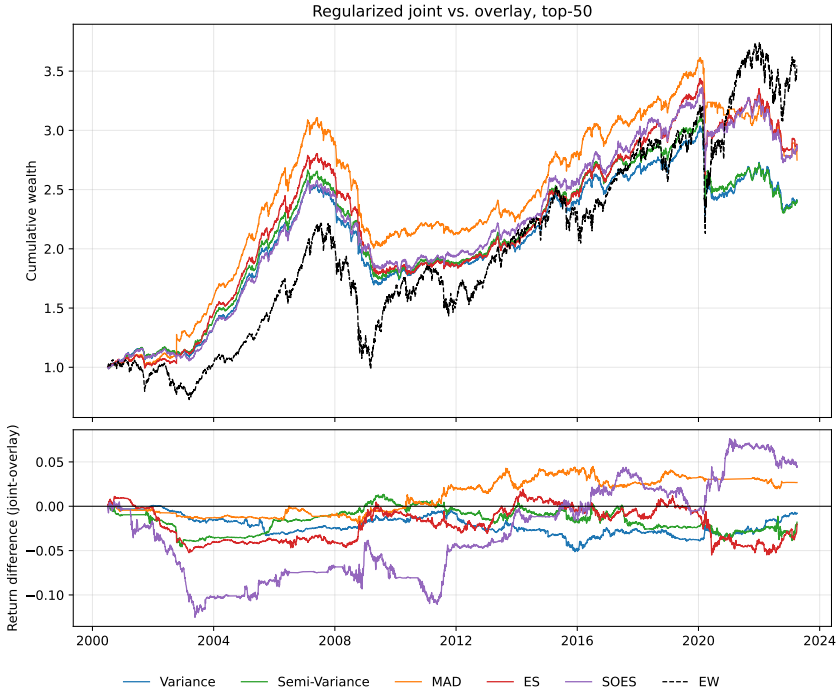


Figure 5.5: Cumulative performance of the regularized joint optimizers and their return difference to the overlay optimizers in the medium investable universe (top 50 stocks per economy).

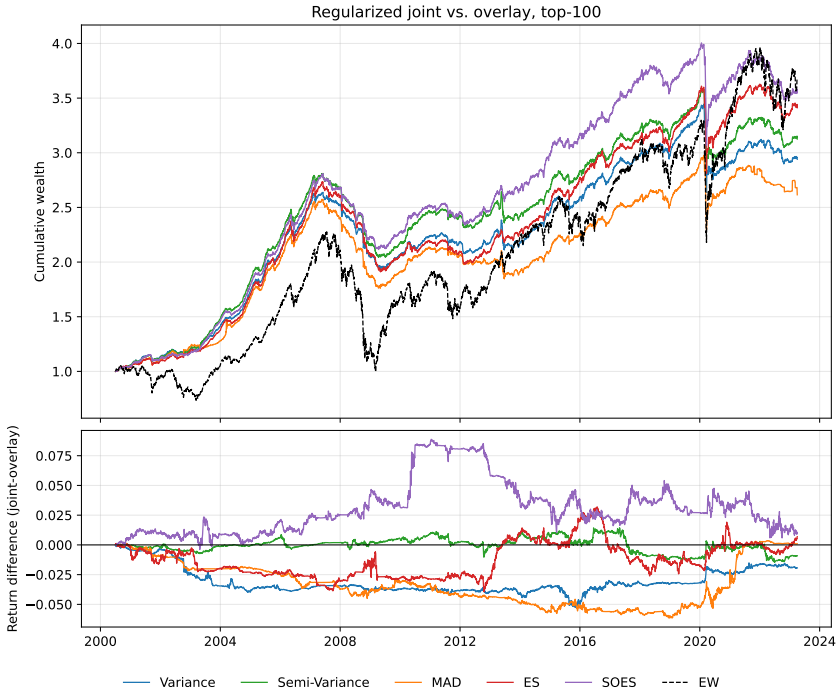


Figure 5.6: Cumulative performance of the regularized joint optimizers and their return difference to the overlay optimizers in the large investable universe (top 100 stocks per economy).

Figures 5.4 to 5.6 show the out-of-sample performance of the regularized optimizers in the small, medium and large investable universes, respectively. In each figure, the top panel shows the cumulative wealth of the joint optimizers and the equally weighted portfolio, and the bottom panel shows the cumulative sum of the daily return differences between the joint and the overlay optimizers. The portfolios are rebalanced quarterly and reported net of transaction costs. The corresponding performance metrics are reported in Tables 5.3 to 5.5.

After regularization, the performance of the Joint and Overlay optimizers remains very similar. Their Sharpe ratios in Tables 5.3 to 5.5 differ only marginally for every risk measure and universe size, with the sign of the difference changing across risk measures within each universe. The Joint optimizers again achieve a lower volatility in every measure-universe combination. The only exception being second-order Expected Shortfall in the small universe, where Overlay beats Joint on every metric. The bottom panel of Figure 5.4 shows that this happens almost entirely after the onset of the COVID-19 crisis in 2020, so it is more reflective of this period rather than a persistent difference. The bottom panels of Figures 5.5 and 5.6 show that the cumulative return differences in the medium and large universes remain small for all risk measures throughout the sample. The main effect of regularization on the comparison is on currency turnover. Without regularization, the currency turnover of Joint exceeds that of Overlay by 11% to 82%, while with regularization the two differ by at most 7%. The penalties on the currency weights thus shrink the larger currency positions of the joint program to about the size of those of the overlay program. This does not have an effect on their risk-adjusted returns, which remain very similar, with no optimizer consistently outperforming the other.

The results so far share a common feature. All optimizers minimize a risk measure without any return term in the objective, and while this delivers substantially lower volatility and drawdowns than the equally weighted portfolio, it also gives up potential excess asset and currency returns. A natural extension is therefore to add a mean term to the objective, which lets the optimizers trade off expected return against risk. A risk-return objective is also where joint optimization could be more beneficial, as the joint program can account for the

dependence structure between asset and currency returns when choosing positions that raise the expected return.

5.3 Mean-Risk Optimization

In this section, we add a mean term to the objective of the unregularized joint and overlay programs of Section 5.1. Each program now minimizes the risk measure with an added expected portfolio return scaled by $1/\gamma$, where γ is a risk aversion coefficient, as introduced in Section 3.2. In the overlay program, the asset step includes the expected asset returns and the currency step the expected excess currency returns. The expected asset returns are averaged across the stocks of each economy and shrunk toward zero, while the expected excess currency returns are set to the negative forward premium, that is, the difference between the foreign and the US interest rate. The choice of γ for each risk measure is described in Appendix **??**. All other aspects of the backtests remain unchanged. Table 5.6 reports the change in the Sharpe ratio from adding the mean term, and Table 5.7 reports the performance metrics.

Change in Sharpe Ratio from Adding a Mean Term

	Overlay			Joint		
Risk Measure	**Small**	**Medium**	**Large**	**Small**	**Medium**	**Large**
Variance	+0.23	+0.24	+0.22	+0.21	+0.24	+0.25
Semi-Variance	+0.30	+0.16	+0.12	+0.30	+0.17	+0.18
MAD	+0.18	+0.15	+0.14	+0.19	+0.21	+0.18
ES	+0.04	−0.01	0.00	+0.05	+0.05	+0.04
SOES	+0.09	+0.04	+0.02	+0.10	+0.14	+0.08
Deviation measures	+0.24	+0.18	+0.16	+0.23	+0.21	+0.20
Tail measures	+0.07	+0.02	+0.01	+0.08	+0.10	+0.06

Table 5.6: Change in the out-of-sample Sharpe ratio of the Overlay and Joint optimizers from adding a mean term

Table 5.6 reports the change in the out-of-sample Sharpe ratio from adding a mean term to the objective, computed as the Sharpe ratio of each mean-risk optimizer in Table 5.7 minus that of the unregularized version in Table 5.1, for the Overlay and Joint optimizers in the small, medium and large investable universes. Rows correspond to the risk measure in the objective, that is, variance, semi-variance, mean absolute deviation (MAD), Expected Shortfall (ES) and second-order Expected Shortfall (SOES), and the last two rows report the average change across the deviation measures (variance, semi-variance and MAD) and across the tail measures (ES and SOES).

Looking at Table 5.6, we see that adding a mean term to the objective raises the Sharpe ratios of the Overlay and Joint optimizers in 28 of the 30 combinations of risk measure and universe size. Unlike regularization, which was most beneficial in the large universe, the addition of a mean term improves the Sharpe ratios by a similar amount in all three universes. Minimizing risk alone ignores expected returns, and the mean term recovers part of the excess return the optimizers otherwise give up. Though we can see that the improvement is very uneven across risk measures. The Sharpe ratios of the deviation measures rise by 0.16 to 0.24 on average, while those of the tail measures rise by at most 0.10.

This split is in line with how the risk measures handle the mean of the portfolio return. Expected Shortfall and second-order Expected Shortfall are translation equivariant, so minimizing them already rewards portfolios with a higher mean. The two tail measures therefore carry an implicit mean term, which the explicit one only strengthens. The deviation measures are

invariant to a shift in the portfolio return, so for them the addition of a mean term introduces new information in the optimization problem. The average currency turnover is consistent with this interpretation. In the medium and large universes, the mean term raises the currency turnover of the variance and semi-variance optimizers to two to three times its previous level, while for the tail-measure optimizers it remains stable around its previous level. We analyze the implicit risk-return tradeoff of the tail measures in more detail in Section 6.2.

Mean-Risk Backtest Performance across Universe Sizes

Panel A: Small Investable Universe (top 10 stocks per economy)											
		Variance		**Semi-Variance**		**MAD**		**ES**		**SOES**	
Metric	**EW**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**
ER (%)	1.77	3.21	3.13	3.75	3.75	2.81	3.05	0.96	1.01	1.51	1.03
Vol (%)	13.96	8.56	8.40	9.85	9.81	8.31	8.13	8.20	7.95	9.19	8.75
MD (%)	54.77	26.41	25.91	30.52	30.17	26.52	25.58	29.58	27.57	33.77	37.06
SR	0.13	0.37	0.37	0.38	0.38	0.34	0.37	0.12	0.13	0.16	0.12
SOR	0.17	0.51	0.51	0.53	0.53	0.46	0.51	0.16	0.17	0.23	0.16
Turnover^a	0.30	0.94	0.92	1.17	1.22	1.13	1.13	1.30	1.28	1.39	1.48
Turnover^c	0.92	2.36	2.41	2.65	2.61	2.14	2.25	2.15	2.35	2.25	2.42
Panel B: Medium Investable Universe (top 50 stocks per economy)											
		Variance		**Semi-Variance**		**MAD**		**ES**		**SOES**	
Metric	**EW**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**
ER (%)	4.09	4.54	4.66	4.78	5.21	3.59	3.82	3.15	3.26	2.67	3.54
Vol (%)	12.50	6.37	6.33	7.79	7.95	5.68	5.57	5.68	5.58	5.89	5.70
MD (%)	55.75	37.66	36.85	40.93	41.50	36.19	37.92	39.52	38.75	34.11	32.47
SR	0.33	0.71	0.74	0.61	0.66	0.63	0.69	0.55	0.58	0.45	0.62
SOR	0.44	0.96	1.00	0.83	0.90	0.87	0.95	0.75	0.79	0.61	0.86
Turnover^a	0.27	1.17	1.21	1.43	1.52	1.64	1.69	1.78	1.80	2.13	2.25
Turnover^c	0.92	2.62	2.71	3.15	3.10	1.60	2.01	0.99	1.72	1.03	1.92
Panel C: Large Investable Universe (top 100 stocks per economy)											
		Variance		**Semi-Variance**		**MAD**		**ES**		**SOES**	
Metric	**EW**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**	**Overlay**	**Joint**
ER (%)	4.26	5.02	5.26	5.57	6.39	3.20	3.25	3.30	3.66	2.51	2.90
Vol (%)	11.66	5.73	5.76	7.36	7.59	4.79	4.68	4.59	4.57	4.74	4.72
MD (%)	56.26	31.32	30.62	37.45	37.04	31.43	33.29	31.65	30.47	31.65	29.57
SR	0.37	0.87	0.91	0.76	0.84	0.67	0.70	0.72	0.80	0.53	0.61
SOR	0.49	1.17	1.23	1.03	1.16	0.89	0.91	0.95	1.07	0.70	0.82
Turnover^a	0.28	1.24	1.29	1.56	1.65	1.74	1.85	1.89	1.92	2.14	2.20
Turnover^c	0.91	2.88	2.95	3.40	3.36	1.37	1.94	0.97	1.49	1.00	1.80

Table 5.7: Performance metrics of the mean-risk optimizers in the small, medium and large investable universes

Table 5.7 reports out-of-sample performance metrics for the mean-risk optimizers in the three investable universes, with portfolios rebalanced quarterly. Panel A covers the small, Panel B the medium and Panel C the large investable universe. Each panel reports the unregularized sequential currency overlay (Overlay) and joint (Joint) optimizers with an added mean term, together with the equally weighted portfolio (EW) as a benchmark. Columns are grouped by the risk measure, that is, variance, semi-variance, mean absolute deviation (MAD), Expected Shortfall (ES) and second-order Expected Shortfall (SOES). The realized out-of-sample performance metrics shown are the annualized excess return (ER), volatility (Vol), maximum drawdown (MD), the Sharpe ratio (SR), the Sortino ratio (SOR), and average asset and currency turnover, respectively. The sample period runs from June 1990 to June 2023, and the first out-of-sample portfolio is formed in June 2000. More details on the optimization approaches, backtest configuration, and performance metrics are provided in Chapters 3 and 4.

As we can see in Table 5.7, the higher Sharpe ratios come from substantially higher excess returns achieved by the mean-risk optimizers. The excess returns are increased by 1.9 to 3.6 percentage points for the variance and semi-variance optimizers and by 0.8 to 1.7 percentage points for the mean absolute deviation optimizers, but by at most 0.8 percentage points for the tail measures. This comes at the cost of an increase in volatility, which rises by 0.6 to 3.4 percentage points for the variance and semi-variance optimizers, while for mean absolute deviation and the tail measures it changes by less than half a percentage point. Even with the increase, the volatility of the variance and semi-variance optimizers remains at between 50% and 70% of that of the equally weighted portfolio, and the maximum drawdowns of all mean-risk optimizers stay below 70% of that of the equally weighted portfolio. The mean term thus raises returns without giving up the risk reduction that distinguishes the optimizers from the equally weighted portfolio. As a result, every mean-risk optimizer outperforms the equally weighted portfolio on a risk-adjusted basis in the medium and large universes. In the small universe, the deviation measures reach Sharpe ratios of 0.34 to 0.38 against 0.13 for the equally weighted portfolio, while the tail measures perform similarly at 0.12 to 0.16. For the deviation measures, the underperformance of the pure risk-minimizing optimizers in the small universe is therefore not a failure of risk reduction, but a consequence of ignoring expected returns. The Sortino ratios show the same pattern, so a comparison focused on downside risk does not change the conclusion.

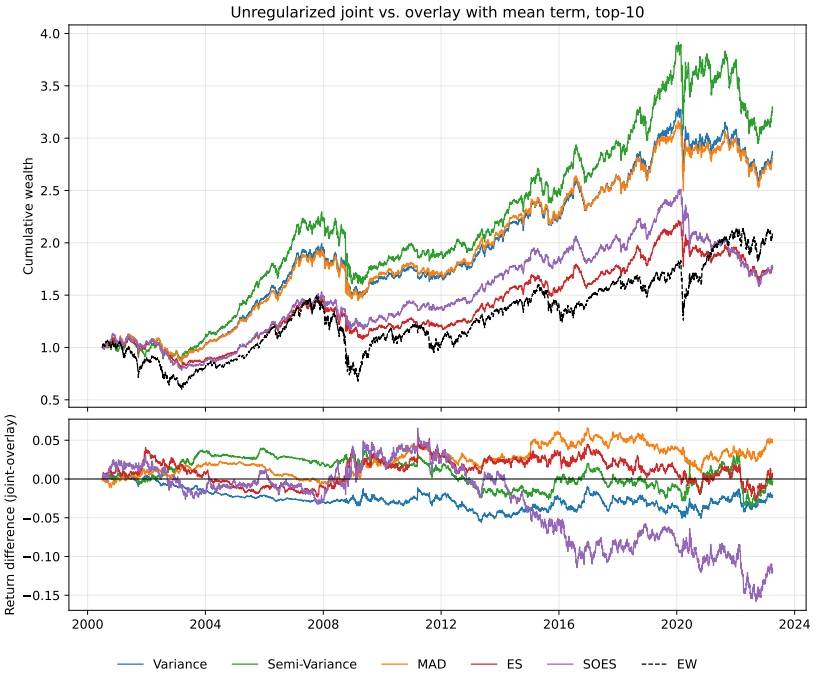


Figure 5.7: Cumulative performance of the mean-risk joint optimizers and their return difference to the overlay optimizers in the small investable universe (top 10 stocks per economy).

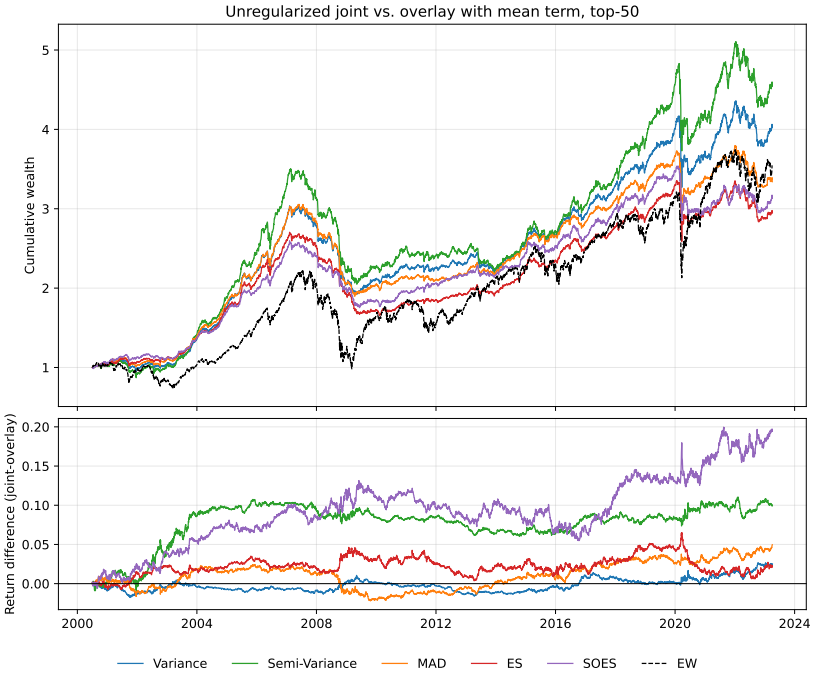


Figure 5.8: Cumulative performance of the mean-risk joint optimizers and their return difference to the overlay optimizers in the medium investable universe (top 50 stocks per economy).

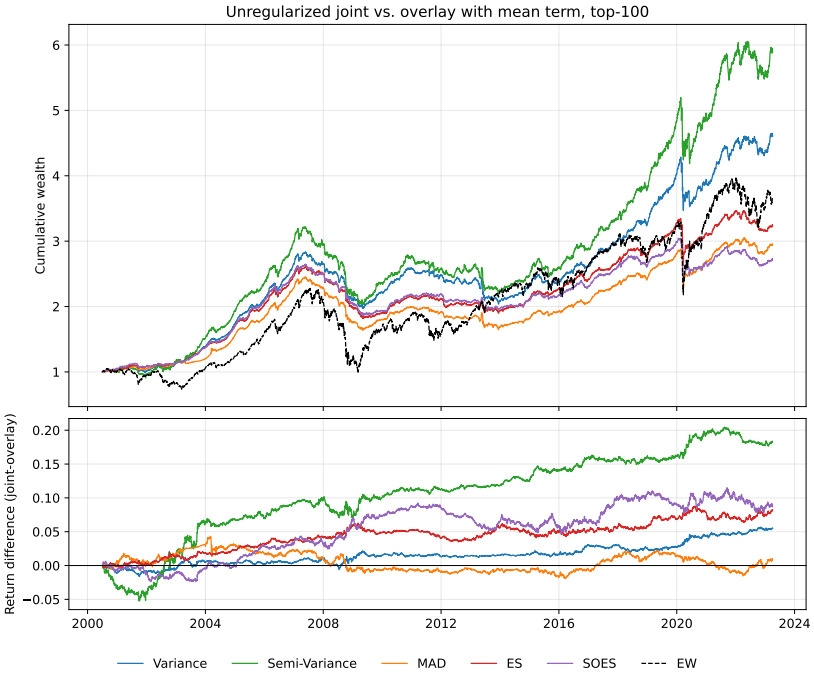


Figure 5.9: Cumulative performance of the mean-risk joint optimizers and their return difference to the overlay optimizers in the large investable universe (top 100 stocks per economy).

Figures 5.7 to 5.9 show the out-of-sample performance of the mean-risk optimizers in the small, medium and large investable universes, respectively. In each figure, the top panel shows the cumulative wealth of the joint optimizers and the equally weighted portfolio, and the bottom panel shows the cumulative sum of the daily return differences between the joint and the overlay optimizers. The portfolios are rebalanced quarterly and reported net of transaction costs. The corresponding performance metrics are reported in Panels A, B and C of Table 5.7, respectively.

The addition of a mean term also changes the comparison between the Joint and Overlay optimizers. In the previous sections, neither optimization approach consistently outperformed the other. Now Joint attains a higher Sharpe ratio than Overlay for every risk measure in the medium and large universes, by 0.03 to 0.17 and by about 0.07 on average. The Sortino ratios confirm this, with Joint performing better, by 0.02 to 0.25. The bottom panels of Figures 5.8 and 5.9 show that this advantage builds up gradually over the sample rather than in a single period, with the cumulative return difference reaching about 0.18 for semi-variance in the large universe and 0.19 for second-order Expected Shortfall in the medium universe by the end of the sample. At the same time, Joint attains a lower volatility than Overlay for all risk measures, by up to 0.44 percentage points. The lower volatility of the Joint optimizers is a persistent finding across all backtests in Sections 5.1 and 5.2. Intuitively, the joint program can choose assets and currencies that offset each other's risk, while the overlay program picks its assets before considering the currencies. Without a mean term, this flexibility is used only to reduce risk, which lowers volatility but leaves returns largely unchanged. With a mean term, the joint program can also weigh the expected returns of assets and currencies against each other, keeping currency exposure where it is expected to pay and hedging it where it is not.

5.4 Conclusion

All optimizers considered in this chapter achieve a much lower volatility and smaller drawdowns than the equally weighted portfolio. Without a mean term, the excess return of the Joint and Overlay optimizers is also lower. They minimize risk without any return term in the objective and have a lower exposure to the common movement of the equity markets, so they earn a smaller share of the equity risk premium. In the medium and large universes, the optimizers achieve a higher Sharpe ratio than the equally weighted portfolio, but in the small universe they do not outperform it consistently. With fewer stocks to choose from, the optimizers lose more excess return than they reduce volatility, while a larger universe offers more low-risk stocks, and the reduction in volatility outweighs the lost return.

Regularization increases the Sharpe ratios, especially in the large universe. The penalties limit the extreme positions that the unregularized programs fit to noise in the lookback window, which matters more when the number of assets is large relative to the length of the window. The regularized Joint and Overlay optimizers perform on par with the Full optimizer, which keeps the currency exposure fully hedged. Without a mean term in the objective, actively trading the currencies therefore adds little compared to full hedging in terms of returns, likely because the currency weights are then chosen only to reduce risk and not to earn the excess currency returns. Compared with the Zero optimizer, which keeps the full currency exposure, managing the currencies pays off in the medium and large universes, where equity risk is reduced further and currency risk makes up a larger share of the portfolio's total risk.

Adding a mean term to the objective raises the Sharpe ratios in almost every case and by a similar amount in all three universes. The gain comes from higher excess returns, while the volatility and drawdowns stay well below those of the equally weighted portfolio. The most significant improvement is for the deviation measures. Since they are invariant to a shift in the portfolio return, minimizing them ignores expected returns entirely, and the mean term adds information the optimizers did not use before. Expected Shortfall and second-order Expected Shortfall are translation equivariant and already reward a higher mean, so the explicit mean term changes them much less.

The Joint and Overlay optimizers perform very similarly when they only minimize risk. The Joint optimizer achieves a lower volatility and, without regularization, a higher currency

turnover. With regularization, the currency turnover decreases to about the size of those of the Overlay. In both cases, neither approach achieves consistently higher risk-adjusted returns. This changes when a mean term is added. The Joint optimizer then has higher Sharpe and Sortino ratios than the Overlay for every risk measure in the medium and large universes, and this advantage is consistent over the sample. Once expected returns enter the objective, the joint program can optimize the expected returns of assets and currencies against each other, which the overlay program cannot do since it fixes the asset weights first. This is the setting in which active currency management adds value, consistent with the result that it adds little when the currency weights are only optimized to reduce risk.

The lower volatility of the Joint optimizer is the most consistent difference between the two approaches. We would expect this, since both programs minimize risk and the joint program can always choose the same portfolio as the overlay program. Intuitively, the joint program can choose assets and currencies that offset each other's risk, while the overlay program picks its assets before considering the currencies, and this matters most in the small universe, where the number of assets is small relative to the number of currencies. In the next chapter, we therefore focus on comparing the optimizers in terms of reducing the risk measure they are optimizing for.

Chapter 6

Joint versus Sequential Optimization

In Chapter 5 we showed that the joint and overlay programs achieve similar returns, and that the most consistent difference between them is the lower realized volatility of the joint optimizers. In this chapter, we analyze the performance of the joint and overlay programs in terms of risk minimization. We first show in Section 6.1 that the joint program has an advantage over the overlay program by construction in terms of minimizing its target risk measure, since the overlay program's solution is feasible for the joint program. We define the in-sample and out-of-sample gaps between the two programs and investigate the realized advantage of joint optimization across the five risk measures for investable universes of different sizes. The results, as in Chapter 5, separate the risk measures into two groups, deviation and tail measures. In Section 6.2 we look at the tail measures more closely. We perform the same gap analysis across different confidence levels and show that the in-sample advantage of joint optimization grows with the tail depth while the out-of-sample advantage does not. This identifies a tail-fitting mechanism that does not generalize. We then transform the tail measures into deviation measures by de-meaning them and show that the in-sample advantage begins to fall with universe size, which identifies an implicit risk-return tradeoff as a second cause of the tail measures in-sample advantage. Finally, in Section 6.3 we focus on the deviation measures and argue, with the two mechanisms from Section 6.2, why a larger percentage of the in-sample advantage of joint optimization is realized out of sample.

6.1 Advantage of Joint Optimization for Risk Minimization

In sample, joint optimization provides a benefit over sequential optimization by construction when it comes to minimizing the target risk measure. We write the weight vector as $\boldsymbol{\theta} = (\boldsymbol{x}, \boldsymbol{\psi})$ and let $\Theta = \mathcal{X} \times \mathbb{R}^M$ denote the feasible set defined by the budget and leverage constraints as in Chapter 3. For a rebalancing period $t \in \{1, \ldots, T\}$, let ρ_t^{IS} be the in-sample risk over the lookback window and ρ_t^{OOS} the out-of-sample realized risk over the next holding period. Let $\hat{\boldsymbol{\theta}}_t$ denote the weight vector an optimizer returns on the lookback window, and $\rho_t^{\text{OOS}}(\hat{\boldsymbol{\theta}}_t)$ and $\rho_t^{\text{IS}}(\hat{\boldsymbol{\theta}}_t)$ are the risk these weight vectors realize. The two optimization programs are written as

$$\hat{\boldsymbol{\theta}}_{J,t} \in \operatorname*{arg\,min}_{\boldsymbol{\theta} \in \Theta} \rho_t^{\text{IS}}(\boldsymbol{\theta}),$$

and

$$\hat{\boldsymbol{x}}_{O,t} \in \operatorname*{arg\,min}_{\boldsymbol{x} \in \mathcal{X}} \rho_{A,t}^{\text{IS}}(\boldsymbol{x}), \quad \hat{\boldsymbol{\psi}}_{O,t} \in \operatorname*{arg\,min}_{\boldsymbol{\psi} \in \mathbb{R}^M} \rho_t^{\text{IS}}(\hat{\boldsymbol{x}}_{O,t}, \boldsymbol{\psi}),$$

where $\rho_{A,t}^{\text{IS}}(\boldsymbol{x}) \coloneqq \rho_t^{\text{IS}}(\boldsymbol{x}, \mathbf{0})$ is the in-sample risk of the fully hedged asset portfolio, and $\hat{\boldsymbol{\theta}}_{i,t} = (\hat{\boldsymbol{x}}_{i,t}, \hat{\boldsymbol{\psi}}_{i,t})$, $i \in \{J, O\}$. The subscript J denotes the joint optimizer and O the sequential

overlay optimizer. The overlay optimizer is a restricted minimizer. It first minimizes over $\boldsymbol{x}$ against a different objective, then over $\boldsymbol{\psi}$ with $\hat{\boldsymbol{x}}_{O,t}$ held fixed.

The overlay optimizer's output satisfies the joint program's constraints, so $\hat{\boldsymbol{\theta}}_{O,t} \in \Theta$, and its second step is a partial minimization of the same ρ_t^{IS}. Hence

$$\rho_t^{\mathrm{IS}}(\hat{\boldsymbol{\theta}}_{J,t}) = \min_{\boldsymbol{\theta}\in\Theta} \rho_t^{\mathrm{IS}}(\boldsymbol{\theta}) \;\leq\; \min_{\boldsymbol{\psi}\in\mathbb{R}^M} \rho_t^{\mathrm{IS}}(\hat{\boldsymbol{x}}_{O,t}, \boldsymbol{\psi}) = \rho_t^{\mathrm{IS}}(\hat{\boldsymbol{\theta}}_{O,t}).$$

We define the in-sample gap between the two optimizers as

$$g_t^{\mathrm{IS}} \coloneqq \rho_t^{\mathrm{IS}}(\hat{\boldsymbol{\theta}}_{O,t}) - \rho_t^{\mathrm{IS}}(\hat{\boldsymbol{\theta}}_{J,t}) \;\geq\; 0.$$

The in-sample gap is non-negative and measures the advantage of joint optimization over sequential optimization in the lookback window. Naturally, the question we are interested in is whether this advantage is realized out of sample.

We define the out-of-sample error of a fitted $\hat{\boldsymbol{\theta}}_t$ as

$$\epsilon_t(\hat{\boldsymbol{\theta}}_t) \coloneqq \rho_t^{\mathrm{OOS}}(\hat{\boldsymbol{\theta}}_t) - \rho_t^{\mathrm{IS}}(\hat{\boldsymbol{\theta}}_t),$$

the amount by which realized risk in the holding quarter exceeds the risk the optimizer achieved in sample.[1] Then the out-of-sample gap is defined as

$$\begin{aligned} g_t^{\mathrm{OOS}} &\coloneqq \rho_t^{\mathrm{OOS}}(\hat{\boldsymbol{\theta}}_{O,t}) - \rho_t^{\mathrm{OOS}}(\hat{\boldsymbol{\theta}}_{J,t}) \\ &= \underbrace{\left[\rho_t^{\mathrm{IS}}(\hat{\boldsymbol{\theta}}_{O,t}) - \rho_t^{\mathrm{IS}}(\hat{\boldsymbol{\theta}}_{J,t})\right]}_{g_t^{\mathrm{IS}} \geq 0} - \underbrace{\left[\epsilon_t(\hat{\boldsymbol{\theta}}_{J,t}) - \epsilon_t(\hat{\boldsymbol{\theta}}_{O,t})\right]}_{\Delta\epsilon_t}. \end{aligned}$$

The difference between the out-of-sample errors $\Delta\epsilon_t$ can be either positive or negative. For the variance, the in-sample risk of an optimized portfolio is known to understate its true risk, and the more weights are fitted relative to the sample size, the larger this understatement becomes [6]. Since the joint program fits all of its weights to the same in-sample objective, we expect it to adapt more closely to the noise in the lookback window than the overlay program, and hence $\Delta\epsilon_t$ to be typically positive. If $\Delta\epsilon_t$ exceeds g_t^{IS}, the overlay program could in some cases realize lower risk than the joint program out of sample. We also expect $\Delta\epsilon_t$ to be measure-dependent. Tail measures, in our case Expected Shortfall and second-order Expected Shortfall, are estimated from a smaller fraction of the sample, whereas the deviation measures, variance and mean absolute deviation use every observation, and semi-variance those below the mean. For a given sample size, we therefore expect $\Delta\epsilon_t$ to be larger and more volatile for the tail measures.

We perform the following gap analysis to investigate the in-sample and out-of-sample performance of the joint and overlay programs with respect to the risk measure they are minimizing. The in-sample gap is guaranteed to be non-negative only when each program minimizes the risk measure it is scored on, so we consider the unregularized formulation from Chapter 3. At each of the $T = 91$ quarterly rebalancing dates the joint and the overlay optimizers, for a given risk measure, are solved on the same ten-year lookback window, and under the same budget and leverage constraints. The resulting weights are then held fixed and scored twice, once on the lookback window itself and once on the following holding quarter, which gives ρ_t^{IS} and ρ_t^{OOS} for each of the two programs. Both gaps are reported relative to the risk of the overlay optimizer

[1] This quantity is the portfolio analogue of the optimism of the training error, see for example Section 7.4 in [9].

$$\tilde{g}_t^{\mathrm{IS}} := \frac{g_t^{\mathrm{IS}}}{\rho_t^{\mathrm{IS}}(\hat{\boldsymbol{\theta}}_{O,t})} \qquad \text{and} \qquad \tilde{g}_t^{\mathrm{OOS}} := \frac{g_t^{\mathrm{OOS}}}{\rho_t^{\mathrm{OOS}}(\hat{\boldsymbol{\theta}}_{O,t})},$$

so they are to be interpreted as the risk reduction achieved by the joint optimizer over the overlay optimizer, expressed as a percentage of the latter's risk. For each risk measure and universe size we report the mean of the relative gaps across the rebalancing dates together with a 95% bootstrap confidence interval based on 10,000 resamples.

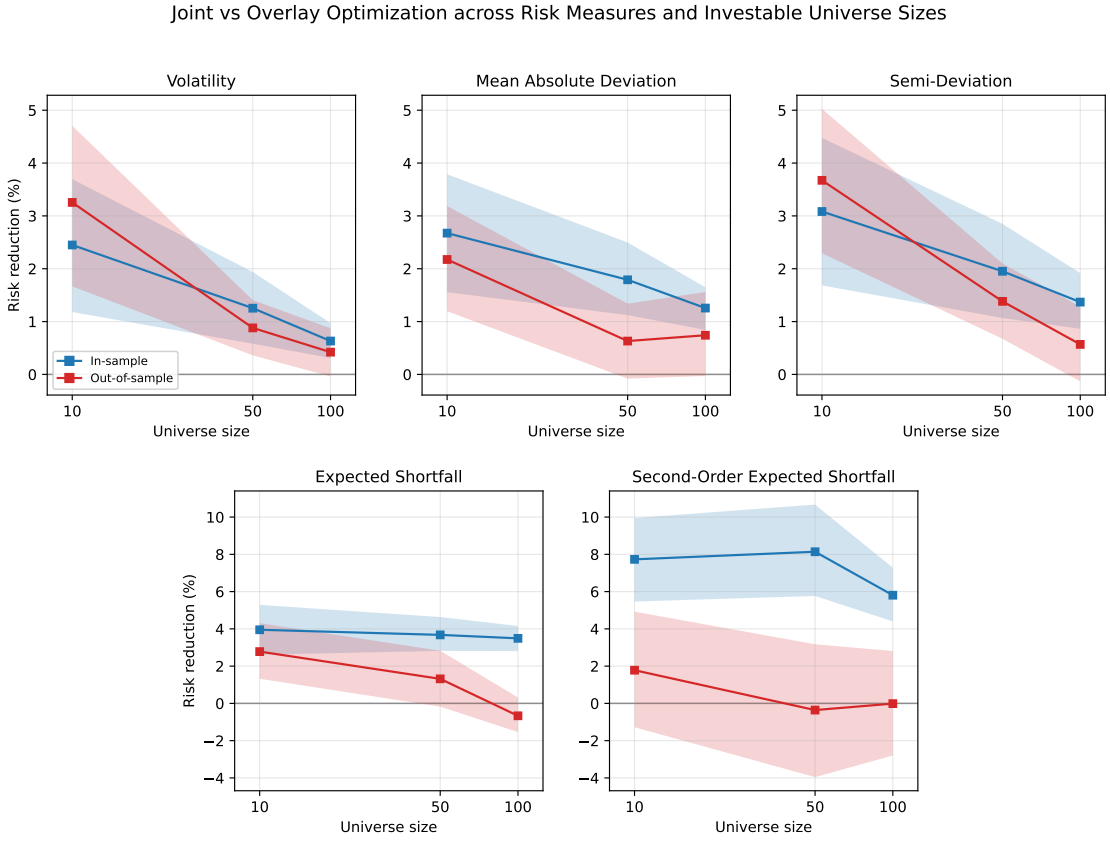


Figure 6.1: In-sample and out-of-sample relative gaps between the unregularized joint and overlay optimizers, by risk measure and universe size.

Figure 6.1 reports the mean in-sample and out-of-sample relative gaps between the unregularized joint and overlay optimizers, expressed as a percentage of the overlay optimizer's risk. Each panel reports one risk measure. The horizontal axis is the size of the investable universe and the vertical axis the risk reduction the joint program achieves over the overlay program. Squares mark the mean across the $T = 91$ quarterly rebalancing dates, in blue for the in-sample gap $\tilde{g}_t^{\mathrm{IS}}$ and in red for the realized gap $\tilde{g}_t^{\mathrm{OOS}}$. Shaded areas mark the 95% bootstrapped confidence intervals.

We see in Figure 6.1 that the in-sample gap of joint optimization is positive for all risk measures across all universe sizes, and that it is largest for the two considered tail measures. The out-of-sample gap is smaller than the in-sample gap for almost all risk measures and universe sizes, the only exception being the variance and semi-variance measures in the small universe. Out of sample, we see a negative dependence on the size of the investable universe. In the small universe the realized out-of-sample advantage of joint optimization is in the range between 1.78% and 3.67%, with second-order Expected Shortfall having the smallest benefit and semi-variance having the largest. In the medium universe, the advantage drops to zero for second-order Expected Shortfall, and remains around 1% for the other measures, and in the large universe the advantage for Expected Shortfall drops slightly below zero while the remaining measures stay below 1%.

The in-sample series separates the two groups of risk measures the same way as in Chapter 5, that is, whether they are deviation or tail measures. The in-sample advantage falls as the universe grows, from 2.45% to 0.63% for variance, from 2.67% to 1.26% for mean absolute deviation and from 3.08% to 1.37% for semi-variance, whereas for the two tail measures it remains larger and more stable, from 3.95% to 3.49% for Expected Shortfall and from 7.73% to 5.81% for second-order Expected Shortfall. For the deviation measures the in-sample and out-of-sample series fall together, whereas for the tail measures, the out-of-sample series falls to zero, and the in-sample series remains large and positive. We visualize the difference between the two series in Figure 6.2.

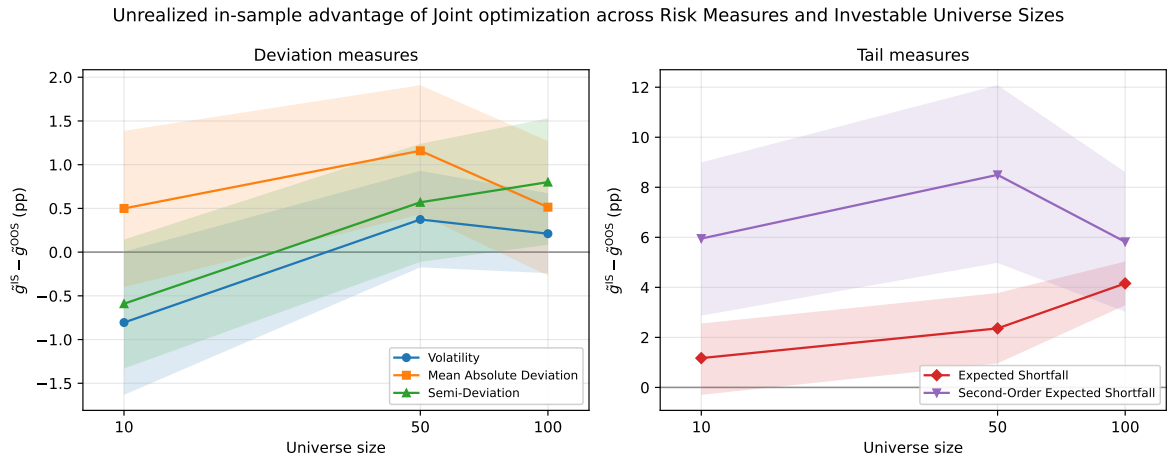


Figure 6.2: Unrealized in-sample advantage of unregularized joint optimization, $\tilde{g}_t^{\text{IS}} - \tilde{g}_t^{\text{OOS}}$, by risk measure and universe size.

Figure 6.2 reports $\tilde{g}_t^{\text{IS}} - \tilde{g}_t^{\text{OOS}}$ for the unregularized optimizers, the percentage advantage the joint optimizer holds in sample minus the percentage advantage it realizes out of sample. For every risk measure, we mark the mean difference across the 91 quarterly rebalancing dates, with shaded areas marking the 95% bootstrap intervals obtained the same way as in Figure 6.1. Being a difference of two ratios, the quantity reported is measured in percentage points. The deviation and the tail measures are shown on separate panels, with independent vertical scales, because the latter are an order of magnitude larger.

The first thing we notice in Figure 6.2 is that the bootstrapped confidence intervals for the two tail measures are wider than those for the three deviation measures, averaging 4.37 against 1.41 percentage points. This is expected, since the tail measures are estimated on a smaller fraction of the sample. The second thing is the size of the unrealized advantage itself, which is between 1.17 and 8.50 percentage points for the two tail measures while the three deviation measures are between -0.81 and 1.16 percentage points. This, together with the in-sample series of Figure 6.1 (which for the tail measures stays between 3.49% and 8.14% at every universe size), shows that the tail measures have a larger in-sample advantage coming from joint optimization, that is not realized out of sample, whereas for the deviation measures the advantage that is achieved in sample is largely realized out of sample but drops with universe size. The two groups produce different results, and we analyze the mechanism behind them separately in Sections 6.2 and 6.3.

6.2 Tail Measures

6.2.1 Level Analysis

To investigate the mechanism behind the in-sample advantage of joint optimization for the two tail measures, Expected Shortfall and second-order Expected Shortfall, we perform the same gap analysis as in Section 6.1 across a range of confidence levels. We report the results in Figures 6.3 and 6.4.

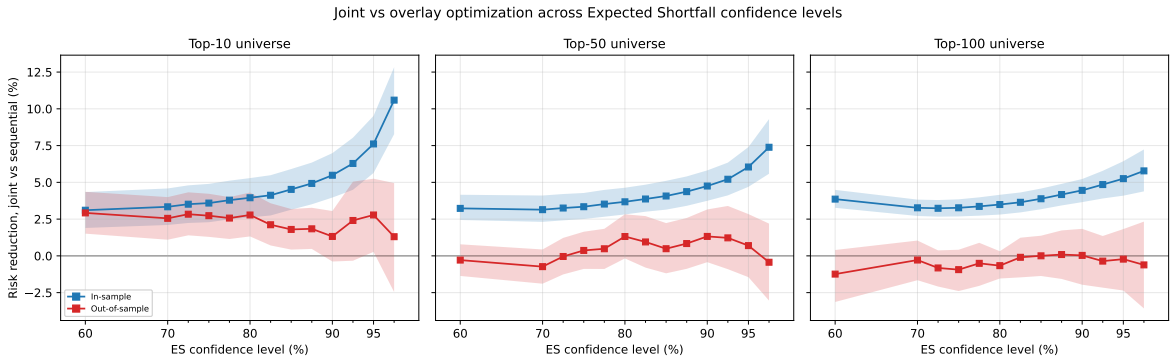


Figure 6.3: In-sample and out-of-sample relative gaps between the unregularized joint and overlay Expected Shortfall optimizers, by confidence level and universe size.

Figure 6.3 reports the mean in-sample and out-of-sample relative gaps between the unregularized joint and overlay Expected Shortfall optimizers, expressed as a percentage of the overlay optimizer's risk, with the confidence level on the horizontal axis and one panel per universe size. Squares mark the mean across the $T = 91$ quarterly rebalancing dates, in blue for the in-sample gap $\tilde{g}_t^{\mathrm{IS}}$ and in red for the realized gap $\tilde{g}_t^{\mathrm{OOS}}$. Shaded areas mark the 95% bootstrapped confidence intervals. The figures cover thirteen confidence levels from 60% to 97.5%.

We see in Figure 6.3 that the in-sample advantage of joint optimization for minimizing Expected Shortfall grows with the confidence level and the growth is more pronounced in the smaller universe. The out-of-sample advantage is much smaller and does not rise with the confidence level at any universe size. In the small universe, the in-sample advantage rises from 3.11% at confidence level 60% to 10.59% at level 97.5%, while the out-of-sample advantage stays in the 1.30 to 2.92% range. In the larger universes the in-sample advantage grows at a lower rate, going from 3.23% to 7.39% in the medium and from 3.86% to 5.78% in the large universe, while the realized advantage never exceeds 1.32% and 0.09% respectively.

Expected Shortfall at confidence level α is the average loss over the worst $1 - \alpha$ fraction of the lookback window, which is roughly 1,000 observations at 60% and 63 at 97.5%. Which observations fall in the tail depends largely on the portfolio weights. The joint program chooses the asset and currency weights together, so it can select assets whose worst days the currency positions can offset. The overlay program fixes the asset weights of the fully hedged portfolio, and its currency step can adjust M weights against a largely fixed set of the worst performing days. As we increase the confidence level, the number of observations in the tail shrinks, so the joint program's currency weights can cover a larger share of the worst performing days, and the in-sample advantage grows. The growth is largest in the small universe, where the currency weights are a larger share of all the decision variables. In larger universes, the overlay program's asset step has more weights to diversify, so the joint program's advantage decreases. Since the weights are fitted to specific worst performing days of the lookback window that do not necessarily occur in the next period, the mechanism generalizes poorly out of sample.



Figure 6.4: In-sample and out-of-sample relative gaps between the unregularized joint and overlay second-order Expected Shortfall optimizers, by confidence level and universe size.

Figure 6.4 reports the mean in-sample and out-of-sample relative gaps between the unregularized joint and overlay second-order Expected Shortfall optimizers, expressed as a percentage of the overlay optimizer's risk, with the confidence level on the horizontal axis and one panel per universe size. Squares mark the mean across the $T = 91$ quarterly rebalancing dates, in blue for the in-sample gap $\tilde{g}_t^{\mathrm{IS}}$ and in red for the realized gap $\tilde{g}_t^{\mathrm{OOS}}$. Shaded areas mark the 95% bootstrapped confidence intervals. The figures cover thirteen confidence levels from 60% to 97.5%.

In Figure 6.4 we see the same pattern for second-order Expected Shortfall. The in-sample advantage of joint optimization grows with the confidence level and the growth is faster in the smaller universe. We also observe that above certain confidence levels, the in-sample advantage becomes constant. The out-of-sample advantage is much smaller and does not rise with the confidence level at any universe size. In the small universe, the in-sample advantage rises from 4.00% at level 60% to 16.28% at level 92.5% and from there remains constant. The out-of-sample advantage stays in the -0.45 to 3.88% range. In the larger universes the in-sample advantage grows at a lower rate and becomes constant at lower confidence levels, going from 3.44% to 8.52% in the medium and from 3.83% to 5.81% in the large universe, while the out-of-sample advantage fluctuates around zero.

Second-order Expected Shortfall squares the tail observations, so more weight is placed on the most extreme observations. At level α it therefore behaves like Expected Shortfall at a higher confidence level, which is why the in-sample advantage is larger at every confidence level and every universe size. As we increase the confidence level, the squared weighting concentrates on a smaller and more stable set of the worst performing days, and above a certain level, raising the confidence level further no longer changes which observations dominate the objective, and the in-sample series becomes constant. In larger universes this happens at lower confidence levels, since the overlay program's asset step has more weights to fit those days. Out of sample, the same reasoning as for Expected Shortfall applies, and the realized advantage does not rise with the confidence level at any universe size and is only present in the small universe.

6.2.2 De-meaning the Tail Measures

As discussed in Chapter 3, Expected Shortfall and second-order Expected Shortfall are coherent (strictly expectation bounded) risk measures rather than deviation measures. Unlike variance, mean absolute deviation and semi-variance, they are not invariant to translation, so a portfolio with a higher mean return is scored favorably to one with a lower mean return, everything else being equal. The objectives therefore carry an implicit risk-return tradeoff that the three deviation measures do not. We can transform a coherent strictly expectation bounded risk measure $\mathcal{R}$ into a deviation measure $\mathcal{D}$ via a simple de-meaning operation. More precisely

Theorem D.3 *Let $\mathcal{R}$ be a coherent strictly expectation bounded risk measure, and let $\mathcal{D}$ be the functional defined by*

$$\mathcal{D}(X) \coloneqq \mathcal{R}(X - \mathbb{E}[X]).$$

Then $\mathcal{D}$ is a deviation measure.

The proof of Theorem D.3 and further details can be found in Appendix D. We perform the same gap analysis as in Section 6.1 on the de-meaned Expected Shortfall and second-order Expected Shortfall optimization programs, and report the results in Figure 6.5.

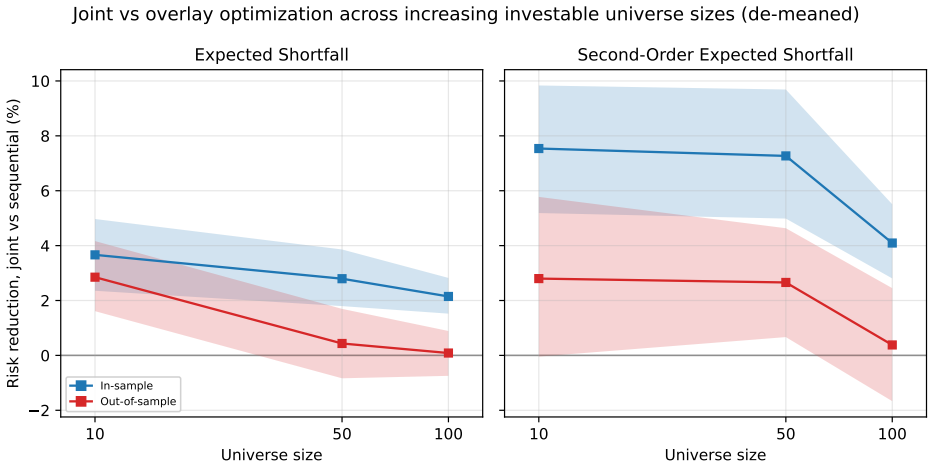


Figure 6.5: In-sample and out-of-sample relative gaps between the unregularized joint and overlay de-meaned Expected Shortfall and second-order Expected Shortfall optimizers, by universe size.

Figure 6.5 reports the mean in-sample and out-of-sample relative gaps between the unregularized joint and overlay de-meaned Expected Shortfall and second-order Expected Shortfall optimizers, expressed as a percentage of the overlay optimizer's risk, with the universe size on the horizontal axis. Squares mark the mean across the $T = 91$ quarterly rebalancing dates, in blue for the in-sample gap $\tilde{g}_t^{\mathrm{IS}}$ and in red for the realized gap $\tilde{g}_t^{\mathrm{OOS}}$. Shaded areas mark the 95% bootstrapped confidence intervals.

We see in Figure 6.5 that de-meaning the optimizers steepens the decline of the in-sample relative gaps. For Expected Shortfall the in-sample gap falls from 3.66% to 2.15% across the three universes, against 3.95% to 3.49% for the original tail measure, and for second-order Expected Shortfall from 7.54% to 4.09% against 7.73% to 5.81%. The drop in the gap across universes is 41.3% as opposed to 11.6% for Expected Shortfall and 45.7% as opposed to 24.8% for second-order Expected Shortfall. This indicates that the implicit risk-return tradeoff is one of the driving factors for the tail measures in-sample advantage as the universe grows. With it removed, they begin to decay in the way the deviation measures do. In the large universe the de-meaned Expected Shortfall gap is still 2.15% against, for example, 0.63% for variance, so the mean term accounts for only part of the in-sample advantage.

6.3 Deviation Measures

For the three considered deviation measures the advantage of joint optimization is smaller in sample than for the tail measures, and most of it survives out of sample. If we compare the two series from Figure 6.1 measure by measure, the realized gap $\tilde{g}^{\mathrm{OOS}}$ as a fraction of the in-sample gap $\tilde{g}^{\mathrm{IS}}$ is 133%, 70% and 67% for variance across the small, medium and large universes, 81%, 35% and 59% for mean absolute deviation, and 119%, 71% and 42% for semi-variance. For the tail measures the same ratio is 70%, 36% and −19% for Expected Shortfall and 23%, −4% and 0% for second-order Expected Shortfall.

Variance, semi-variance and mean absolute deviation are estimated from all observations in the window (half for semi-variance), so the objective is estimated more accurately and the out-of-sample gap follows the in-sample gap more closely. Expected Shortfall and its second-order variant are determined by a small fraction of the observations at the confidence levels considered, so the same gap is a noisier quantity. Figures 6.3 and 6.4 make this explicit, as the tail is made thinner the in-sample gap grows while the realized gap stays flat or decays. The deviation measures are also translation invariant, so they carry no implicit risk-return tradeoff, and the joint program cannot produce an in-sample advantage by tilting the weights towards a high realized mean. The tail measures can, and Figure 6.5 shows that removing the mean term accounts for part of the in-sample advantage across universes. So deviation measures are left with a small advantage that shrinks both in sample and out of sample as the investable universe grows.

6.4 Conclusion

Even though joint optimization does not consistently outperform in terms of returns, it has an advantage over sequential optimization by construction when it comes to minimizing the targeted risk measure. In sample, the advantage for the deviation measures falls with universe size, while for the two considered tail measures it stays stable. Out of sample, the advantage falls with universe size for all five considered risk measures.

The tail measures in-sample advantage is driven by two mechanisms. The first is a tail-fitting mechanism, which lets the joint program align the currency weights with the specific worst performing days of the lookback window, and which does not generalize out of sample. The second is an implicit risk-return tradeoff, which lets the joint program produce an in-sample advantage by tilting the weights towards a high realized mean. Once de-meaned, the tail measures in-sample advantage begins to fall with universe size, in the same way as the deviation measures. Neither of the two mechanisms is available to the deviation measures, which is why their in-sample advantage was smaller to begin with and a larger fraction of it survives out of sample.

The out-of-sample advantage of joint optimization is present only in the small universe, where the currency weights are a larger share of the decision variables. In larger universes the overlay program's asset step has enough weights to fit the objective on its own, and the joint program's advantage drops to zero. Joint optimization is therefore beneficial for risk reduction when the number of assets is small relative to the number of currencies. These results agree with the findings in Chapter 5, where the joint program's lower realized volatility was largest in the small universe.

Chapter 7

Conclusion

In this thesis, we studied how risk-based portfolio optimization performs for an investor who holds assets denominated in several currencies, and whether optimizing the asset and currency weights jointly improves on the industry-standard sequential currency overlay. We considered five risk measures, variance, semi-variance, mean absolute deviation, Expected Shortfall and second-order Expected Shortfall, and investable universes of three sizes, built from the largest stocks of 14 economies by market capitalization. In Chapter 5 we backtested the joint and overlay programs in an unregularized, a regularized and a mean-risk formulation against the equally weighted portfolio and two asset-only benchmarks, and in Chapter 6 we compared how well the two programs minimize their target risk measure in and out of sample. We find that the optimizers achieve a much lower risk than the equally weighted portfolio in every universe and consistently achieve higher risk-adjusted returns in the medium and large universes. The benefit of joint over sequential optimization is limited to two settings, objectives with a mean term in the medium and large universes, and risk minimization in the small universe.

7.1 Findings

Out of sample, the unregularized risk-minimizing optimizers reduce volatility by 34% to 64% relative to the equally weighted portfolio, and its maximum drawdown by up to 30 percentage points. Without a return term in the objective and with a lower exposure to the common movement of the equity markets, they earn a smaller share of the equity risk premium and lower excess returns. On a risk-adjusted basis, they outperform the equally weighted portfolio, with Sharpe ratios of up to 70% higher in the medium universe and up to 105% higher in the large universe, but do not consistently outperform in the small universe. The results are consistent with the low-volatility anomaly empirical finding, where low-risk stocks tend to earn higher risk-adjusted returns than high-risk stocks. The optimizers tilt the portfolio toward low-risk stocks, and a larger universe offers more of them, while in the small universe they give up more return than they remove risk.

Regularizing the optimizers further improves their risk-adjusted performance, most in the large universe, where the Sharpe ratios increase by 0.17 on average. The penalties limit the extreme positions that the unregularized programs fit to noise in the lookback window. Both the regularized Joint and Overlay optimizers perform on par with the Full optimizer, which keeps the currency exposure fully hedged, so actively trading the currencies adds little over full hedging. The Zero optimizer, which leaves the currency exposure unhedged, outperforms them only in the small universe. In the medium and large universes, where currency risk makes up a larger share of the portfolio's total risk, managing it actively is more effective

than leaving it unhedged, and the regularized Joint and Overlay optimizers achieve better risk-adjusted performance than the Zero optimizer.

Adding a mean term to the optimizers transforms the problem into a risk-return optimization problem. It has a large effect on risk-adjusted performance, increasing the Sharpe ratios by 0.16 to 0.24 on average for the deviation measures and by 0.01 to 0.10 for the tail measures. This is not surprising, as deviation measures are invariant to a shift in the portfolio return, so minimizing them ignores expected returns entirely, whereas Expected Shortfall and second-order Expected Shortfall already reward a higher mean. The mean term also increases exposure to higher-yielding currencies, so the currency positions likely earn part of the carry premium that pure risk minimization gives up. The optimizers earn higher risk-adjusted returns than the equally weighted portfolio in the medium and large universes, and the deviation measures also in the small universe, at about 40% to 75% of its volatility and maximum drawdown.

Comparing all backtests in Tables 5.1, 5.3 to 5.5 and 5.7, the mean-variance Joint optimizer performs best out of sample. It more than doubles the Sharpe ratio of the equally weighted portfolio across all universe sizes and outperforms the Full and Zero asset-only benchmarks by a significant margin. Although several of the risk-minimizing optimizers have smaller drawdowns, the mean-variance Joint optimizer still outperforms them on a risk-adjusted basis. The optimizer benefits from two sources of improvement from its other formulations. Variance is a deviation measure, which gains the most from the mean term, and the joint program can weigh the expected returns of assets and currencies against each other as compared to overlay.

The benefit of optimizing assets and currencies jointly rather than in two steps depends on whether expected returns are part of the objective. When the optimizers only minimize risk, the Joint and Overlay optimizers perform very similarly, with or without regularization, with joint optimization consistently achieving lower volatility, although predominantly in the small universe. With a mean term in the objective, the Joint optimizers achieve higher Sharpe and Sortino ratios than the Overlay for every risk measure in the medium and large universes.

Since joint optimization consistently achieves lower risk than sequential optimization, we compare the two optimization approaches on the risk measure they minimize in Chapter 6. In sample, the joint program achieves lower risk by construction. Out of sample, this advantage is realized only in the small universe, where the currency weights are a larger share of the decision variables. In the medium and large universes it drops to zero, as the asset step of the overlay program has enough weights to fit the objective on its own. Since the dependence between asset and currency returns barely changes across universe sizes, this decline is driven by the number of assets relative to the number of currencies and not by a change in the dependence structure.

For the deviation measures, most of the small in-sample advantage in risk reduction is realized out of sample, while the larger in-sample advantage of the tail measures is not. We identify two mechanisms that drive this discrepancy, a tail-fitting mechanism, in which the joint program aligns the currency weights with the specific worst days of the lookback window, and an implicit risk-return tradeoff, through which the tail measures reward a high realized mean. Both mechanisms fail to generalize out of sample.

Our results agree with the finding of [30] that the benefit of joint optimization is most pronounced in small universes, which we find for risk minimization, and with the finding of [31] that joint optimization outperforms the overlay in a mean-variance setting, which we find in the medium and large universes once a mean term is added.

7.2 Limitations and Future Research

Our analysis covers one base currency, the US dollar, 14 economies, and one out-of-sample period from 2000 to 2023, so the results reflect the dependence between asset and currency returns for a US investor over this period. For an investor based in another currency, such as the euro, the Swiss franc or the Japanese yen, the excess currency returns and their dependence with the asset returns differ, and repeating the analysis for other base currencies would show how much the results depend on this structure.

The backtests are also based on simplified trading assumptions, namely constant transaction costs, no market impact, and are conducted with quarterly rebalancing, among others. Since the unregularized and mean-risk Joint optimizers have a much higher currency turnover than the Overlay in most cases, higher trading costs, more frequent rebalancing and market impact would likely reduce the advantage of the Joint optimizers over the Overlay. A more realistic analysis could include a dynamic transaction cost model and multiple rebalancing frequencies, and test the sensitivity of the results to these assumptions.

Two further extensions follow from the results. First, the tail measures are estimated from the empirical distribution of the lookback window, and are sensitive to estimation error. Estimators or formulations that stabilize the tail, such as the extreme value approach proposed in [19], could improve the out-of-sample performance. Second, the explanations given for the optimizers' performance, a lower market exposure, a tilt towards low-risk stocks and currency carry, could be further investigated by regressing their returns on market, low-volatility, carry and other factors.

Appendices

Appendix A

Optimization Derivation: Expected Shortfall and Second-Order Expected Shortfall

In this appendix, we derive the optimization problems used in the main text, building up from the Expected Shortfall objective to the implementable quadratic programs solved in the empirical analysis. In Section A.1 we follow [28] to derive the equivalence between direct minimization of Expected Shortfall and an auxiliary formulation suitable for linear programming. Following [30], in Section A.2 we embed this equivalence in a mean-ES objective, by replacing the conditional expectation with its empirical counterpart, and linearizing the terms to assemble the base unregularized portfolio optimization problem as a linear program. In Section A.3 we introduce three regularization techniques, an L_1 penalty, an L_2 penalty, and a leverage constraint, giving us the base regularized quadratic program for portfolio optimization. In Section A.4 we extend the framework to the multi-currency setting under the separate (currency overlay) approach, deriving the asset and the currency quadratic programs. Finally, in Section A.5 we derive the joint quadratic program formulation, in which asset and currency weights are optimized simultaneously through a single Expected Shortfall term that captures the combined return distribution.

A.1 Auxiliary Formulations of Expected Shortfall and Second-Order Expected Shortfall

We derive the equivalence between the direct minimization of Expected Shortfall and the auxiliary formulation proposed by Rockafellar and Uryasev [28], and introduce second-order Expected Shortfall, which Krokhmal [13] defines directly through an auxiliary formulation of the same type. The two formulations allow for writing the portfolio optimization problem as a linear program for Expected Shortfall and as a second-order cone program for second-order Expected Shortfall.

A.1.1 Expected Shortfall

We assume that the loss distribution is continuous and has a strictly increasing cumulative distribution function, meaning the cumulative distribution function has neither jumps nor flat sections, and additionally that its probability density function exists. Under these assumptions, the Value-at-Risk (VaR) defined below is the unique point at which the cumulative distribution function equals the confidence level [24]. The assumptions simplify the derivation but are not needed for the result. Rockafellar and Uryasev [27] show that the equivalence derived below

holds for general loss distributions, including the empirical distribution used in Section A.2, when Expected Shortfall is defined as the average of the quantiles of the loss distribution over all levels above the confidence level [1].

Define the loss random variable at time $t+1$, given portfolio weights $\boldsymbol{x}_t \in \mathbb{R}^N$, as $L(\boldsymbol{x}_t) = -\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1}$, where $\boldsymbol{R}_{t+1} \in \mathbb{R}^N$ is the asset return vector. Let $f(l)$ denote the probability density function of the loss $L(\boldsymbol{x}_t)$. The VaR at confidence level β is defined as the β-quantile of the loss distribution

$$\mathrm{VaR}_\beta(\boldsymbol{x}_t) = \min\{\alpha : \mathbb{P}(L(\boldsymbol{x}_t) \leq \alpha) \geq \beta\}.$$

The Expected Shortfall at confidence level β is defined as the conditional expectation of losses exceeding VaR_β

$$\mathrm{ES}_\beta(\boldsymbol{x}_t) = \mathbb{E}_t[L(\boldsymbol{x}_t) \mid L(\boldsymbol{x}_t) \geq \mathrm{VaR}_\beta(\boldsymbol{x}_t)] = \frac{1}{1-\beta}\int_{\mathrm{VaR}_\beta(\boldsymbol{x}_t)}^{\infty} l f(l)\, dl.$$

We now consider the auxiliary function F_β defined as

$$\begin{aligned} F_\beta(\boldsymbol{x}_t, \alpha) &= \alpha + \frac{1}{1-\beta}\mathbb{E}_t[(L(\boldsymbol{x}_t) - \alpha)^+] \\ &= \alpha + \frac{1}{1-\beta}\int_\alpha^\infty (l-\alpha) f(l)\, dl. \end{aligned} \tag{A.1}$$

Before minimizing (A.1) with respect to α, we verify convexity. Using Leibniz's rule, the first and second derivatives of F_β with respect to α are

$$\begin{aligned} \frac{\partial F_\beta}{\partial \alpha} &= 1 + \frac{1}{1-\beta}\frac{\partial}{\partial \alpha}\int_\alpha^\infty (l-\alpha) f(l)\, dl \\ &= 1 + \frac{1}{1-\beta}\left[-f(\alpha)\cdot 0 + \int_\alpha^\infty (-1) f(l)\, dl\right] \\ &= 1 - \frac{1}{1-\beta}\int_\alpha^\infty f(l)\, dl, \\ \frac{\partial^2 F_\beta}{\partial \alpha^2} &= \frac{1}{1-\beta} f(\alpha). \end{aligned}$$

Since the probability density function $f(\alpha) \geq 0$ and $\beta < 1$, the second derivative is non-negative. Thus, (A.1) is convex in α, ensuring that the first-order condition is a necessary and sufficient condition for a global minimum.

Setting the derivative to zero gives the optimal α^*

$$\begin{aligned} \frac{\partial F_\beta}{\partial \alpha} &= 0 \\ 1 - \frac{1}{1-\beta}\int_{\alpha^*}^\infty f(l)\, dl &= 0 \\ 1 - \frac{1}{1-\beta}\mathbb{P}(L(\boldsymbol{x}_t) \geq \alpha^*) &= 0 \\ \mathbb{P}(L(\boldsymbol{x}_t) \geq \alpha^*) &= 1-\beta \\ \mathbb{P}(L(\boldsymbol{x}_t) \leq \alpha^*) &= \beta. \end{aligned}$$

We see that the optimal α^* that minimizes the auxiliary function for a fixed $\boldsymbol{x}_t$ is exactly the Value-at-Risk at level β.

Substituting $\alpha^* = \mathrm{VaR}_\beta(\boldsymbol{x}_t)$ back into (A.1) we obtain

$$\begin{aligned} F_\beta(\boldsymbol{x}_t, \mathrm{VaR}_\beta(\boldsymbol{x}_t)) &= \mathrm{VaR}_\beta(\boldsymbol{x}_t) + \frac{1}{1-\beta}\int_{\mathrm{VaR}_\beta(\boldsymbol{x}_t)}^{\infty} (l - \mathrm{VaR}_\beta(\boldsymbol{x}_t)) f(l)\, dl \\ &= \mathrm{VaR}_\beta(\boldsymbol{x}_t) + \frac{1}{1-\beta}\left(\int_{\mathrm{VaR}_\beta(\boldsymbol{x}_t)}^{\infty} l f(l)\, dl - \mathrm{VaR}_\beta(\boldsymbol{x}_t)\int_{\mathrm{VaR}_\beta(\boldsymbol{x}_t)}^{\infty} f(l)\, dl\right). \end{aligned}$$

Recognizing that $\int_{\mathrm{VaR}_\beta(\boldsymbol{x}_t)}^{\infty} f(l)\, dl = 1-\beta$, we have

$$\begin{aligned} F_\beta(\boldsymbol{x}_t, \mathrm{VaR}_\beta(\boldsymbol{x}_t)) &= \mathrm{VaR}_\beta(\boldsymbol{x}_t) + \frac{1}{1-\beta}\int_{\mathrm{VaR}_\beta(\boldsymbol{x}_t)}^{\infty} l f(l)\, dl - \frac{1}{1-\beta}\mathrm{VaR}_\beta(\boldsymbol{x}_t)(1-\beta) \\ &= \mathrm{VaR}_\beta(\boldsymbol{x}_t) + \mathrm{ES}_\beta(\boldsymbol{x}_t) - \mathrm{VaR}_\beta(\boldsymbol{x}_t) \\ &= \mathrm{ES}_\beta(\boldsymbol{x}_t). \end{aligned}$$

Thus, minimizing $\mathrm{ES}_\beta(\boldsymbol{x}_t)$ with respect to $\boldsymbol{x}_t$ is equivalent to minimizing $F_\beta(\boldsymbol{x}_t, \alpha)$ jointly with respect to $\boldsymbol{x}_t$ and α

$$\min_{\boldsymbol{x}_t} \mathrm{ES}_\beta(\boldsymbol{x}_t) = \min_{\boldsymbol{x}_t, \alpha}\left(\alpha + \frac{1}{1-\beta}\mathbb{E}_t[(-\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1} - \alpha)^+]\right). \tag{A.2}$$

Finally, we show that $F_\beta(\boldsymbol{x}_t, \alpha)$ is jointly convex in $(\boldsymbol{x}_t, \alpha)$. The inner map $(\boldsymbol{x}_t, \alpha) \mapsto (-\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1} - \alpha)$ is affine and the function $z \mapsto z^+$ is convex and nondecreasing. The composition of a convex function with an affine map is convex [2, Section 3.2.2], thus $(-\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1} - \alpha)^+$ is jointly convex in $(\boldsymbol{x}_t, \alpha)$. Expectation is a linear operator which preserves convexity and the sum of convex functions is convex.

A.1.2 Second-Order Expected Shortfall

Second-order Expected Shortfall belongs to the family of higher moment risk measures introduced by Krokhmal [13]. It is obtained from the auxiliary function F_β by replacing the expectation of the excess loss $(L(\boldsymbol{x}_t) - \alpha)^+$ with its L^2 norm.[1] This gives the auxiliary function

$$G_\beta(\boldsymbol{x}_t, \alpha) = \alpha + \frac{1}{1-\beta}\left\|(L(\boldsymbol{x}_t) - \alpha)^+\right\|_{L^2},$$

and second-order Expected Shortfall at confidence level β is defined as its minimum over α

$$\mathrm{SOES}_\beta(\boldsymbol{x}_t) = \min_\alpha G_\beta(\boldsymbol{x}_t, \alpha). \tag{A.3}$$

The L^2 norm squares the excess losses, so large losses receive more weight than in Expected Shortfall. Second-order Expected Shortfall is a coherent risk measure, which we prove in Proposition D.5.

Unlike Expected Shortfall, second-order Expected Shortfall is defined directly through the minimization over α, so there is no equivalence to derive. Instead, we only have to verify that G_β is jointly convex in $(\boldsymbol{x}_t, \alpha)$.

We argued for the case of F_β that the map $(\boldsymbol{x}_t, \alpha) \mapsto (-\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1} - \alpha)^+$ is convex. The L^2 norm is a convex function [2, Section 3.1.5], and nondecreasing for non-negative random variables. The composition of a convex nondecreasing function with a convex function is again convex [2,

[1] We define the L^2 norm in the standard way as $\|Z\|_{L^2} = (\mathbb{E}_t[Z^2])^{1/2}$ for a random variable Z.

Section 3.2.4]. Hence $\|(L(\boldsymbol{x}_t) - \alpha)^+\|_{L^2}$ is jointly convex in $(\boldsymbol{x}_t, \alpha)$, and the sum of convex functions is convex.

By definition (A.3), minimizing $\mathrm{SOES}_\beta(\boldsymbol{x}_t)$ with respect to $\boldsymbol{x}_t$ is equivalent to minimizing $G_\beta(\boldsymbol{x}_t, \alpha)$ jointly with respect to $\boldsymbol{x}_t$ and α

$$\min_{\boldsymbol{x}_t} \mathrm{SOES}_\beta(\boldsymbol{x}_t) = \min_{\boldsymbol{x}_t, \alpha} \left(\alpha + \frac{1}{1-\beta} \|(-\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1} - \alpha)^+\|_{L^2} \right). \tag{A.4}$$

Since both F_β and G_β are jointly convex, any local minimum is global, ensuring that both optimization problems can be solved reliably with standard convex optimization methods.

A.2 Representation for Portfolio Optimization

Let $\rho_\beta \in \{\mathrm{ES}_\beta, \mathrm{SOES}_\beta\}$ denote either of the two risk measures, written as a function of the portfolio return. The portfolio optimization problem is

$$\boldsymbol{x}_t^* = \arg\min_{\boldsymbol{x}_t} \rho_\beta(\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1}), \quad \text{s.t. } \boldsymbol{x}_t^\top \mathbf{1}_N = 1, \tag{A.5}$$

where the budget constraint enforces full investment. By (A.2) and (A.4), problem (A.5) is equivalent to minimizing $F_\beta(\boldsymbol{x}_t, \alpha)$ or $G_\beta(\boldsymbol{x}_t, \alpha)$ jointly with respect to $\boldsymbol{x}_t$ and α, which we write as a linear program for Expected Shortfall in Section A.2.1 and as a second-order cone program for second-order Expected Shortfall in Section A.2.2. Each program in this appendix defines its own decision vector, cost vector and penalty and constraint matrices, so the same symbols denote slightly different objects in different programs, and we refer to a specific program by its equation number.

A.2.1 Expected Shortfall

Let $\boldsymbol{y}_1, \ldots, \boldsymbol{y}_q \in \mathbb{R}^N$ denote q historical realizations of the asset return vector $\boldsymbol{R}_{t+1}$. We approximate the conditional expectation in F_β by the empirical average over this sample

$$\mathbb{E}_t[(-\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1} - \alpha)^+] \approx \frac{1}{q} \sum_{k=1}^{q} (-\boldsymbol{x}_t^\top \boldsymbol{y}_k - \alpha)^+,$$

so that the portfolio optimization problem under Expected Shortfall becomes

$$\boldsymbol{x}_t^*, \alpha^* = \arg\min_{\boldsymbol{x}_t, \alpha} \left\{ \alpha + \frac{1}{q(1-\beta)} \sum_{k=1}^{q} (-\boldsymbol{x}_t^\top \boldsymbol{y}_k - \alpha)^+ \right\}, \quad \text{s.t. } \boldsymbol{x}_t^\top \mathbf{1}_N = 1. \tag{A.6}$$

The summand $(-\boldsymbol{x}_t^\top \boldsymbol{y}_k - \alpha)^+$ is convex but non-linear. We linearize it by introducing, for each observation k, an auxiliary variable $u_k \geq 0$ subject to

$$u_k \geq -\boldsymbol{x}_t^\top \boldsymbol{y}_k - \alpha \qquad \Longleftrightarrow \qquad \boldsymbol{x}_t^\top \boldsymbol{y}_k + \alpha + u_k \geq 0. \tag{A.7}$$

Replacing each summand by u_k, the objective becomes $\alpha + \frac{1}{q(1-\beta)} \sum_{k=1}^{q} u_k$. It is strictly increasing in each u_k and minimized, so every u_k is driven to its lower bound at the optimum, $u_k^* = (-\boldsymbol{x}_t^\top \boldsymbol{y}_k - \alpha)^+$. The slack constraints therefore reproduce the term $(-\boldsymbol{x}_t^\top \boldsymbol{y}_k - \alpha)^+$ while keeping all relations linear.

By collecting the decision variables in the vector

$$\tilde{\boldsymbol{x}}_t = \begin{bmatrix} x_{1,t}, \ldots, x_{N,t}, & \alpha, & u_1, \ldots, u_q \end{bmatrix}^\top = \begin{bmatrix} \boldsymbol{x}_t^\top, & \alpha, & \boldsymbol{u}^\top \end{bmatrix}^\top \in \mathbb{R}^{(N+q+1)},$$

and constants in the vector

$$\boldsymbol{c} = \left[\mathbf{0}_N^\top, \quad 1, \quad \frac{1}{q(1-\beta)}\mathbf{1}_q^\top\right]^\top \in \mathbb{R}^{(N+q+1)},$$

the portfolio optimization problem (A.6) can be written as a linear program

$$\begin{aligned} \tilde{\boldsymbol{x}}_t^* = \arg\min_{\tilde{\boldsymbol{x}}_t} \quad & \boldsymbol{c}^\top \tilde{\boldsymbol{x}}_t \\ \text{s.t.} \quad & \boldsymbol{D}\tilde{\boldsymbol{x}}_t = \begin{bmatrix}\mathbf{1}_N^\top & \mathbf{0}_{q+1}^\top\end{bmatrix} \tilde{\boldsymbol{x}}_t = E = 1, \\ & \boldsymbol{A}\tilde{\boldsymbol{x}}_t = \begin{bmatrix}-\boldsymbol{R}^{\text{data}} & -\mathbf{1}_q & -\boldsymbol{I}_q\end{bmatrix} \tilde{\boldsymbol{x}}_t \leq \boldsymbol{b} = \mathbf{0}_q, \\ & \alpha \geq 0, \quad u_k \geq 0 \quad \text{for } k = 1, \ldots, q. \end{aligned} \tag{A.8}$$

Here $\boldsymbol{D} \in \mathbb{R}^{1\times(N+q+1)}$ encodes the budget constraint, while $\boldsymbol{A} \in \mathbb{R}^{q\times(N+q+1)}$ and $\boldsymbol{b} \in \mathbb{R}^q$ encode the q inequality constraints arising from the linearization (A.7). The matrix $\boldsymbol{R}^{\text{data}} \in \mathbb{R}^{q\times N}$ is comprised of the sampled return vectors $\boldsymbol{y}_k^\top$ row by row, $\mathbf{0}_q$ is the q-dimensional zero vector, and $\boldsymbol{I}_q$ the q-dimensional identity matrix.

Problem (A.8) has a linear objective subject to one linear equality and q linear inequalities, and is thus a linear program solvable with standard methods. This formulation serves as the base unregularized portfolio optimization problem under Expected Shortfall.

A.2.2 Second-Order Expected Shortfall

Using the same historical realizations $\boldsymbol{y}_1, \ldots, \boldsymbol{y}_q$ of Section A.2.1, we approximate the L^2 norm in G_β by its empirical average over the sample

$$\left\|(-\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1} - \alpha)^+\right\|_{L^2} \approx \left(\frac{1}{q}\sum_{k=1}^{q}\left((-\boldsymbol{x}_t^\top \boldsymbol{y}_k - \alpha)^+\right)^2\right)^{1/2},$$

so that the portfolio optimization problem under second-order Expected Shortfall becomes

$$\begin{aligned} \boldsymbol{x}_t^*, \alpha^* = \arg\min_{\boldsymbol{x}_t, \alpha} \left\{\alpha + \frac{1}{1-\beta}\left(\frac{1}{q}\sum_{k=1}^{q}\left((-\boldsymbol{x}_t^\top \boldsymbol{y}_k - \alpha)^+\right)^2\right)^{1/2}\right\}, \\ \text{s.t. } \boldsymbol{x}_t^\top \mathbf{1}_N = 1. \end{aligned} \tag{A.9}$$

As in Section A.2.1, we replace each term $(-\boldsymbol{x}_t^\top \boldsymbol{y}_k - \alpha)^+$ by u_k subject to (A.7). The empirical norm then becomes $\left(\frac{1}{q}\sum_{k=1}^{q} u_k^2\right)^{1/2} = \frac{1}{\sqrt{q}}\|\boldsymbol{u}\|_2$, and the objective becomes $\alpha + \frac{1}{(1-\beta)\sqrt{q}}\|\boldsymbol{u}\|_2$. The Euclidean norm is strictly increasing in every positive u_k, so, as for Expected Shortfall, every u_k is driven to its lower bound at the optimum, $u_k^* = (-\boldsymbol{x}_t^\top \boldsymbol{y}_k - \alpha)^+$.

The objective is still not linear, so we introduce one more auxiliary variable s subject to the constraint $\|\boldsymbol{u}\|_2 \leq s$ and minimize $\alpha + \frac{1}{(1-\beta)\sqrt{q}}s$ instead. Since $\frac{1}{(1-\beta)\sqrt{q}} > 0$, $s^* = \|\boldsymbol{u}^*\|_2$ at the optimum, and the objective is linear in the decision variables. The constraint $\|\boldsymbol{u}\|_2 \leq s$ requires the vector $(\boldsymbol{u}, s) \in \mathbb{R}^{q+1}$ to lie in the second-order cone, the set of vectors whose last entry is at least the Euclidean norm of the remaining entries [2, Section 2.2.3]. This set is a cone, since it contains every non-negative multiple of its elements, and it is convex, so the constraint keeps the optimization problem convex. An optimization problem with a linear objective, linear constraints and second-order cone constraints is a second-order cone program [2, Section 4.4.2].

By collecting the decision variables in the vector

$$\tilde{\boldsymbol{x}}_t = \begin{bmatrix} \boldsymbol{x}_t^\top, & \alpha, & \boldsymbol{u}^\top, & s \end{bmatrix}^\top \in \mathbb{R}^{(N+q+2)},$$

and constants in the vector

$$\boldsymbol{c} = \begin{bmatrix} \mathbf{0}_N^\top, & 1, & \mathbf{0}_q^\top, & \dfrac{1}{(1-\beta)\sqrt{q}} \end{bmatrix}^\top \in \mathbb{R}^{(N+q+2)},$$

the portfolio optimization problem (A.9) can be written as a second-order cone program

$$\begin{aligned} \tilde{\boldsymbol{x}}_t^* = \arg\min_{\tilde{\boldsymbol{x}}_t} \quad & \boldsymbol{c}^\top \tilde{\boldsymbol{x}}_t \\ \text{s.t.} \quad & \boldsymbol{D}\tilde{\boldsymbol{x}}_t = \begin{bmatrix} \mathbf{1}_N^\top & \mathbf{0}_{q+2}^\top \end{bmatrix} \tilde{\boldsymbol{x}}_t = E = 1, \\ & \boldsymbol{A}\tilde{\boldsymbol{x}}_t = \begin{bmatrix} -\boldsymbol{R}^{\text{data}} & -\mathbf{1}_q & -\boldsymbol{I}_q & \mathbf{0}_q \end{bmatrix} \tilde{\boldsymbol{x}}_t \leq \boldsymbol{b} = \mathbf{0}_q, \\ & \|\boldsymbol{u}\|_2 \leq s, \\ & \alpha \geq 0, \quad u_k \geq 0 \quad \text{for } k = 1, \ldots, q. \end{aligned} \tag{A.10}$$

Here $\boldsymbol{D} \in \mathbb{R}^{1\times(N+q+2)}$ and $\boldsymbol{A} \in \mathbb{R}^{q\times(N+q+2)}$ are the constraint matrices of (A.8) with a zero column added for the variable s. Problem (A.10) has a linear objective subject to one linear equality, q linear inequalities and one second-order cone constraint, and is thus a second-order cone program. This formulation serves as the base unregularized portfolio optimization problem under second-order Expected Shortfall.

A.3 Regularization

The base problems (A.8) and (A.10) are known to be numerically unstable in high-dimensional settings, producing extreme and concentrated weights [5, 12]. We stabilize them with three regularization techniques, an L^1-norm penalty that promotes sparsity, an L^2-norm penalty that promotes shrinkage, and an asset leverage constraint that bounds gross exposure. Each is appended to the base problems while keeping them convex.

A.3.1 L^1 Regularization

The L^1 penalty $\lambda_1\|\boldsymbol{x}_t\|_1 = \lambda_1 \sum_{i=1}^N |x_{i,t}|$, with $\lambda_1 \geq 0$, is non-linear through the absolute value. Following [29], we linearize it by splitting each weight into its positive and negative parts via the auxiliary variables $\Delta_{i,t}^+ \geq 0$ and $\Delta_{i,t}^- \geq 0$

$$x_{i,t} = \Delta_{i,t}^+ - \Delta_{i,t}^-, \qquad |x_{i,t}| = \Delta_{i,t}^+ + \Delta_{i,t}^-. \tag{A.11}$$

The second identity holds at the optimum. Minimizing the L^1 penalty drives the $\Delta_{i,t}^+$ and $\Delta_{i,t}^-$ to their smallest feasible values, so at most one is non-zero, giving $\Delta_{i,t}^+ = \max(x_{i,t}, 0)$ and $\Delta_{i,t}^- = \max(-x_{i,t}, 0)$. Writing the positive and negative parts as vectors $\boldsymbol{\Delta}_t^+ = [\Delta_{1,t}^+, \ldots, \Delta_{N,t}^+]^\top$ and $\boldsymbol{\Delta}_t^- = [\Delta_{1,t}^-, \ldots, \Delta_{N,t}^-]^\top$, the L^1 penalty is written as

$$\lambda_1\|\boldsymbol{x}_t\|_1 = \lambda_1 \mathbf{1}_N^\top(\boldsymbol{\Delta}_t^+ + \boldsymbol{\Delta}_t^-).$$

We append $\boldsymbol{\Delta}_t^-$ and $\boldsymbol{\Delta}_t^+$ to the decision vector, with the penalty coefficient λ_1 on each of their entries in the cost vector. The decomposition (A.11) is enforced as N additional linear equality constraints, $\boldsymbol{x}_t + \boldsymbol{\Delta}_t^- - \boldsymbol{\Delta}_t^+ = \mathbf{0}_N$, alongside the original budget constraint.

A.3.2 L^2 Regularization

The L^2 penalty $\lambda_2\|\boldsymbol{x}_t\|_2^2 = \lambda_2\boldsymbol{x}_t^\top\boldsymbol{x}_t$, with $\lambda_2 \geq 0$, is quadratic in the weights and is written as a quadratic form

$$\lambda_2\|\boldsymbol{x}_t\|_2^2 = \boldsymbol{x}_t^\top \boldsymbol{Q}\,\boldsymbol{x}_t, \qquad \boldsymbol{Q} = \lambda_2\boldsymbol{I}_N.$$

Embedding this in the decision vector $\tilde{\boldsymbol{x}}_t$, the penalty becomes $\tilde{\boldsymbol{x}}_t^\top\tilde{\boldsymbol{Q}}\,\tilde{\boldsymbol{x}}_t$, where $\tilde{\boldsymbol{Q}}$ places the block $\lambda_2\boldsymbol{I}_N$ on the asset weights and zeros for the rest of the variables. Adding this term to the linear objective of the base problems makes the objective quadratic, and since $\tilde{\boldsymbol{Q}}$ is positive semi-definite, it remains convex.

A.3.3 Leverage Constraint

The leverage constraint $\|\boldsymbol{x}_t\|_1 \leq \ell^a$, with $\ell^a \geq 1$, bounds the portfolio's gross exposure and guarantees a bounded feasible region even when $\lambda_1 = 0$. Reusing the decomposition (A.11), it is expressed as a single linear inequality

$$\|\boldsymbol{x}_t\|_1 = \mathbf{1}_N^\top(\boldsymbol{\Delta}_t^+ + \boldsymbol{\Delta}_t^-) \leq \ell^a. \tag{A.12}$$

The choice $\ell^a = 1$ rules out leverage entirely and is equivalent to a hard L^1 constraint.

A.3.4 Expected Shortfall

Adding all three regularization terms to the base problem (A.8) extends the decision vector to

$$\tilde{\boldsymbol{x}}_t = \left[\boldsymbol{x}_t^\top,\quad \alpha,\quad \boldsymbol{u}^\top,\quad (\boldsymbol{\Delta}_t^-)^\top,\quad (\boldsymbol{\Delta}_t^+)^\top\right]^\top \in \mathbb{R}^{3N+q+1},$$

with the cost vector

$$\tilde{\boldsymbol{c}} = \left[\mathbf{0}_N^\top,\quad 1,\quad \tfrac{1}{q(1-\beta)}\mathbf{1}_q^\top,\quad \lambda_1\mathbf{1}_N^\top,\quad \lambda_1\mathbf{1}_N^\top\right]^\top$$

and the penalty matrix

$$\tilde{\boldsymbol{Q}} = \begin{bmatrix} \lambda_2\boldsymbol{I}_N & \mathbf{0}_{N\times(2N+q+1)} \\ \mathbf{0}_{(2N+q+1)\times N} & \mathbf{0}_{(2N+q+1)\times(2N+q+1)} \end{bmatrix}.$$

The portfolio optimization problem under Expected Shortfall then becomes a quadratic program with linear constraints

$$\begin{aligned} \tilde{\boldsymbol{x}}_t^* &= \arg\min_{\tilde{\boldsymbol{x}}_t} \quad \tilde{\boldsymbol{c}}^\top\tilde{\boldsymbol{x}}_t + \tilde{\boldsymbol{x}}_t^\top\tilde{\boldsymbol{Q}}\,\tilde{\boldsymbol{x}}_t, \\ \text{s.t.}\quad \tilde{\boldsymbol{D}}\tilde{\boldsymbol{x}}_t &= \begin{bmatrix} \mathbf{1}_N^\top & \mathbf{0}_{q+1}^\top & \mathbf{0}_N^\top & \mathbf{0}_N^\top \\ \boldsymbol{I}_N & \mathbf{0}_{N\times(q+1)} & \boldsymbol{I}_N & -\boldsymbol{I}_N \end{bmatrix}\tilde{\boldsymbol{x}}_t = \tilde{\boldsymbol{E}} = \begin{bmatrix} 1 \\ \mathbf{0}_N \end{bmatrix}, \\ \tilde{\boldsymbol{A}}\tilde{\boldsymbol{x}}_t &= \begin{bmatrix} -\boldsymbol{R}^{\text{data}} & -\mathbf{1}_q & -\boldsymbol{I}_q & \mathbf{0}_{q\times N} & \mathbf{0}_{q\times N} \\ \mathbf{0}_N^\top & 0 & \mathbf{0}_q^\top & \mathbf{1}_N^\top & \mathbf{1}_N^\top \end{bmatrix}\tilde{\boldsymbol{x}}_t \leq \boldsymbol{b} = \begin{bmatrix} \mathbf{0}_q \\ \ell^a \end{bmatrix}, \\ \alpha \geq 0,\quad u_k \geq 0,\quad \Delta_{i,t}^- \geq 0,\quad \Delta_{i,t}^+ \geq 0 \quad &\text{for } i = 1,\ldots,N,\ k = 1,\ldots,q. \end{aligned} \tag{A.13}$$

The first row of $\tilde{\boldsymbol{D}}$ enforces the budget constraint and the remaining N rows enforce the weight decomposition (A.11). The first q rows of $\tilde{\boldsymbol{A}}$ encode the inequality constraints from the linearization (A.7), extended with zeros for the new variables, and the final row of $\tilde{\boldsymbol{A}}$ imposes the leverage constraint (A.12).

A.3.5 Second-Order Expected Shortfall

Adding the same three regularization terms to the base problem (A.10) extends the decision vector to

$$\tilde{\boldsymbol{x}}_t = \left[\boldsymbol{x}_t^\top,\ \ \alpha,\ \ \boldsymbol{u}^\top,\ \ (\boldsymbol{\Delta}_t^-)^\top,\ \ (\boldsymbol{\Delta}_t^+)^\top,\ \ s\right]^\top \in \mathbb{R}^{3N+q+2},$$

with the cost vector

$$\tilde{\boldsymbol{c}} = \left[\mathbf{0}_N^\top,\ \ 1,\ \ \mathbf{0}_q^\top,\ \ \lambda_1 \mathbf{1}_N^\top,\ \ \lambda_1 \mathbf{1}_N^\top,\ \ \tfrac{1}{(1-\beta)\sqrt{q}}\right]^\top$$

and the penalty matrix

$$\tilde{\boldsymbol{Q}} = \begin{bmatrix} \lambda_2 \boldsymbol{I}_N & \mathbf{0}_{N\times(2N+q+2)} \\ \mathbf{0}_{(2N+q+2)\times N} & \mathbf{0}_{(2N+q+2)\times(2N+q+2)} \end{bmatrix}.$$

The portfolio optimization problem under second-order Expected Shortfall then becomes a second-order cone program with a convex quadratic objective

$$\begin{aligned} \tilde{\boldsymbol{x}}_t^* = \arg\min_{\tilde{\boldsymbol{x}}_t} \quad & \tilde{\boldsymbol{c}}^\top \tilde{\boldsymbol{x}}_t + \tilde{\boldsymbol{x}}_t^\top \tilde{\boldsymbol{Q}}\, \tilde{\boldsymbol{x}}_t, \\ \text{s.t.} \quad \tilde{\boldsymbol{D}}\tilde{\boldsymbol{x}}_t &= \begin{bmatrix} \mathbf{1}_N^\top & \mathbf{0}_{q+1}^\top & \mathbf{0}_N^\top & \mathbf{0}_N^\top & 0 \\ \boldsymbol{I}_N & \mathbf{0}_{N\times(q+1)} & \boldsymbol{I}_N & -\boldsymbol{I}_N & \mathbf{0}_N \end{bmatrix} \tilde{\boldsymbol{x}}_t = \tilde{\boldsymbol{E}} = \begin{bmatrix} 1 \\ \mathbf{0}_N \end{bmatrix}, \\ \tilde{\boldsymbol{A}}\tilde{\boldsymbol{x}}_t &= \begin{bmatrix} -\boldsymbol{R}^{\text{data}} & -\mathbf{1}_q & -\boldsymbol{I}_q & \mathbf{0}_{q\times N} & \mathbf{0}_{q\times N} & \mathbf{0}_q \\ \mathbf{0}_N^\top & 0 & \mathbf{0}_q^\top & \mathbf{1}_N^\top & \mathbf{1}_N^\top & 0 \end{bmatrix} \tilde{\boldsymbol{x}}_t \le \boldsymbol{b} = \begin{bmatrix} \mathbf{0}_q \\ \ell^a \end{bmatrix}, \\ & \|\boldsymbol{u}\|_2 \le s, \\ \alpha \ge 0, \quad u_k \ge 0, \quad & \Delta_{i,t}^- \ge 0, \quad \Delta_{i,t}^+ \ge 0 \quad \text{for } i = 1,\dots,N,\ k = 1,\dots,q. \end{aligned} \tag{A.14}$$

Here $\tilde{\boldsymbol{D}}$ and $\tilde{\boldsymbol{A}}$ are the constraint matrices of (A.13) with a zero column added for the variable s, and the cone constraint is the one of (A.10).

A.4 Sequential Optimization

So far, we have formulated the portfolio optimization problem over the asset weights only. An international portfolio, however, holds assets and currency forward positions, so constructing it also requires selecting currency exposures. In the sequential portfolio optimization approach, the asset weights and currency exposures are optimized in two steps. First, the asset weights $\boldsymbol{x}_t \in \mathbb{R}^N$ are optimized on the fully hedged returns, then, conditional on the optimal $\boldsymbol{x}_t^*$, the currency exposures $\boldsymbol{\psi}_t \in \mathbb{R}^M$ are optimized. The two optimization steps are given by

$$\begin{aligned} \boldsymbol{x}_t^* = \arg\min_{\boldsymbol{x}_t} & \left\{ \rho_\beta\left(\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1}^{fh}\right) + \lambda_1^a \|\boldsymbol{x}_t\|_1 + \lambda_2^a \|\boldsymbol{x}_t\|_2^2 \right\} \\ & \text{s.t. } \boldsymbol{x}_t^\top \mathbf{1}_N = 1,\ \|\boldsymbol{x}_t\|_1 \le \ell^a, \end{aligned} \tag{A.15}$$

$$\boldsymbol{\psi}_t^* \mid \boldsymbol{x}_t^* = \arg\min_{\boldsymbol{\psi}_t} \left\{ \rho_\beta\left(\boldsymbol{x}_t^{*\top} \boldsymbol{R}_{t+1}^{fh} + \boldsymbol{\psi}_t^\top (\boldsymbol{e}_{t+1} - \boldsymbol{f}_t)\right) + \lambda_1^c \|\boldsymbol{\psi}_t\|_1 + \lambda_2^c \|\boldsymbol{\psi}_t\|_2^2 \right\}. \tag{A.16}$$

The regularization parameters are kept separate across assets $(\lambda_1^a, \lambda_2^a)$ and currencies $(\lambda_1^c, \lambda_2^c)$.

The fully hedged return vector $\boldsymbol{R}_{t+1}^{fh}$ (defined in Section 3.1.3) takes the place of the asset return vector $\boldsymbol{R}_{t+1}$ of Sections A.1–A.3, so the asset step (A.15) is identical to the regularized programs (A.13) for Expected Shortfall and (A.14) for second-order Expected Shortfall, and

is obtained from them by substituting the matrix of historically sampled fully hedged returns $\boldsymbol{R}^{fh} \in \mathbb{R}^{q\times N}$ for $\boldsymbol{R}^{\text{data}}$ and the asset-specific penalties λ_1^a, λ_2^a for λ_1, λ_2.

In the currency step (A.16), the asset weights $\boldsymbol{x}_t^*$ are fixed and the currency exposures $\boldsymbol{\psi}_t$ are optimized. The currency exposure $\psi_{c,t}$ is the net exposure of the portfolio to the foreign currency c as defined in Section 3.1.3, and each unit of net exposure earns the excess currency return $e_{c,t+1} - f_{c,t}$. Using the same auxiliary representation for the risk measures and the empirical approximation as in Section A.2, we introduce α_ψ and auxiliary variables $u_{k,\psi} \geq 0$, and the slack constraint (A.7) becomes

$$u_{k,\psi} \geq -\boldsymbol{\psi}_t^\top(\boldsymbol{e}_k - \boldsymbol{f}_k) - \boldsymbol{x}_t^{*\top}\boldsymbol{R}_k^{fh} - \alpha_\psi \quad \Longleftrightarrow \quad -\boldsymbol{\psi}_t^\top(\boldsymbol{e}_k - \boldsymbol{f}_k) - \alpha_\psi - u_{k,\psi} \leq \boldsymbol{x}_t^{*\top}\boldsymbol{R}_k^{fh},$$

where $\boldsymbol{R}_k^{fh}$ is the k-th sampled fully hedged return vector, and $\boldsymbol{e}_k$ and $\boldsymbol{f}_k$ are the k-th sampled currency return and forward premium vectors. We stack the sampled excess currency returns $(\boldsymbol{e}_k - \boldsymbol{f}_k)^\top$ row by row into the matrix $\boldsymbol{R}^c \in \mathbb{R}^{q\times M}$, and collect the fixed asset-portfolio returns in $\boldsymbol{R}_P^{fh} = \boldsymbol{R}^{fh}\boldsymbol{x}_t^* \in \mathbb{R}^q$.

The currency step has no budget constraint, since the currency overlay is a zero-investment portfolio whose domestic exposure $\psi_{0,t} = -\mathbf{1}_M^\top\boldsymbol{\psi}_t$ is inferred residually, and no leverage constraint is imposed on the currency exposures. We linearize the currency L^1 penalty with the decomposition $\psi_{c,t} = \Delta^+_{\psi,c,t} - \Delta^-_{\psi,c,t}$ as in Section A.3.1, and write the currency L^2 penalty as a quadratic form with the block $\lambda_2^c\boldsymbol{I}_M$ as in Section A.3.2.

A.4.1 Currency Overlay under Expected Shortfall

For Expected Shortfall, the decision vector of the currency step is

$$\tilde{\boldsymbol{\psi}}_t = \left[\boldsymbol{\psi}_t^\top, \quad \alpha_\psi, \quad \boldsymbol{u}_\psi^\top, \quad (\boldsymbol{\Delta}^-_{\psi,t})^\top, \quad (\boldsymbol{\Delta}^+_{\psi,t})^\top\right]^\top \in \mathbb{R}^{3M+q+1},$$

with the cost vector

$$\tilde{\boldsymbol{c}} = \left[\mathbf{0}_M^\top, \quad 1, \quad \tfrac{1}{q(1-\beta)}\mathbf{1}_q^\top, \quad \lambda_1^c\mathbf{1}_M^\top, \quad \lambda_1^c\mathbf{1}_M^\top\right]^\top$$

and the penalty matrix

$$\tilde{\boldsymbol{Q}} = \begin{bmatrix} \lambda_2^c\boldsymbol{I}_M & \mathbf{0}_{M\times(2M+q+1)} \\ \mathbf{0}_{(2M+q+1)\times M} & \mathbf{0}_{(2M+q+1)\times(2M+q+1)} \end{bmatrix}.$$

The currency overlay optimization under Expected Shortfall then becomes a quadratic program with linear constraints

$$\begin{aligned} \tilde{\boldsymbol{\psi}}_t^* &= \arg\min_{\tilde{\boldsymbol{\psi}}_t} \quad \tilde{\boldsymbol{c}}^\top\tilde{\boldsymbol{\psi}}_t + \tilde{\boldsymbol{\psi}}_t^\top\tilde{\boldsymbol{Q}}\,\tilde{\boldsymbol{\psi}}_t, \\ \text{s.t.} \quad \tilde{\boldsymbol{D}}\tilde{\boldsymbol{\psi}}_t &= \left[\boldsymbol{I}_M \quad \mathbf{0}_{M\times(q+1)} \quad \boldsymbol{I}_M \quad -\boldsymbol{I}_M\right]\tilde{\boldsymbol{\psi}}_t = \tilde{\boldsymbol{E}} = \mathbf{0}_M, \\ \tilde{\boldsymbol{A}}\tilde{\boldsymbol{\psi}}_t &= \left[-\boldsymbol{R}^c \quad -\mathbf{1}_q \quad -\boldsymbol{I}_q \quad \mathbf{0}_{q\times M} \quad \mathbf{0}_{q\times M}\right]\tilde{\boldsymbol{\psi}}_t \leq \boldsymbol{b} = \boldsymbol{R}_P^{fh}, \\ \alpha_\psi \geq 0, \quad u_{k,\psi} \geq 0, \quad \Delta^-_{\psi,c,t} \geq 0, \quad \Delta^+_{\psi,c,t} \geq 0 \quad &\text{for } c = 1,\ldots,M,\ k = 1,\ldots,q. \end{aligned} \tag{A.17}$$

The M rows of $\tilde{\boldsymbol{D}}$ enforce the weight decomposition (A.11) for the currency exposures, and the q rows of $\tilde{\boldsymbol{A}}$ encode the slack constraints of the currency step, with the fixed asset-portfolio returns $\boldsymbol{R}_P^{fh}$ on the right-hand side.

A.4.2 Currency Overlay under Second-Order Expected Shortfall

For second-order Expected Shortfall, we introduce, as in Section A.2.2, one more auxiliary variable s_ψ subject to the cone constraint $\|\boldsymbol{u}_\psi\|_2 \leq s_\psi$, which extends the decision vector to

$$\tilde{\boldsymbol{\psi}}_t = \Big[\boldsymbol{\psi}_t^\top, \quad \alpha_\psi, \quad \boldsymbol{u}_\psi^\top, \quad (\boldsymbol{\Delta}_{\psi,t}^-)^\top, \quad (\boldsymbol{\Delta}_{\psi,t}^+)^\top, \quad s_\psi\Big]^\top \in \mathbb{R}^{3M+q+2},$$

with the cost vector

$$\tilde{\boldsymbol{c}} = \Big[\boldsymbol{0}_M^\top, \quad 1, \quad \boldsymbol{0}_q^\top, \quad \lambda_1^c \boldsymbol{1}_M^\top, \quad \lambda_1^c \boldsymbol{1}_M^\top, \quad \tfrac{1}{(1-\beta)\sqrt{q}}\Big]^\top$$

and the penalty matrix

$$\tilde{\boldsymbol{Q}} = \begin{bmatrix} \lambda_2^c \boldsymbol{I}_M & \boldsymbol{0}_{M\times(2M+q+2)} \\ \boldsymbol{0}_{(2M+q+2)\times M} & \boldsymbol{0}_{(2M+q+2)\times(2M+q+2)} \end{bmatrix}.$$

The currency overlay optimization under second-order Expected Shortfall then becomes a second-order cone program with a convex quadratic objective

$$\begin{aligned} \tilde{\boldsymbol{\psi}}_t^* = \underset{\tilde{\boldsymbol{\psi}}_t}{\arg\min} \quad & \tilde{\boldsymbol{c}}^\top \tilde{\boldsymbol{\psi}}_t + \tilde{\boldsymbol{\psi}}_t^\top \tilde{\boldsymbol{Q}}\, \tilde{\boldsymbol{\psi}}_t, \\ \text{s.t.} \quad & \tilde{\boldsymbol{D}}\tilde{\boldsymbol{\psi}}_t = \Big[\boldsymbol{I}_M \quad \boldsymbol{0}_{M\times(q+1)} \quad \boldsymbol{I}_M \quad -\boldsymbol{I}_M \quad \boldsymbol{0}_M\Big]\, \tilde{\boldsymbol{\psi}}_t = \tilde{\boldsymbol{E}} = \boldsymbol{0}_M, \\ & \tilde{\boldsymbol{A}}\tilde{\boldsymbol{\psi}}_t = \Big[-\boldsymbol{R}^c \quad -\boldsymbol{1}_q \quad -\boldsymbol{I}_q \quad \boldsymbol{0}_{q\times M} \quad \boldsymbol{0}_{q\times M} \quad \boldsymbol{0}_q\Big]\, \tilde{\boldsymbol{\psi}}_t \leq \boldsymbol{b} = \boldsymbol{R}_P^{fh}, \\ & \|\boldsymbol{u}_\psi\|_2 \leq s_\psi, \\ & \alpha_\psi \geq 0, \quad u_{k,\psi} \geq 0, \quad \Delta_{\psi,c,t}^- \geq 0, \quad \Delta_{\psi,c,t}^+ \geq 0 \quad \text{for } c = 1,\dots,M,\ k = 1,\dots,q. \end{aligned} \tag{A.18}$$

Here $\tilde{\boldsymbol{D}}$ and $\tilde{\boldsymbol{A}}$ are the constraint matrices of (A.17) with a zero column added for the variable s_ψ.

A.5 Joint Optimization

In the joint portfolio optimization approach, the asset weights and currency exposures are optimized simultaneously rather than sequentially. We stack them into the combined decision vector $\boldsymbol{\theta}_t = [\boldsymbol{x}_t^\top, \boldsymbol{\psi}_t^\top]^\top \in \mathbb{R}^{N+M}$ and the corresponding returns into the combined return vector $\boldsymbol{R}_{t+1}^{\text{all}} = [\boldsymbol{R}_{t+1}^{fh\top}, (\boldsymbol{e}_{t+1} - \boldsymbol{f}_t)^\top]^\top \in \mathbb{R}^{N+M}$, so that $\boldsymbol{\theta}_t^\top \boldsymbol{R}_{t+1}^{\text{all}} = \boldsymbol{x}_t^\top \boldsymbol{R}_{t+1}^{fh} + \boldsymbol{\psi}_t^\top(\boldsymbol{e}_{t+1} - \boldsymbol{f}_t)$ is the hedged portfolio return of Section A.4. The joint optimization problem is

$$\begin{aligned} \boldsymbol{\theta}_t^* = \underset{\boldsymbol{\theta}_t}{\arg\min} \Big\{ \rho_\beta\Big(\boldsymbol{\theta}_t^\top \boldsymbol{R}_{t+1}^{\text{all}}\Big) + \|\boldsymbol{S}_1\boldsymbol{\theta}_t\|_1 + \|\boldsymbol{S}_2\boldsymbol{\theta}_t\|_2^2 \Big\} \\ \text{s.t. } \boldsymbol{\theta}_t^\top \boldsymbol{q}_{N,M} = 1, \quad \|\boldsymbol{L}^a\boldsymbol{\theta}_t\|_1 \leq \ell^a. \end{aligned} \tag{A.19}$$

The following matrices and vector adapt the problem to the joint optimization setting

$$\boldsymbol{q}_{N,M} = \begin{bmatrix} \boldsymbol{1}_N \\ \boldsymbol{0}_M \end{bmatrix}, \quad \boldsymbol{L}^a = \begin{bmatrix} \boldsymbol{I}_N & \boldsymbol{0}_{N\times M} \\ \boldsymbol{0}_{M\times N} & \boldsymbol{0}_{M\times M} \end{bmatrix},$$

$$\boldsymbol{S}_1 = \begin{bmatrix} \lambda_1^a \boldsymbol{I}_N & \boldsymbol{0}_{N\times M} \\ \boldsymbol{0}_{M\times N} & \lambda_1^c \boldsymbol{I}_M \end{bmatrix}, \quad \boldsymbol{S}_2 = \begin{bmatrix} \sqrt{\lambda_2^a}\, \boldsymbol{I}_N & \boldsymbol{0}_{N\times M} \\ \boldsymbol{0}_{M\times N} & \sqrt{\lambda_2^c}\, \boldsymbol{I}_M \end{bmatrix}.$$

The vector $\boldsymbol{q}_{N,M}$ applies the budget constraint to the asset weights only, so that $\boldsymbol{\theta}_t^\top \boldsymbol{q}_{N,M} = \boldsymbol{x}_t^\top \boldsymbol{1}_N = 1$, leaving the currency exposures as a zero-investment overlay. The matrix $\boldsymbol{L}^a$ applies

the leverage constraint to the asset weights only, giving $\|\boldsymbol{L}^a\boldsymbol{\theta}_t\|_1 = \|\boldsymbol{x}_t\|_1$. The diagonal regularization matrices $\boldsymbol{S}_1$ and $\boldsymbol{S}_2$ allow the asset and currency penalties to differ within a single objective,

$$\|\boldsymbol{S}_1\boldsymbol{\theta}_t\|_1 = \lambda_1^a\|\boldsymbol{x}_t\|_1 + \lambda_1^c\|\boldsymbol{\psi}_t\|_1, \qquad \|\boldsymbol{S}_2\boldsymbol{\theta}_t\|_2^2 = \lambda_2^a\|\boldsymbol{x}_t\|_2^2 + \lambda_2^c\|\boldsymbol{\psi}_t\|_2^2.$$

Let $\boldsymbol{R}^{\text{all}} = [\,\boldsymbol{R}^{fh} \;\; \boldsymbol{R}^c\,] \in \mathbb{R}^{q\times(N+M)}$ denote the matrix of jointly sampled asset and currency returns, whose k-th row $(\boldsymbol{R}_k^{\text{all}})^\top = [(\boldsymbol{R}_k^{fh})^\top, (\boldsymbol{e}_k - \boldsymbol{f}_k)^\top]$ is the k-th historical realization of $\boldsymbol{R}_{t+1}^{\text{all}}$. We replace the expectation in F_β and the L^2 norm in G_β by their empirical counterparts as in Section A.2, with the joint return $\boldsymbol{\theta}_t^\top\boldsymbol{R}_{t+1}^{\text{all}}$ in place of $\boldsymbol{x}_t^\top\boldsymbol{R}_{t+1}$. Introducing α and the auxiliary variables $u_k \geq 0$ for $k = 1,\dots,q$ as in Section A.2, the slack constraint (A.7) becomes

$$u_k \geq -\boldsymbol{\theta}_t^\top\boldsymbol{R}_k^{\text{all}} - \alpha \qquad\Longleftrightarrow\qquad \boldsymbol{\theta}_t^\top\boldsymbol{R}_k^{\text{all}} + \alpha + u_k \geq 0.$$

To linearize the L^1 penalty, we decompose all $N+M$ components of $\boldsymbol{\theta}_t$ into positive and negative parts as in Section A.3.1, $\boldsymbol{\theta}_t = \boldsymbol{\Delta}_t^+ - \boldsymbol{\Delta}_t^-$ with $\boldsymbol{\Delta}_t^\pm = [(\boldsymbol{\Delta}_{x,t}^\pm)^\top, (\boldsymbol{\Delta}_{\psi,t}^\pm)^\top]^\top$, so that at the optimum

$$\begin{aligned}\|\boldsymbol{S}_1\boldsymbol{\theta}_t\|_1 &= \lambda_1^a\mathbf{1}_N^\top(\boldsymbol{\Delta}_{x,t}^+ + \boldsymbol{\Delta}_{x,t}^-) + \lambda_1^c\mathbf{1}_M^\top(\boldsymbol{\Delta}_{\psi,t}^+ + \boldsymbol{\Delta}_{\psi,t}^-),\\ \|\boldsymbol{L}^a\boldsymbol{\theta}_t\|_1 &= \boldsymbol{q}_{N,M}^\top(\boldsymbol{\Delta}_t^+ + \boldsymbol{\Delta}_t^-) \leq \ell^a.\end{aligned}$$

The L^2 penalty is the quadratic form $\|\boldsymbol{S}_2\boldsymbol{\theta}_t\|_2^2 = \boldsymbol{\theta}_t^\top\boldsymbol{S}_2^2\,\boldsymbol{\theta}_t$, where $\boldsymbol{S}_2^2$ is diagonal with the blocks $\lambda_2^a\boldsymbol{I}_N$ and $\lambda_2^c\boldsymbol{I}_M$, which enter the penalty matrix $\tilde{\boldsymbol{Q}}$ of the two programs below.

A.5.1 Joint Optimization under Expected Shortfall

For Expected Shortfall, the decision vector is

$$\tilde{\boldsymbol{\theta}}_t = \Big[\,\boldsymbol{x}_t^\top,\;\; \boldsymbol{\psi}_t^\top,\;\; \alpha,\;\; \boldsymbol{u}^\top,\;\; (\boldsymbol{\Delta}_{x,t}^-)^\top,\;\; (\boldsymbol{\Delta}_{\psi,t}^-)^\top,\;\; (\boldsymbol{\Delta}_{x,t}^+)^\top,\;\; (\boldsymbol{\Delta}_{\psi,t}^+)^\top\,\Big]^\top \in \mathbb{R}^{3N+3M+q+1},$$

with the cost vector

$$\tilde{\boldsymbol{c}} = \Big[\,\mathbf{0}_{N+M}^\top,\;\; 1,\;\; \tfrac{1}{q(1-\beta)}\mathbf{1}_q^\top,\;\; \lambda_1^a\mathbf{1}_N^\top,\;\; \lambda_1^c\mathbf{1}_M^\top,\;\; \lambda_1^a\mathbf{1}_N^\top,\;\; \lambda_1^c\mathbf{1}_M^\top\,\Big]^\top$$

and the penalty matrix

$$\tilde{\boldsymbol{Q}} = \begin{bmatrix} \lambda_2^a\boldsymbol{I}_N & \mathbf{0}_{N\times M} & \mathbf{0}_{N\times(2N+2M+q+1)} \\ \mathbf{0}_{M\times N} & \lambda_2^c\boldsymbol{I}_M & \mathbf{0}_{M\times(2N+2M+q+1)} \\ \mathbf{0}_{(2N+2M+q+1)\times N} & \mathbf{0}_{(2N+2M+q+1)\times M} & \mathbf{0}_{(2N+2M+q+1)\times(2N+2M+q+1)} \end{bmatrix}.$$

The joint portfolio optimization problem under Expected Shortfall then becomes a quadratic program with linear constraints

$$\begin{aligned} \tilde{\boldsymbol{\theta}}_t^* &= \arg\min_{\tilde{\boldsymbol{\theta}}_t} \quad \tilde{\boldsymbol{c}}^\top\tilde{\boldsymbol{\theta}}_t + \tilde{\boldsymbol{\theta}}_t^\top\tilde{\boldsymbol{Q}}\,\tilde{\boldsymbol{\theta}}_t, \\ \text{s.t.}\quad \tilde{\boldsymbol{D}}\tilde{\boldsymbol{\theta}}_t &= \begin{bmatrix} \boldsymbol{q}_{N,M}^\top & \mathbf{0}_{q+1}^\top & \mathbf{0}_{N+M}^\top & \mathbf{0}_{N+M}^\top \\ \boldsymbol{I}_{N+M} & \mathbf{0}_{(N+M)\times(q+1)} & \boldsymbol{I}_{N+M} & -\boldsymbol{I}_{N+M} \end{bmatrix}\tilde{\boldsymbol{\theta}}_t = \tilde{\boldsymbol{E}} = \begin{bmatrix} 1 \\ \mathbf{0}_{N+M} \end{bmatrix}, \\ \tilde{\boldsymbol{A}}\tilde{\boldsymbol{\theta}}_t &= \begin{bmatrix} -\boldsymbol{R}^{\text{all}} & -\mathbf{1}_q & -\boldsymbol{I}_q & \mathbf{0}_{q\times(N+M)} & \mathbf{0}_{q\times(N+M)} \\ \mathbf{0}_{N+M}^\top & 0 & \mathbf{0}_q^\top & \boldsymbol{q}_{N,M}^\top & \boldsymbol{q}_{N,M}^\top \end{bmatrix}\tilde{\boldsymbol{\theta}}_t \leq \boldsymbol{b} = \begin{bmatrix} \mathbf{0}_q \\ \ell^a \end{bmatrix}, \\ &\alpha \geq 0, \quad u_k \geq 0 \quad \text{for } k = 1,\dots,q, \\ &\Delta_{x,i,t}^\pm \geq 0, \quad \Delta_{\psi,c,t}^\pm \geq 0 \quad \text{for } i = 1,\dots,N,\; c = 1,\dots,M. \end{aligned} \tag{A.20}$$

In $\tilde{\boldsymbol{D}}$, the first row enforces the asset budget constraint $\boldsymbol{x}_t^\top\mathbf{1}_N = 1$, and the remaining $N+M$ rows enforce the weight decomposition (A.11) for all components of $\boldsymbol{\theta}_t$. In $\tilde{\boldsymbol{A}}$, the first q rows encode the linearization inequalities (A.7) applied to the joint return matrix $\boldsymbol{R}^{\text{all}}$. The final row imposes the asset-only leverage constraint (A.12), giving $\mathbf{1}_N^\top(\boldsymbol{\Delta}_{x,t}^+ + \boldsymbol{\Delta}_{x,t}^-) = \|\boldsymbol{x}_t\|_1 \leq \ell^a$.

A.5.2 Joint Optimization under Second-Order Expected Shortfall

For second-order Expected Shortfall, we introduce, as in Section A.2.2, one more auxiliary variable s subject to the cone constraint $\|\boldsymbol{u}\|_2 \leq s$, which extends the decision vector to

$$\tilde{\boldsymbol{\theta}}_t = \Big[\, \boldsymbol{x}_t^\top, \;\; \boldsymbol{\psi}_t^\top, \;\; \alpha, \;\; \boldsymbol{u}^\top, \;\; (\boldsymbol{\Delta}_{x,t}^-)^\top, \;\; (\boldsymbol{\Delta}_{\psi,t}^-)^\top, \;\; (\boldsymbol{\Delta}_{x,t}^+)^\top, \;\; (\boldsymbol{\Delta}_{\psi,t}^+)^\top, \;\; s \,\Big]^\top \in \mathbb{R}^{3N+3M+q+2},$$

with the cost vector

$$\tilde{\boldsymbol{c}} = \Big[\, \mathbf{0}_{N+M}^\top, \;\; 1, \;\; \mathbf{0}_q^\top, \;\; \lambda_1^a \mathbf{1}_N^\top, \;\; \lambda_1^c \mathbf{1}_M^\top, \;\; \lambda_1^a \mathbf{1}_N^\top, \;\; \lambda_1^c \mathbf{1}_M^\top, \;\; \tfrac{1}{(1-\beta)\sqrt{q}} \,\Big]^\top$$

and the penalty matrix

$$\tilde{\boldsymbol{Q}} = \begin{bmatrix} \lambda_2^a \boldsymbol{I}_N & \mathbf{0}_{N\times M} & \mathbf{0}_{N\times(2N+2M+q+2)} \\ \mathbf{0}_{M\times N} & \lambda_2^c \boldsymbol{I}_M & \mathbf{0}_{M\times(2N+2M+q+2)} \\ \mathbf{0}_{(2N+2M+q+2)\times N} & \mathbf{0}_{(2N+2M+q+2)\times M} & \mathbf{0}_{(2N+2M+q+2)\times(2N+2M+q+2)} \end{bmatrix}.$$

The joint portfolio optimization problem under second-order Expected Shortfall then becomes a second-order cone program with a convex quadratic objective

$$\begin{aligned} \tilde{\boldsymbol{\theta}}_t^* &= \arg\min_{\tilde{\boldsymbol{\theta}}_t} \quad \tilde{\boldsymbol{c}}^\top \tilde{\boldsymbol{\theta}}_t + \tilde{\boldsymbol{\theta}}_t^\top \tilde{\boldsymbol{Q}}\, \tilde{\boldsymbol{\theta}}_t, \\ \text{s.t.} \quad \tilde{\boldsymbol{D}}\tilde{\boldsymbol{\theta}}_t &= \begin{bmatrix} \boldsymbol{q}_{N,M}^\top & \mathbf{0}_{q+1}^\top & \mathbf{0}_{N+M}^\top & \mathbf{0}_{N+M}^\top & 0 \\ \boldsymbol{I}_{N+M} & \mathbf{0}_{(N+M)\times(q+1)} & \boldsymbol{I}_{N+M} & -\boldsymbol{I}_{N+M} & \mathbf{0}_{N+M} \end{bmatrix} \tilde{\boldsymbol{\theta}}_t = \tilde{\boldsymbol{E}} = \begin{bmatrix} 1 \\ \mathbf{0}_{N+M} \end{bmatrix}, \\ \tilde{\boldsymbol{A}}\tilde{\boldsymbol{\theta}}_t &= \begin{bmatrix} -\boldsymbol{R}^{\text{all}} & -\mathbf{1}_q & -\boldsymbol{I}_q & \mathbf{0}_{q\times(N+M)} & \mathbf{0}_{q\times(N+M)} & \mathbf{0}_q \\ \mathbf{0}_{N+M}^\top & 0 & \mathbf{0}_q^\top & \boldsymbol{q}_{N,M}^\top & \boldsymbol{q}_{N,M}^\top & 0 \end{bmatrix} \tilde{\boldsymbol{\theta}}_t \leq \boldsymbol{b} = \begin{bmatrix} \mathbf{0}_q \\ \ell^a \end{bmatrix}, \qquad \text{(A.21)} \\ &\|\boldsymbol{u}\|_2 \leq s, \\ &\qquad \alpha \geq 0, \quad u_k \geq 0 \quad \text{for } k = 1, \ldots, q, \\ &\qquad \Delta_{x,i,t}^{\pm} \geq 0, \quad \Delta_{\psi,c,t}^{\pm} \geq 0 \quad \text{for } i = 1, \ldots, N,\ c = 1, \ldots, M. \end{aligned}$$

Here $\tilde{\boldsymbol{D}}$ and $\tilde{\boldsymbol{A}}$ are the constraint matrices of (A.20) with a zero column added for the variable s.

Appendix B

Optimization Derivation: Variance and Semi-Variance

In this appendix, we derive the optimization problems used in the main text for mean-variance portfolio optimization. We start with the standard mean-variance formulation and then introduce regularization techniques and estimators to stabilize the optimization. We then extend the formulation to the multi-currency setting, considering separate optimization for assets and currency overlay, and finally the joint optimization problem. In Section B.1, we follow [31, 18] and present the quadratic program representation of the mean-variance portfolio optimization problem.

B.1 Quadratic Program Representation of the Mean-Variance Portfolio Optimization

The mean-variance portfolio optimization of [18] trades off the expected portfolio return against its return variance. The portfolio return $\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1}$ is linear in the asset weights, so its variance is the quadratic form $\mathrm{Var}(\boldsymbol{x}_t^\top \boldsymbol{R}_{t+1}) = \boldsymbol{x}_t^\top \boldsymbol{\Sigma}_t \boldsymbol{x}_t$, where $\boldsymbol{R}_{t+1} \in \mathbb{R}^N$ is the asset return vector and $\boldsymbol{\Sigma}_t$ its covariance matrix. Following the international mean-variance formulation of [31], we write the portfolio optimization problem as a minimization that trades off expected return against variance through a relative risk-aversion coefficient $\gamma > 0$

$$\boldsymbol{x}_t^* = \arg\min_{\boldsymbol{x}_t} \left\{ -\boldsymbol{\mu}_t^\top \boldsymbol{x}_t + \frac{\gamma}{2} \boldsymbol{x}_t^\top \boldsymbol{\Sigma}_t \boldsymbol{x}_t \right\}, \quad \text{s.t. } \boldsymbol{x}_t^\top \mathbf{1}_N = 1, \tag{B.1}$$

where $\boldsymbol{\mu}_t$ denotes the asset return mean estimator, $\boldsymbol{\Sigma}_t$ the asset return covariance estimator, and the budget constraint $\boldsymbol{x}_t^\top \mathbf{1}_N = 1$ enforces full investment.

In contrast to the Expected Shortfall objective (**??**), which depends on the unknown conditional expectation $\mathbb{E}_t[\cdot]$ and requires the auxiliary representation of Section A.1 together with a sample approximation to obtain a tractable form, the variance term in (B.1) is already expressed in closed quadratic form. Problem (B.1) is therefore already formulated as a quadratic program.

The objective in (B.1) is convex. The first term $-\boldsymbol{\mu}_t^\top \boldsymbol{x}_t$ is linear, hence convex, and the quadratic term $\frac{\gamma}{2} \boldsymbol{x}_t^\top \boldsymbol{\Sigma}_t \boldsymbol{x}_t$ has Hessian $\gamma \boldsymbol{\Sigma}_t$. Since a covariance matrix is positive semi-definite and $\gamma > 0$, this Hessian is positive semi-definite, so the quadratic term, and with it the full objective, is convex [2]. Any local minimum is therefore global, ensuring the problem can be solved reliably with standard convex optimization methods.

The mean and covariance terms are unknown and replaced by their sample estimators $\boldsymbol{\mu}_t$ and $\boldsymbol{\Sigma}_t$. Because the variance term is expressed in closed form, this step introduces no auxiliary

variables and leaves the decision vector at dimension N, independent of the number of historical observations q used to estimate it. This is different from the empirical approximation of the Expected Shortfall objective in Section A.2, where the conditional expectation is replaced by an average over q historical return samples and q slack variables u_k enlarge the problem to dimension $N + q + 1$. The sample covariance estimator we use is positive semi-definite, so the convexity established above is preserved under estimation. The choice of estimators is discussed in Section B.2.2.

Collecting the asset weights in the decision vector $\boldsymbol{x}_t \in \mathbb{R}^N$ and defining the cost vector and the penalty matrix

$$\boldsymbol{c} = -\boldsymbol{\mu}_t \in \mathbb{R}^N, \qquad \boldsymbol{Q} = \gamma \boldsymbol{\Sigma}_t \in \mathbb{R}^{N \times N},$$

the mean-variance problem (B.1) is written as the quadratic program

$$\begin{aligned} \boldsymbol{x}_t^* = \arg\min_{\boldsymbol{x}_t} \quad & \boldsymbol{c}^\top \boldsymbol{x}_t + \tfrac{1}{2}\, \boldsymbol{x}_t^\top \boldsymbol{Q}\, \boldsymbol{x}_t, \\ \text{s.t.} \quad & \boldsymbol{D}\boldsymbol{x}_t = \mathbf{1}_N^\top \boldsymbol{x}_t = E = 1. \end{aligned} \tag{B.2}$$

Here $\boldsymbol{D} = \mathbf{1}_N^\top \in \mathbb{R}^{1 \times N}$ encodes the budget constraint with right-hand side $E = 1$.

Problem (B.2) has a convex quadratic objective subject to a single linear equality constraint, and is thus a quadratic program solvable with standard methods. This formulation serves as the base unregularized portfolio optimization problem under the mean-variance criterion. Regularization terms and additional constraints are added in Section B.2.

B.2 Regularization

The base problem (B.2) is well-documented to suffer from instability. The weights are sensitive to estimation error in $\boldsymbol{\mu}_t$ and $\boldsymbol{\Sigma}_t$ and tend to take extreme, concentrated values, so that portfolios formed from the raw sample moments perform poorly out of sample [22, 31]. We stabilize the problem with two complementary approaches. First, we regularize the weights directly through an L_1-norm penalty that promotes sparsity, an L_2-norm penalty that promotes shrinkage, and a leverage constraint that bounds gross exposure. These are the same three techniques appended to the Expected Shortfall optimizationprogram in Section A.3. Second, we shrink the input estimators $\boldsymbol{\mu}_t$ and $\boldsymbol{\Sigma}_t$ themselves toward more stable targets. Both preserve the convex quadratic structure of (B.2).

B.2.1 L_1, L_2 Regularization and Leverage Constraint

The L_1 penalty $\lambda_1 \|\boldsymbol{x}_t\|_1 = \lambda_1 \sum_{i=1}^N |x_{i,t}|$, with $\lambda_1 \geq 0$, is linearized just like in Appendix A.3.1. Each weight is split into its positive and negative parts through the auxiliary variables $\Delta_{i,t}^+ \geq 0$ and $\Delta_{i,t}^- \geq 0$, giving

$$\lambda_1 \|\boldsymbol{x}_t\|_1 = \lambda_1 \mathbf{1}_N^\top (\boldsymbol{\Delta}_t^+ + \boldsymbol{\Delta}_t^-).$$

The leverage constraint $\|\boldsymbol{x}_t\|_1 \leq \ell$, with $\ell \geq 1$, reuses the same decomposition and is the single linear inequality

$$\|\boldsymbol{x}_t\|_1 = \mathbf{1}_N^\top (\boldsymbol{\Delta}_t^+ + \boldsymbol{\Delta}_t^-) \leq \ell.$$

As before, the choice $\ell = 1$ rules out leverage entirely and is equivalent to a hard L_1 constraint.

The L_2 penalty $\lambda_2 \|\boldsymbol{x}_t\|_2^2 = \lambda_2 \boldsymbol{x}_t^\top \boldsymbol{x}_t$, with $\lambda_2 \geq 0$, is the quadratic form $\frac{1}{2}\boldsymbol{x}_t^\top (2\lambda_2 \boldsymbol{I}_N)\boldsymbol{x}_t$, as in Appendix A.3.2. Here the mean-variance and Expected Shortfall cases differ. In the Expected Shortfall program, the L_2 term is what turns the linear program (A.8) into a quadratic program. The mean-variance objective (B.1) is already quadratic through the variance term,

so the L_2 penalty does not change the program class. Instead, it adds to the existing quadratic form and combines with the covariance into the single term

$$\tfrac{\gamma}{2}\,\boldsymbol{x}_t^\top \boldsymbol{\Sigma}_t \boldsymbol{x}_t + \lambda_2 \boldsymbol{x}_t^\top \boldsymbol{x}_t = \tfrac{1}{2}\,\boldsymbol{x}_t^\top \left(\gamma \boldsymbol{\Sigma}_t + 2\lambda_2 \boldsymbol{I}_N\right) \boldsymbol{x}_t.$$

Because $\gamma\boldsymbol{\Sigma}_t$ is positive semi-definite and $2\lambda_2\boldsymbol{I}_N$ is positive definite for $\lambda_2 > 0$, the combined matrix is positive definite, so the penalized objective is strictly convex and admits a unique minimizer even when the covariance estimator $\boldsymbol{\Sigma}_t$ is singular, which occurs when the number of assets exceeds the number of observations. The combined term also makes the stabilizing role of the penalty explicit, since $\gamma\boldsymbol{\Sigma}_t + 2\lambda_2\boldsymbol{I}_N = \gamma(\boldsymbol{\Sigma}_t + \frac{2\lambda_2}{\gamma}\boldsymbol{I}_N)$ shrinks the covariance toward a scaled identity. The L_2 penalty can therefore be interpreted as a covariance shrinkage toward the identity [16], complementing the estimator shrinkage introduced in Section B.2.2.

B.2.2 Estimators

The program (B.2) depends on the unknown mean $\boldsymbol{\mu}_t$ and covariance $\boldsymbol{\Sigma}_t$. Replacing them by the sample mean $\bar{\boldsymbol{\mu}}_t$ and the sample covariance $\boldsymbol{S}_t$ is the approach, which is known to produce unstable and poorly diversified portfolios, as the optimization amplifies estimation error in these inputs [22, 14]. We instead replace each input by a shrinkage estimator, $\boldsymbol{\mu}_t = \boldsymbol{\mu}_t^{BS}$ and $\boldsymbol{\Sigma}_t = \boldsymbol{\Sigma}_t^{LW}$, following [31].

For estimating the mean, we use the Bayes–Stein estimator introduced in [11], which shrinks the sample mean towards a fixed shrinkage target. The shrinkage target is the return of the minimum-variance portfolio

$$\mu_0 = \frac{\mathbf{1}_N^\top \boldsymbol{S}_t^{-1} \bar{\boldsymbol{\mu}}_t}{\mathbf{1}_N^\top \boldsymbol{S}_t^{-1} \mathbf{1}_N},$$

and the Bayes–Stein estimator is the convex combination

$$\boldsymbol{\mu}_t^{BS} = (1-\hat{\omega})\,\bar{\boldsymbol{\mu}}_t + \hat{\omega}\,\mu_0\,\mathbf{1}_N, \qquad \hat{\omega} = \frac{N+2}{(N+2) + q\,(\bar{\boldsymbol{\mu}}_t - \mu_0 \mathbf{1}_N)^\top \boldsymbol{S}_t^{-1} (\bar{\boldsymbol{\mu}}_t - \mu_0 \mathbf{1}_N)},$$

where q is the number of observations in the estimation window, N the number of assets, and $\hat{\omega} \in [0,1]$ the shrinkage intensity. The intensity increases toward one as the sample size gets smaller, or as the variance of the sample means around μ_0 decreases. This pulls the estimator toward the shrinkage target when the data is least informative. The derivation of the estimator is given in [11].

For estimating the covariance matrix, we use the Ledoit–Wolf linear shrinkage estimator introduced in [15], which shrinks the sample covariance toward a shrinkage target. The Ledoit–Wolf estimator is the convex combination

$$\boldsymbol{\Sigma}_t^{LW} = \hat{v}\,\boldsymbol{F}_t + (1-\hat{v})\,\boldsymbol{S}_t,$$

with shrinkage intensity $\hat{v} \in [0,1]$. The shrinkage target $\boldsymbol{F}_t$ is the constant-correlation matrix, which retains the sample variances and replaces every pairwise correlation by the average sample correlation

$$F_{ii,t} = S_{ii,t}, \qquad F_{ij,t} = \bar{\rho}_t \sqrt{S_{ii,t}\,S_{jj,t}} \;\; (i \neq j), \qquad \bar{\rho}_t = \frac{2}{N(N-1)} \sum_{i<j} \frac{S_{ij,t}}{\sqrt{S_{ii,t}\,S_{jj,t}}}.$$

The intensity $\hat{v}$ is chosen to minimize the expected Frobenius distance between $\boldsymbol{\Sigma}_t^{LW}$ and the true covariance matrix, its closed form is derived in [15]. Both $\boldsymbol{F}_t$ and $\boldsymbol{S}_t$ are positive semi-definite, so $\boldsymbol{\Sigma}_t^{LW}$ is positive semi-definite and the convexity established in Section B.1 is preserved under estimation.

The estimators above are stated for a single asset block. In the multi-currency setting, the asset and currency means and covariances, and in the separate approach the covariance of the augmented currency-overlay problem, are constructed analogously, we defer their specification to Section B.3 and Section B.4.

B.2.3 Quadratic Program Representation of the Regularized Portfolio Optimization

Adding the three weight-level regularizers to the base problem (B.1) gives the regularized mean-variance problem

$$\begin{aligned} \boldsymbol{x}_t^* = \arg\min_{\boldsymbol{x}_t} \left\{ -\boldsymbol{\mu}_t^{BS\top} \boldsymbol{x}_t + \tfrac{\gamma}{2}\, \boldsymbol{x}_t^\top \boldsymbol{\Sigma}_t^{LW} \boldsymbol{x}_t + \lambda_1 \|\boldsymbol{x}_t\|_1 + \lambda_2 \|\boldsymbol{x}_t\|_2^2 \right\}, \\ \text{s.t. } \boldsymbol{x}_t^\top \mathbf{1}_N = 1,\ \|\boldsymbol{x}_t\|_1 \le \ell, \end{aligned} \tag{B.3}$$

Following the L_1 decomposition, we extend the decision vector to carry the positive and negative parts of the weights

$$\tilde{\boldsymbol{x}}_t = \left[\boldsymbol{x}_t^\top,\ \ (\boldsymbol{\Delta}_t^-)^\top,\ \ (\boldsymbol{\Delta}_t^+)^\top \right]^\top \in \mathbb{R}^{3N},$$

and define the cost vector and the penalty matrix

$$\tilde{\boldsymbol{c}} = \left[-\boldsymbol{\mu}_t^{BS\top},\ \ \lambda_1 \mathbf{1}_N^\top,\ \ \lambda_1 \mathbf{1}_N^\top \right]^\top, \qquad \tilde{\boldsymbol{Q}} = \begin{bmatrix} \gamma \boldsymbol{\Sigma}_t^{LW} + 2\lambda_2 \boldsymbol{I}_N & \mathbf{0}_{N \times 2N} \\ \mathbf{0}_{2N \times N} & \mathbf{0}_{2N \times 2N} \end{bmatrix}.$$

The regularized mean-variance portfolio optimization problem is then the quadratic program

$$\begin{aligned} \tilde{\boldsymbol{x}}_t^* &= \arg\min_{\tilde{\boldsymbol{x}}_t} \quad \tilde{\boldsymbol{c}}^\top \tilde{\boldsymbol{x}}_t + \tfrac{1}{2}\, \tilde{\boldsymbol{x}}_t^\top \tilde{\boldsymbol{Q}}\, \tilde{\boldsymbol{x}}_t, \\ \text{s.t.} \quad \tilde{\boldsymbol{D}} \tilde{\boldsymbol{x}}_t &= \begin{bmatrix} \mathbf{1}_N^\top & \mathbf{0}_N^\top & \mathbf{0}_N^\top \\ \boldsymbol{I}_N & \boldsymbol{I}_N & -\boldsymbol{I}_N \end{bmatrix} \tilde{\boldsymbol{x}}_t = \tilde{\boldsymbol{E}} = \begin{bmatrix} 1 \\ \mathbf{0}_N \end{bmatrix}, \\ \tilde{\boldsymbol{A}} \tilde{\boldsymbol{x}}_t &= \begin{bmatrix} \mathbf{0}_N^\top & \mathbf{1}_N^\top & \mathbf{1}_N^\top \end{bmatrix} \tilde{\boldsymbol{x}}_t \le b = \ell, \\ \Delta_{i,t}^- \ge 0, &\quad \Delta_{i,t}^+ \ge 0 \quad \text{for } i = 1, \dots, N. \end{aligned} \tag{B.4}$$

The first row of $\tilde{\boldsymbol{D}}$ enforces the budget constraint and the remaining N rows enforce the weight decomposition (A.11). The single row of $\tilde{\boldsymbol{A}}$ imposes the leverage constraint (A.12), giving $\mathbf{1}_N^\top(\boldsymbol{\Delta}_t^+ + \boldsymbol{\Delta}_t^-) = \|\boldsymbol{x}_t\|_1 \le \ell$.

Compared with the regularized Expected Shortfall program (A.13), the mean-variance program is smaller. Its decision vector has dimension $3N$ rather than $3N + q + 1$. This formulation serves as the regularized base problem under the mean-variance criterion. The asset and currency extensions under the separate and joint approaches are derived in Section B.3 and Section B.4.

B.3 Separate (Currency Overlay) Optimization

In the separate (currency overlay) portfolio optimization approach, the asset and currency allocations are optimized in two sequential steps. First, the asset weights $\boldsymbol{x}_t \in \mathbb{R}^N$ are optimized on the fully hedged returns, then, conditional on the optimal $\boldsymbol{x}_t^*$, the currency exposures $\boldsymbol{\psi}_t \in \mathbb{R}^M$ are optimized. Let

$$\begin{aligned} \boldsymbol{\mu}^{fh} &= \mathbb{E}_t[\boldsymbol{R}_{t+1}^{fh}] \in \mathbb{R}^N, & \boldsymbol{\Sigma}^a &= \mathrm{Var}_t(\boldsymbol{R}_{t+1}^{fh}) \in \mathbb{R}^{N \times N}, \\ \boldsymbol{\mu}^c &= \mathbb{E}_t[\boldsymbol{e}_{t+1} - \boldsymbol{f}_t] \in \mathbb{R}^M, & \boldsymbol{\Sigma}^c &= \mathrm{Var}_t(\boldsymbol{e}_{t+1} - \boldsymbol{f}_t) \in \mathbb{R}^{M \times M}, \end{aligned}$$

denote the expected fully hedged asset return and excess currency return together with their covariance matrices, and let $\boldsymbol{\Sigma}^{ac} = \mathrm{Cov}_t(\boldsymbol{R}_{t+1}^{fh},\, \boldsymbol{e}_{t+1} - \boldsymbol{f}_t) \in \mathbb{R}^{N\times M}$ denote the cross-covariance between fully hedged asset returns and excess currency returns. Each moment is the corresponding shrinkage estimator of Section B.2.2. The two optimization steps are then given by

$$\begin{aligned} \boldsymbol{x}_t^* = \arg\min_{\boldsymbol{x}_t} \Big\{ -\boldsymbol{x}_t^\top \boldsymbol{\mu}^{fh} + \tfrac{\gamma}{2}\, \boldsymbol{x}_t^\top \boldsymbol{\Sigma}^a \boldsymbol{x}_t + \lambda_1^a \|\boldsymbol{x}_t\|_1 + \lambda_2^a \|\boldsymbol{x}_t\|_2^2 \Big\} \\ \text{s.t. } \boldsymbol{x}_t^\top \mathbf{1}_N = 1,\ \|\boldsymbol{x}_t\|_1 \le \ell, \end{aligned} \tag{B.5}$$

$$\begin{aligned} \boldsymbol{\psi}_t^* \mid \boldsymbol{x}_t^* = \arg\min_{\boldsymbol{\psi}_t} \Big\{ -\boldsymbol{\psi}_t^\top \boldsymbol{\mu}^c + \tfrac{\gamma}{2}\, \mathrm{Var}_t\Big(\boldsymbol{x}_t^{*\top} \boldsymbol{R}_{t+1}^{fh} + \boldsymbol{\psi}_t^\top (\boldsymbol{e}_{t+1} - \boldsymbol{f}_t) \Big) \\ + \lambda_1^c \|\boldsymbol{\psi}_t\|_1 + \lambda_2^c \|\boldsymbol{\psi}_t\|_2^2 \Big\}. \end{aligned} \tag{B.6}$$

The regularization parameters are kept separate across assets $(\lambda_1^a, \lambda_2^a)$ and currencies $(\lambda_1^c, \lambda_2^c)$.

B.3.1 Asset Optimization

The asset optimization step (B.5) is identical to the regularized single-currency program (B.4). It carries a budget constraint, L_1 and L_2 penalty terms, and a leverage constraint. It is therefore obtained directly from (B.4) by substituting the fully hedged asset moments $\boldsymbol{\mu}^{fh}$ for $\boldsymbol{\mu}_t$ and $\boldsymbol{\Sigma}^a$ for $\boldsymbol{\Sigma}_t$, together with the asset-specific penalties λ_1^a, λ_2^a. With the augmented vector

$$\tilde{\boldsymbol{x}}_t = \Big[\boldsymbol{x}_t^\top,\ \ (\boldsymbol{\Delta}_{x,t}^-)^\top,\ \ (\boldsymbol{\Delta}_{x,t}^+)^\top \Big]^\top \in \mathbb{R}^{3N},$$

the cost vector and penalty matrix are

$$\tilde{\boldsymbol{c}} = \Big[-(\boldsymbol{\mu}^{fh})^\top,\ \ \lambda_1^a \mathbf{1}_N^\top,\ \ \lambda_1^a \mathbf{1}_N^\top \Big]^\top, \qquad \tilde{\boldsymbol{Q}} = \begin{bmatrix} \gamma \boldsymbol{\Sigma}^a + 2\lambda_2^a \boldsymbol{I}_N & \mathbf{0}_{N\times 2N} \\ \mathbf{0}_{2N\times N} & \mathbf{0}_{2N\times 2N} \end{bmatrix},$$

and the equality, inequality, and sign constraints are exactly those of (B.4) with $\boldsymbol{\Sigma}^a$ and $\boldsymbol{\mu}^{fh}$ in place of $\boldsymbol{\Sigma}_t$ and $\boldsymbol{\mu}_t$. The step yields the optimal asset weights $\boldsymbol{x}_t^*$, which are then fixed in the currency overlay optimization.

B.3.2 Currency Overlay Optimization

Conditional on $\boldsymbol{x}_t^*$, the currency step (B.6) optimizes the mean-variance objective of the combined hedged return $\boldsymbol{x}_t^{*\top} \boldsymbol{R}_{t+1}^{fh} + \boldsymbol{\psi}_t^\top (\boldsymbol{e}_{t+1} - \boldsymbol{f}_t)$ over the currency exposures $\boldsymbol{\psi}_t$. Unlike the Expected Shortfall in the currency overlay of **??**, the variance decomposes in closed form

$$\mathrm{Var}_t\Big(\boldsymbol{x}_t^{*\top} \boldsymbol{R}_{t+1}^{fh} + \boldsymbol{\psi}_t^\top (\boldsymbol{e}_{t+1} - \boldsymbol{f}_t) \Big) = \boldsymbol{x}_t^{*\top} \boldsymbol{\Sigma}^a \boldsymbol{x}_t^* + 2\, \boldsymbol{x}_t^{*\top} \boldsymbol{\Sigma}^{ac} \boldsymbol{\psi}_t + \boldsymbol{\psi}_t^\top \boldsymbol{\Sigma}^c \boldsymbol{\psi}_t.$$

The term $\boldsymbol{x}_t^{*\top} \boldsymbol{\Sigma}^a \boldsymbol{x}_t^*$ does not depend on $\boldsymbol{\psi}_t$ and is therefore dropped from the minimization. Substituting the remaining terms into (B.6), the currency overlay objective becomes

$$\boldsymbol{\psi}_t^* \mid \boldsymbol{x}_t^* = \arg\min_{\boldsymbol{\psi}_t} \Big\{ -\boldsymbol{\psi}_t^\top \boldsymbol{\mu}^c + \tfrac{\gamma}{2}\, \boldsymbol{\psi}_t^\top \boldsymbol{\Sigma}^c \boldsymbol{\psi}_t + \gamma\, \boldsymbol{x}_t^{*\top} \boldsymbol{\Sigma}^{ac} \boldsymbol{\psi}_t + \lambda_1^c \|\boldsymbol{\psi}_t\|_1 + \lambda_2^c \|\boldsymbol{\psi}_t\|_2^2 \Big\}. \tag{B.7}$$

The fixed asset weights are present through the linear cross-covariance term $\gamma\, \boldsymbol{x}_t^{*\top} \boldsymbol{\Sigma}^{ac} \boldsymbol{\psi}_t$. Collecting the two linear terms shows that the coupling is equivalent to shifting the currency mean, $-\boldsymbol{\psi}_t^\top \boldsymbol{\mu}^c + \gamma\, \boldsymbol{x}_t^{*\top} \boldsymbol{\Sigma}^{ac} \boldsymbol{\psi}_t = -\boldsymbol{\psi}_t^\top (\boldsymbol{\mu}^c - \gamma\, (\boldsymbol{\Sigma}^{ac})^\top \boldsymbol{x}_t^*)$. The fixed asset portfolio therefore lowers the effective expected currency return by $\gamma\, (\boldsymbol{\Sigma}^{ac})^\top \boldsymbol{x}_t^*$. This effect is similar to the

one in **??**, where the fixed asset weights shift the right-hand side of the slack constraints to $\boldsymbol{R}_P^{fh} = \boldsymbol{R}^{fh}\boldsymbol{x}_t^*$.

We linearize the currency L_1 penalty with the decomposition $\psi_{c,t} = \Delta^+_{\psi,c,t} - \Delta^-_{\psi,c,t}$ as in Section A.3.1, and form the augmented vector

$$\tilde{\boldsymbol{\psi}}_t = \left[\boldsymbol{\psi}_t^\top, \quad (\boldsymbol{\Delta}^-_{\psi,t})^\top, \quad (\boldsymbol{\Delta}^+_{\psi,t})^\top\right]^\top \in \mathbb{R}^{3M}.$$

The cost vector and penalty matrix are

$$\tilde{\boldsymbol{c}} = \left[-(\boldsymbol{\mu}^c)^\top + \gamma\, \boldsymbol{x}_t^{*\top}\boldsymbol{\Sigma}^{ac}, \quad \lambda_1^c \mathbf{1}_M^\top, \quad \lambda_1^c \mathbf{1}_M^\top\right]^\top, \qquad \tilde{\boldsymbol{Q}} = \begin{bmatrix} \gamma\boldsymbol{\Sigma}^c + 2\lambda_2^c \boldsymbol{I}_M & \mathbf{0}_{M\times 2M} \\ \mathbf{0}_{2M\times M} & \mathbf{0}_{2M\times 2M} \end{bmatrix}.$$

The currency overlay optimization is then the quadratic program

$$\begin{aligned} \tilde{\boldsymbol{\psi}}_t^* = \arg\min_{\tilde{\boldsymbol{\psi}}_t} \quad & \tilde{\boldsymbol{c}}^\top \tilde{\boldsymbol{\psi}}_t + \tfrac{1}{2}\, \tilde{\boldsymbol{\psi}}_t^\top \tilde{\boldsymbol{Q}}\, \tilde{\boldsymbol{\psi}}_t, \\ \text{s.t.} \quad & \tilde{\boldsymbol{D}}\tilde{\boldsymbol{\psi}}_t = \begin{bmatrix} \boldsymbol{I}_M & \boldsymbol{I}_M & -\boldsymbol{I}_M \end{bmatrix} \tilde{\boldsymbol{\psi}}_t = \tilde{\boldsymbol{E}} = \mathbf{0}_M, \\ & \Delta^-_{\psi,c,t} \geq 0, \quad \Delta^+_{\psi,c,t} \geq 0 \quad \text{for } c = 1, \ldots, M. \end{aligned} \tag{B.8}$$

The currency optimization differs from the asset optimization (B.5) in the same ways as in the Expected Shortfall case. First, there is no budget constraint, as the currency portfolio is a zero-investment portfolio, so the domestic exposure $\psi_{0,t}$ is inferred residually and the equality block $\tilde{\boldsymbol{D}}$ enforces only the weight decomposition. Second, no leverage constraint is imposed on the currency exposures. The biggest difference relative to **??** is that the fixed asset weights enter the mean-variance currency step through the linear cross-covariance term in $\tilde{\boldsymbol{c}}$, rather than through a shifted inequality constraint.

As noted by [31], the cross-covariance could alternatively be shrunk by augmenting the currency overlay problem with the fixed asset portfolio as an additional dimension. We use the direct formulation over $\boldsymbol{\psi}_t$, which keeps the currency step as both the asset program (B.4) and the Expected Shortfall currency overlay (A.17).

B.4 Joint Optimization

In the joint portfolio optimization approach, asset and currency weights are optimized simultaneously rather than sequentially. We stack them into the combined decision vector $\boldsymbol{\theta}_t = [\boldsymbol{x}_t^\top, \boldsymbol{\psi}_t^\top]^\top \in \mathbb{R}^{N+M}$ and the combined return vector $\boldsymbol{R}_{t+1} = [\boldsymbol{R}_{t+1}^{fh\top}, (\boldsymbol{e}_{t+1} - \boldsymbol{f}_t)^\top]^\top \in \mathbb{R}^{N+M}$, with joint mean $\boldsymbol{\mu} = \mathbb{E}_t[\boldsymbol{R}_{t+1}] \in \mathbb{R}^{N+M}$ and joint covariance $\boldsymbol{\Sigma} = \mathrm{Var}_t(\boldsymbol{R}_{t+1}) \in \mathbb{R}^{(N+M)\times(N+M)}$. The joint optimization problem is

$$\begin{aligned} \boldsymbol{\theta}_t^* = \arg\min_{\boldsymbol{\theta}_t} \Big\{ & -\boldsymbol{\theta}_t^\top \boldsymbol{\mu} + \tfrac{\gamma}{2}\, \boldsymbol{\theta}_t^\top \boldsymbol{\Sigma}\, \boldsymbol{\theta}_t + \|\boldsymbol{S}_1\boldsymbol{\theta}_t\|_1 + \|\boldsymbol{S}_2\boldsymbol{\theta}_t\|_2^2 \Big\}, \\ & \text{s.t. } \boldsymbol{\theta}_t^\top \boldsymbol{q}_{N,M} = 1, \quad \|\boldsymbol{L}\boldsymbol{\theta}_t\|_1 \leq \ell. \end{aligned} \tag{B.9}$$

The following matrices and vector adapt the problem to the joint optimization setting

$$\boldsymbol{q}_{N,M} = \begin{bmatrix} \mathbf{1}_N \\ \mathbf{0}_M \end{bmatrix}, \quad \boldsymbol{L} = \begin{bmatrix} \boldsymbol{I}_N & \mathbf{0}_{N\times M} \\ \mathbf{0}_{M\times N} & \mathbf{0}_{M\times M} \end{bmatrix}, \quad \boldsymbol{S}_1 = \begin{bmatrix} \lambda_1^a \boldsymbol{I}_N & 0 \\ 0 & \lambda_1^c \boldsymbol{I}_M \end{bmatrix}, \quad \boldsymbol{S}_2 = \begin{bmatrix} \sqrt{\lambda_2^a}\, \boldsymbol{I}_N & 0 \\ 0 & \sqrt{\lambda_2^c}\, \boldsymbol{I}_M \end{bmatrix},$$

like in the Expected Shortfall problem (A.19). The vector $\boldsymbol{q}_{N,M}$ applies the budget constraint to the asset weights only, so that $\boldsymbol{\theta}_t^\top \boldsymbol{q}_{N,M} = \boldsymbol{x}_t^\top \mathbf{1}_N = 1$, leaving the currency weights as a

zero-investment overlay. The matrix $\boldsymbol{L}$ applies the leverage constraint to the asset weights only, giving $\|\boldsymbol{L}\boldsymbol{\theta}_t\|_1 = \|\boldsymbol{x}_t\|_1$. The diagonal regularization matrices $\boldsymbol{S}_1$ and $\boldsymbol{S}_2$ allow the asset and currency penalties to differ within a single objective,

$$\|\boldsymbol{S}_1\boldsymbol{\theta}_t\|_1 = \lambda_1^a\|\boldsymbol{x}_t\|_1 + \lambda_1^c\|\boldsymbol{\psi}_t\|_1, \qquad \|\boldsymbol{S}_2\boldsymbol{\theta}_t\|_2^2 = \lambda_2^a\|\boldsymbol{x}_t\|_2^2 + \lambda_2^c\|\boldsymbol{\psi}_t\|_2^2.$$

The joint covariance is

$$\boldsymbol{\Sigma} = \begin{bmatrix} \boldsymbol{\Sigma}^a & \boldsymbol{\Sigma}^{ac} \\ (\boldsymbol{\Sigma}^{ac})^\top & \boldsymbol{\Sigma}^c \end{bmatrix},$$

with the asset, currency, and cross-covariance matrices $\boldsymbol{\Sigma}^a$, $\boldsymbol{\Sigma}^c$, and $\boldsymbol{\Sigma}^{ac}$ defined as in Section B.3. The portfolio variance then expands as

$$\boldsymbol{\theta}_t^\top \boldsymbol{\Sigma}\, \boldsymbol{\theta}_t = \boldsymbol{x}_t^\top \boldsymbol{\Sigma}^a \boldsymbol{x}_t + 2\, \boldsymbol{x}_t^\top \boldsymbol{\Sigma}^{ac} \boldsymbol{\psi}_t + \boldsymbol{\psi}_t^\top \boldsymbol{\Sigma}^c \boldsymbol{\psi}_t. \tag{B.10}$$

This is the same decomposition that appears in the currency overlay step (B.7). There, however, the asset weights are fixed at $\boldsymbol{x}_t^*$, so the first term is constant and drops from the objective, while the cross term simplifies to the linear form $\gamma\, \boldsymbol{x}_t^{*\top} \boldsymbol{\Sigma}^{ac} \boldsymbol{\psi}_t$. In the joint problem both $\boldsymbol{x}_t$ and $\boldsymbol{\psi}_t$ are free. This is how thejoint approach captures the asset–currency dependence.

The joint mean $\boldsymbol{\mu}$ and covariance $\boldsymbol{\Sigma}$ are the Bayes–Stein and Ledoit–Wolf estimators of Section B.2.2, here applied to the combined return vector $\boldsymbol{R}_{t+1}$. Because the Ledoit–Wolf shrinkage acts on the full combined covariance, the cross-covariance matrix $\boldsymbol{\Sigma}^{ac}$ is shrunk together with the asset and currency matrices. This is in contrast to the separate approach, where the cross-covariance enters the currency overlay only through the fixed asset weights $\boldsymbol{x}_t^{*\top} \boldsymbol{\Sigma}^{ac}$ and is not shrunk [31].

Applying the L_1 decomposition to all $N + M$ components of $\boldsymbol{\theta}_t$, and collecting the asset and currency parts separately, we form the augmented vector

$$\tilde{\boldsymbol{\theta}}_t = \Big[\, \boldsymbol{x}_t^\top, \quad \boldsymbol{\psi}_t^\top, \quad (\boldsymbol{\Delta}_{x,t}^-)^\top, \quad (\boldsymbol{\Delta}_{\psi,t}^-)^\top, \quad (\boldsymbol{\Delta}_{x,t}^+)^\top, \quad (\boldsymbol{\Delta}_{\psi,t}^+)^\top \,\Big]^\top \in \mathbb{R}^{3N+3M}.$$

The cost vector and penalty matrix are

$$\tilde{\boldsymbol{c}} = \Big[\, -\boldsymbol{\mu}^\top, \quad \lambda_1^a \mathbf{1}_N^\top, \quad \lambda_1^c \mathbf{1}_M^\top, \quad \lambda_1^a \mathbf{1}_N^\top, \quad \lambda_1^c \mathbf{1}_M^\top \,\Big]^\top,$$

$$\tilde{\boldsymbol{Q}} = \begin{bmatrix} \gamma\boldsymbol{\Sigma}^a + 2\lambda_2^a \boldsymbol{I}_N & \gamma\boldsymbol{\Sigma}^{ac} & \mathbf{0}_{N\times(2N+2M)} \\ \gamma(\boldsymbol{\Sigma}^{ac})^\top & \gamma\boldsymbol{\Sigma}^c + 2\lambda_2^c \boldsymbol{I}_M & \mathbf{0}_{M\times(2N+2M)} \\ \mathbf{0}_{(2N+2M)\times N} & \mathbf{0}_{(2N+2M)\times M} & \mathbf{0}_{(2N+2M)\times(2N+2M)} \end{bmatrix}.$$

The active block of $\tilde{\boldsymbol{Q}}$ combines the joint covariance with the L_2 penalty as $\gamma\boldsymbol{\Sigma} + 2\boldsymbol{S}_2^2$, where the off-diagonal $\gamma\boldsymbol{\Sigma}^{ac}$ couples the asset and currency weights. The joint portfolio optimization problem is then the quadratic program

$$\begin{aligned} \tilde{\boldsymbol{\theta}}_t^* = \arg\min_{\tilde{\boldsymbol{\theta}}_t} \quad & \tilde{\boldsymbol{c}}^\top \tilde{\boldsymbol{\theta}}_t + \tfrac{1}{2}\, \tilde{\boldsymbol{\theta}}_t^\top \tilde{\boldsymbol{Q}}\, \tilde{\boldsymbol{\theta}}_t, \\ \text{s.t.} \quad \tilde{\boldsymbol{D}}\tilde{\boldsymbol{\theta}}_t = \begin{bmatrix} \boldsymbol{q}_{N,M}^\top & \mathbf{0}_{N+M}^\top & \mathbf{0}_{N+M}^\top \\ \boldsymbol{I}_{N+M} & \boldsymbol{I}_{N+M} & -\boldsymbol{I}_{N+M} \end{bmatrix} \tilde{\boldsymbol{\theta}}_t = \tilde{\boldsymbol{E}} = \begin{bmatrix} 1 \\ \mathbf{0}_{N+M} \end{bmatrix}, \\ \tilde{\boldsymbol{A}}\tilde{\boldsymbol{\theta}}_t = \begin{bmatrix} \mathbf{0}_{N+M}^\top & \boldsymbol{q}_{N,M}^\top & \boldsymbol{q}_{N,M}^\top \end{bmatrix} \tilde{\boldsymbol{\theta}}_t \le b = \ell, \\ \Delta_{x,i,t}^\pm \ge 0, \quad \Delta_{\psi,c,t}^\pm \ge 0, \quad \text{for } i = 1,\dots,N,\ c = 1,\dots,M. \end{aligned} \tag{B.11}$$

In $\tilde{\boldsymbol{D}}$, the first row enforces the asset budget constraint $\boldsymbol{x}_t^\top \mathbf{1}_N = 1$, and the remaining $N + M$ rows enforce the weight decomposition (A.11) for all components of $\boldsymbol{\theta}_t$. The single row of $\tilde{\boldsymbol{A}}$ imposes the asset-only leverage constraint (A.12), giving $\mathbf{1}_N^\top(\boldsymbol{\Delta}_{x,t}^+ + \boldsymbol{\Delta}_{x,t}^-) = \|\boldsymbol{x}_t\|_1 \leq \ell$.

The joint formulation solves a single quadratic program over $\tilde{\boldsymbol{\theta}}_t$, capturing the interaction between asset and currency exposures through the cross-covariance matrix $\boldsymbol{\Sigma}^{ac}$ in $\tilde{\boldsymbol{Q}}$. Compared with the joint Expected Shortfall program (A.20), the mean-variance program is again smaller its decision vector has dimension $3N+3M$ rather than $3N+3M+q+1$, and its only inequality is the leverage constraint. As in the Expected Shortfall case, the joint program is of higher dimension than the two separate programs of Section B.3.

All quadratic programs derived in this appendix share the same convex structure and are solvable with the same quadratic programming solver. Each extension of the program, from the base regularized single-currency formulation (B.4) to the joint multi-currency program (B.11), modifies only the covariance and cost matrices and the constraint blocks that encode the budget, leverage, and weight decomposition, without changing the program structure.

Appendix C

Optimization Derivation: Mean Absolute Deviation

In this appendix, we derive the optimization problems used in the main text for portfolio optimization under the mean absolute deviation (MAD) risk measure. We begin by defining the MAD risk risk measure and formulate the base portoflio optimization problem. We then introduce regularization techniques and extend the formulation to the multi-currency setting, considering sequential overlay optimization and joint optimization approaches. Finally, we discuss the implementation of the optimization problems in Python using the CVXPY library.

C.1 Linear Program Represenation of the Mean Absolute Deviation Portfolio Optimization Problem

C.2 Regularization

C.2.1 L_1, L_2 Regularization and Leverage Constraint

C.2.2 Quadratic Program Representaion of the Regularized Portfolio Optimization Problem

C.3 Sequential Overlay Optimization

C.3.1 Asset Optimizaiton

C.3.2 Currency Overlay Optimization

C.4 Joint Optimization

C.5 Adding the Mean Term

Appendix D

Coherent Risk Measures and Deviation Measures

In this appendix, we define the notion of a deviation measure and a coherent risk measure, and prove that a strictly expectation bounded coherent risk measure can be transformed into a deviation measure by a simple de-meaning operation. This result is used in Section 6.2 to de-mean expected shortfall and second-order expected shortfall risk measures and transform them into deviation measures. Throughout this appendix, let $X \in L^2(\Omega)$ denote a square-integrable random variable on a probability space $(\Omega, \mathcal{F}, \mathbb{P})$.

Definition D.1 (Deviation measure) *A functional $\mathcal{D} : L^2(\Omega) \to [0, \infty]$ is a* deviation measure *if for all $X, X' \in L^2(\Omega)$, every constant C and every $\lambda > 0$,*

(D1)	$\mathcal{D}(X + C) = \mathcal{D}(X)$,
(D2)	$\mathcal{D}(0) = 0 \quad$ *and* $\quad \mathcal{D}(\lambda X) = \lambda \mathcal{D}(X)$,
(D3)	$\mathcal{D}(X + X') \leq \mathcal{D}(X) + \mathcal{D}(X')$,
(D4)	$\mathcal{D}(X) \geq 0$, *with* $\mathcal{D}(X) > 0$ *for nonconstant* X.

The property *(D1)* is the main one that distinguishes deviation measures from coherent risk measures. It is a form of translation invariance.

Definition D.2 (Coherent and strictly expectation bounded risk measure) *A functional $\mathcal{R} : L^2(\Omega) \to (-\infty, \infty]$ is a* coherent risk measure *if*

(R1)	$\mathcal{R}(X + C) = \mathcal{R}(X) - C$,
(R2)	$\mathcal{R}(0) = 0 \quad$ *and* $\quad \mathcal{R}(\lambda X) = \lambda \mathcal{R}(X)$ *for* $\lambda > 0$,
(R3)	$\mathcal{R}(X + X') \leq \mathcal{R}(X) + \mathcal{R}(X')$,
(R4)	$\mathcal{R}(X) \leq \mathcal{R}(X')$ *whenever* $X \geq X'$,

and it is strictly expectation bounded *if in addition*

(R5) $\mathcal{R}(X) > \mathbb{E}[-X] \quad$ *for every nonconstant* X, $\quad$ *and* $\quad \mathcal{R}(X) = \mathbb{E}[-X]$ *for constant* X.

As mentioned above, properties *(R1)* and *(D1)* are incompatible, so no non-degenerate functional can be a risk measure and a deviation measure at once. The de-meaning transformation is what moves between the two classes. We state and prove only the operation for transforming a coherent and strictly expectation bounded risk measure into a deviation measure. This is one direction of the more general result of Theorem 2 in [26].

Theorem D.3 *Let $\mathcal{R}$ be a strictly expectation bounded coherent risk measure. Then*

$$\mathcal{D}(X) \coloneqq \mathcal{R}(X - \mathbb{E}[X]) \tag{D.1}$$

is a deviation measure.

Proof The proof is a straightforward verification of the four properties in Definition D.1. Write $\bar{X} = X - \mathbb{E}[X]$, then $\mathcal{D}(X) = \mathcal{R}(\bar{X})$.

(D1): For any constant C we have $(X+C) - \mathbb{E}[X+C] = \bar{X}$, hence $\mathcal{D}(X+C) = \mathcal{R}(\bar{X}) = \mathcal{D}(X)$.

(D2): Taking $X = 0$ gives $\mathcal{D}(0) = \mathcal{R}(0) = 0$ by *(R2)*. For $\lambda > 0$, $\lambda X - \mathbb{E}[\lambda X] = \lambda \bar{X}$, so again, *(R2)* gives $\mathcal{D}(\lambda X) = \mathcal{R}(\lambda \bar{X}) = \lambda \mathcal{R}(\bar{X}) = \lambda \mathcal{D}(X)$.

(D3): Since $(X + X') - \mathbb{E}[X + X'] = \bar{X} + \bar{X}'$, *(R3)* gives $\mathcal{D}(X + X') = \mathcal{R}(\bar{X} + \bar{X}') \leq \mathcal{R}(\bar{X}) + \mathcal{R}(\bar{X}') = \mathcal{D}(X) + \mathcal{D}(X')$.

(D4): If X is constant then $\bar{X} = 0$ and $\mathcal{D}(X) = 0$. If X is nonconstant then so is $\bar{X}$, and *(R5)* applied to $\bar{X}$ gives $\mathcal{D}(X) = \mathcal{R}(\bar{X}) > \mathbb{E}[-\bar{X}] = 0$. □

The proof uses the properties *(R2)*, *(R3)* and *(R5)* only. Monotonicity *(R4)* is never used, which is the reason the original Theorem in [26] is stated more genetally for strictly expectation bounded risk measures rather than coherent risk measures.

We now show that both expected shortfall and second-order expected shortfall are strictly expectation bounded coherent risk measures, and we can therefore apply Theorem D.3 to transform them into deviation measures. Write $z^+ = \max\{z, 0\}$ and $z^- = \max\{-z, 0\}$, fix a confidence level $\alpha \in (0, 1)$ and set $\kappa \coloneqq (1 - \alpha)^{-1} > 1$. In the Rockafellar–Uryasev form [27] the expected shortfall is

$$\mathrm{ES}_\alpha(X) = \min_{\eta \in \mathbb{R}} \varphi(X, \eta), \tag{D.2}$$

and its second-order analogue, obtained by replacing the expectation in φ with the L^2 norm [13], is

$$\mathrm{SOES}_\alpha(X) = \min_{\eta \in \mathbb{R}} \psi(X, \eta), \tag{D.3}$$

where

$$\varphi(X, \eta) = \eta + \kappa\, \mathbb{E}\big[(-X - \eta)^+\big], \qquad \psi(X, \eta) = \eta + \kappa\, \big\|(-X - \eta)^+\big\|_2.$$

Both $\varphi(X, \cdot)$ and $\psi(X, \cdot)$ are convex, and both diverge as $|\eta| \to \infty$, so the minima in (D.2) and (D.3) are attained. Appendix A.1 provides additional details on the Rockafellar–Uryasev formulation. It is well known that ES_α is a coherent risk measure, see, for example Proposition 3.1 and Corollary 4.3 in [1]. We focus on showing it is strictly expectation bounded.

Proposition D.4 (Expected Shortfall) ES_α *is strictly expectation bounded.*

Proof Fix η, write $\mu = \mathbb{E}[X]$ and put $Y = -X - \eta$, so that $\mathbb{E}[Y] = -\mu - \eta$ and hence $\eta + \mu = -\mathbb{E}[Y]$. Using $\mathbb{E}[Y] = \mathbb{E}[Y^+] - \mathbb{E}[Y^-]$ we write

$$\begin{aligned} \varphi(X, \eta) - \mathbb{E}[-X] &= \eta + \mu + \kappa\, \mathbb{E}[Y^+] \\ &= -\mathbb{E}[Y] + \kappa\, \mathbb{E}[Y^+] \\ &= \mathbb{E}[Y^-] + (\kappa - 1)\, \mathbb{E}[Y^+]. \end{aligned}$$

Both terms on the right are non-negative, the second because $\kappa > 1$, so $\varphi(X, \eta) \geq \mathbb{E}[-X]$ for every η, with equality if and only if $Y = 0$ almost surely, that is $X \equiv -\eta$. If X is nonconstant the inequality is therefore strict at every η, and since the minimum is attained, $\mathrm{ES}_\alpha(X) > \mathbb{E}[-X]$. If $X \equiv C$, then $\eta = -C$ attains equality, so $\mathrm{ES}_\alpha(X) = \mathbb{E}[-X]$. □

Proposition D.5 (Second-order Expected Shortfall) SOES_α *is a coherent, strictly expectation bounded risk measure.*

Proof *(R1):* Substituting $\eta' = \eta + C$ gives $-X - C - \eta = -X - \eta'$ and $\psi(X + C, \eta) = \eta' - C + \kappa\|(-X - \eta')^+\|_2 = \psi(X, \eta') - C$. As η ranges over $\mathbb{R}$ so does η', so minimizing both sides gives $\mathrm{SOES}_\alpha(X + C) = \mathrm{SOES}_\alpha(X) - C$.

(R2): For $X = 0$ we have $\psi(0, \eta) = \eta + \kappa\,\eta^-$, which equals $\eta \geq 0$ for $\eta \geq 0$ and $(1 - \kappa)\eta > 0$ for $\eta < 0$. The minimum is 0, attained at $\eta = 0$. For $\lambda > 0$, substituting $\eta = \lambda\eta'$ and using positive homogeneity of $z \mapsto z^+$ and of the norm,

$$\psi(\lambda X, \lambda\eta') = \lambda\eta' + \kappa\|\lambda(-X - \eta')^+\|_2 = \lambda\left[\eta' + \kappa\|(-X - \eta')^+\|_2\right] = \lambda\,\psi(X, \eta'),$$

and minimizing over η' gives $\mathrm{SOES}_\alpha(\lambda X) = \lambda\,\mathrm{SOES}_\alpha(X)$.

(R3): Let η and η' attain the minima for X and X'. Since $(a + b)^+ \leq a^+ + b^+$ pointwise, and since $\|\cdot\|_2$ is monotone on non-negative random variables and satisfies the triangle inequality,

$$\|(-(X+X')-(\eta+\eta'))^+\|_2 \leq \|(-X-\eta)^+ + (-X'-\eta')^+\|_2 \leq \|(-X-\eta)^+\|_2 + \|(-X'-\eta')^+\|_2.$$

Adding $\eta + \eta'$ and multiplying the norms by κ, $\mathrm{SOES}_\alpha(X + X') \leq \psi(X + X', \eta + \eta') \leq \psi(X, \eta) + \psi(X', \eta') = \mathrm{SOES}_\alpha(X) + \mathrm{SOES}_\alpha(X')$.

(R4): If $X \geq X'$ then $-X - \eta \leq -X' - \eta$, so $(-X - \eta)^+ \leq (-X' - \eta)^+$ pointwise and therefore $\psi(X, \eta) \leq \psi(X', \eta)$ for every η. Taking the minimum gives $\mathrm{SOES}_\alpha(X) \leq \mathrm{SOES}_\alpha(X')$.

(R5): On a probability space $\|Z\|_2 \geq \|Z\|_1$, so for $Z = (-X - \eta)^+ \geq 0$ we have $\|(-X - \eta)^+\|_2 \geq \mathbb{E}[(-X - \eta)^+]$ and hence $\psi(X, \eta) \geq \varphi(X, \eta)$ for every η. Minimizing over η gives $\mathrm{SOES}_\alpha(X) \geq \mathrm{ES}_\alpha(X)$, and Proposition D.4 then gives $\mathrm{SOES}_\alpha(X) > \mathbb{E}[-X]$ for nonconstant X, with equality for constant X by *(R1)* and *(R2)*. □

Appendix E

Numerical Implementation of the Optimization Problems

In this appendix, we provide details on the numerical implementation of the optimization problems discussed in the thesis. We describe the algorithms used, the software libraries employed, and any specific considerations taken into account to ensure efficient and accurate solutions. Additionally, we discuss the computational complexity of the optimization problems and provide insights into the performance of the implemented algorithms across different scenarios.

E.1 The CVXPY Library

The optimization problems were implemented using the CVXPY library in Python, which provides a user-friendly interface for defining and solving convex optimization problems. For the considered risk measures in the thesis, we implemented the optimization problems as linear programs where possible, quadratic programs in most cases due to the L_2 regularization norm, and second-order cone programs for the second-order expected shortfall risk measure.

The CVXPY library is structured as follows: TODO: summarize the documentation of the library.

E.2 Numerical Implementation

The optimization problems derived in Chapter 3 are not assembled in their explicit matrix form in code. Each problem is instead expressed directly in the high-level modeling syntax of `CVXPY`, converted automatically into the standard form, and finally passed to the solver. For example, the L_1 penalty is written with `CVXPY`'s `norm1` atom, so the positive-negative variable split used in the derivation is internal to the solver and never appears explicitly in the code. The resulting problem class depends only on the regularization. The Expected Shortfall optimizer is a linear program when it is left unregularized, since both the Rockafellar–Uryasev objective and its constraints are linear, whereas the L_2 ridge penalty introduces a quadratic term that turns every regularized Expected Shortfall problem into a quadratic program. The mean–variance problems are quadratic programs by construction, through the variance term in their objective.

All programs are solved with Gurobi, a state-of-the-art solver for linear and quadratic programs,

accessed through `CVXPY`.[1] The solver is configured to use the barrier (interior point) method with tight convergence and accuracy tolerances, so that the large number of problems solved across rebalancing dates, universes, and hyperparameter configurations return reproducible results. The exact configuration is

```
problem.solve(
    solver=cp.GUROBI,
    verbose=False,
    BarConvTol=1e-8,
    Method=2,
    ScaleFlag=0,
    AggFill=0,
    PreDual=1,
    Presolve=1,
    Threads=3,
    FeasibilityTol=1e-9,
    OptimalityTol=1e-9,
)
```

The barrier method (`Method=2`) is paired with a barrier convergence tolerance of 10^{-8} and feasibility and optimality tolerances of 10^{-9}, the remaining flags specify the presolve, scaling, and threading behavior, which were optimized for the problem sizes considered in the thesis.

E.3 Computational Complexity and Performance

[1] Gurobi is a commercial solver and requires a license to use, a free fallback solver is available in the project codebase.

Appendix F

Transaction Costs

In this appendix, we detail the modeling of transaction costs in the context of international portfolio optimization. Following [31, 30, 16], and keeping the notation introduced in Chapter 3, we model transaction costs as proportional to the volume of traded assets and currencies, ignoring any fixed trading costs. The return of the hedged portfolio is expressed as the sum of unhedged asset returns and gains from forward currency positions

$$R^h_{P,t+1} = \sum_{i=1}^{N} x_{i,t} R^u_{i,t+1} + \sum_{c=1}^{M} \phi_{c,t}(f_{c,t} - e_{c,t+1}). \tag{3.2}$$

Let τ^{asset}_{t+1} and τ^{curr}_{t+1} denote the total transaction costs associated with asset and currency trades, respectively. The portfolio return net of transaction costs is then given by

$$R^{h,net}_{P,t+1} = R^h_{P,t+1} - \tau^{asset}_{t+1} - \tau^{curr}_{t+1}. \tag{F.1}$$

The total transaction costs are computed as the sum of the costs incurred from trading individual assets and currencies, weighted by the trading volumes

$$\tau^{asset}_{t+1} = \sum_{i=1}^{N} \tau^{asset}_{i,t+1} \left(x_{i,t+1} - x^-_{i,t+1} \right),$$

$$\tau^{curr}_{t+1} = \sum_{c=1}^{M} \tau^{curr}_{c,t+1} \phi_{c,t+1} + \sum_{c=1}^{M} \tau^{curr}_{c,t+1} \left(w_{c,t+1} - w^-_{c,t+1} \right),$$

where $\tau^{asset}_{i,t+1}$ and $\tau^{curr}_{c,t+1}$ are the per-unit transaction costs for asset i and currency c. The terms $x^-_{i,t+1}$ and $w^-_{c,t+1}$ denote the asset weights and implied currency weights immediately prior to rebalancing at time $t+1$ (i.e., the drifted weights).

Asset transaction costs are incurred only when rebalancing individual asset holdings. Currency transaction costs, however, arise from two sources, adjustments in forward hedge positions (assuming the forward contract maturity matches the rebalancing frequency, the old contract expires and a new contract of full notional size $\phi_{c,t+1}$ must be established at each rebalancing date) and currency exchanges required to fund purchases of assets denominated in different currencies ($w_{c,t+1} - w^-_{c,t+1}$). If the portfolio is rebalanced more frequently than the forward contract maturity, existing forward contracts are kept open and only incremental adjustments are traded. In this case, the forward transaction cost term is modeled as the absolute difference, replacing $\sum_{c=1}^{M} \tau^{curr}_{c,t+1} \phi_{c,t+1}$ with $\sum_{c=1}^{M} \tau^{curr}_{c,t+1} (\phi_{c,t+1} - \phi^-_{c,t+1})$, where $\phi^-_{c,t+1}$ represents the drifted hedge weight. Following [31], these costs are aggregated at the currency level, operating under

the assumption that proceeds from selling assets denominated in currency c can fund the purchase of other assets in the same currency without incurring currency conversion fees.

To define the individual per-unit costs $\tau^{asset}_{i,t+1}$ and $\tau^{curr}_{c,t+1}$, we consider an asset i denominated in currency c_i. Let $P^{ask}_{i,t}$, $P^{bid}_{i,t}$, and $P_{i,t}$ represent the ask, bid, and mid prices of the asset, and let $S^{ask}_{c_i,t}$, $S^{bid}_{c_i,t}$, and $S_{c_i,t}$ represent the corresponding currency prices. If asset i is purchased at time t and sold at time $t+1$, its realized log return is given by

$$\begin{aligned}
\log\left(\frac{P^{bid}_{i,t+1}S^{bid}_{c_i,t+1}}{P^{ask}_{i,t}S^{ask}_{c_i,t}}\right) &= \log\left(\frac{\left(P_{i,t+1}-\frac{\delta^{asset}_{i,t+1}}{2}\right)\left(S_{c_i,t+1}-\frac{\delta^{curr}_{c_i,t+1}}{2}\right)}{\left(P_{i,t}+\frac{\delta^{asset}_{i,t}}{2}\right)\left(S_{c_i,t}+\frac{\delta^{curr}_{c_i,t}}{2}\right)}\right) \\
&= \log\left(\frac{P_{i,t+1}\left(1-\frac{\delta^{asset}_{i,t+1}}{2P_{i,t+1}}\right)S_{c_i,t+1}\left(1-\frac{\delta^{curr}_{c_i,t+1}}{2S_{c_i,t+1}}\right)}{P_{i,t}\left(1+\frac{\delta^{asset}_{i,t}}{2P_{i,t}}\right)S_{c_i,t}\left(1+\frac{\delta^{curr}_{c_i,t}}{2S_{c_i,t}}\right)}\right) \\
&= \log\left(\frac{P_{i,t+1}S_{c_i,t+1}}{P_{i,t}S_{c_i,t}}\right)+\log\left(\frac{1-\frac{\delta^{asset}_{i,t+1}}{2P_{i,t+1}}}{1+\frac{\delta^{asset}_{i,t}}{2P_{i,t}}}\right)+\log\left(\frac{1-\frac{\delta^{curr}_{c_i,t+1}}{2S_{c_i,t+1}}}{1+\frac{\delta^{curr}_{c_i,t}}{2S_{c_i,t}}}\right) \\
&= \log\left(\frac{P_{i,t+1}S_{c_i,t+1}}{P_{i,t}S_{c_i,t}}\right)-\log\left(\frac{1+\frac{\delta^{asset}_{i,t}}{2P_{i,t}}}{1-\frac{\delta^{asset}_{i,t+1}}{2P_{i,t+1}}}\right)-\log\left(\frac{1+\frac{\delta^{curr}_{c_i,t}}{2S_{c_i,t}}}{1-\frac{\delta^{curr}_{c_i,t+1}}{2S_{c_i,t+1}}}\right),
\end{aligned}$$

where δ^{asset}_i and $\delta^{curr}_{c_i}$ are the bid-ask spreads. Comparing this to equation (F.1), the transaction cost is precisely the difference between the actual realized log return and mid-price returs for the asset and currency.

Defining half of the bid-ask spreads as

$$\chi^{asset}_{i,t} = \frac{\overbrace{P^{ask}_{i,t}-P^{bid}_{i,t}}^{\delta^{asset}_i}}{2P_{i,t}} \quad \text{and} \quad \chi^{curr}_{c,t} = \frac{\overbrace{S^{ask}_{c,t}-S^{bid}_{c,t}}^{\delta^{curr}_c}}{2S_{c,t}},$$

the per-unit transaction costs for trading asset i and currency c become

$$\tau^{asset}_{i,t+1} = \begin{cases} \log\left(\frac{1+\chi^{asset}_{i,t+1}}{1-\chi^{asset}_{i,t}}\right) & \text{if } x_{i,t+1} > x^{-}_{i,t+1}, \\ \log\left(\frac{1-\chi^{asset}_{i,t}}{1+\chi^{asset}_{i,t+1}}\right) & \text{otherwise,} \end{cases}$$

and

$$\tau^{curr}_{c,t+1} = \begin{cases} \log\left(\frac{1+\chi^{curr}_{c,t+1}}{1-\chi^{curr}_{c,t}}\right) & \text{if } \phi_{c,t+1} > 0 \text{ or } w_{c,t+1} > w^{-}_{c,t+1}, \\ \log\left(\frac{1-\chi^{curr}_{c,t}}{1+\chi^{curr}_{c,t+1}}\right) & \text{otherwise.} \end{cases}$$

For our empirical analysis, we adopt the simplifying assumption that proportional bid-ask spreads remain constant both over time and across the cross-sections of both assets and currencies. We set the half-spreads to $\chi^{asset} = 0.001$ (10 bps) for all equities and $\chi^{curr} = 0.0001$ (1 bp) for all currencies, consistent with empirical finidngs in [10, 23, 21, 7]. Under these assumptions, the transaction costs simplify to constant values of $\tau^{asset} \approx 0.002$ and $\tau^{curr} \approx 0.0002$.

Appendix G

Additional Results

In this appendix, we justify the bounded hyperparameter grids over which the regularization penalties are selected in the empirical analysis of Chapter 5. In Section G.1 we describe the sensitivity analysis used to determine these bounds, following a similar approach as [30].

TODO: Summarize the sparsity analysis and its purpose.

this is from the data chapter, maybe it will come in handy Hyperparameter selection follows the walk-forward grid search described in Section 4.2.1. The tuned hyperparameters are the L_1 and L_2 regularization coefficients applied separately to the asset and currency weights, λ_1^a, λ_2^a, λ_1^c, and λ_2^c, and the grid is the full cartesian product of their individual ranges. Each candidate configuration is scored by computing, on the validation window, the objective the optimizer minimizes, evaluated at the weights that configuration produces. The other parameters specific to the optimziation problems (these include the risk-aversion coefficient, the leverage budget, and the net currency exposure bounds) are fixed a priori rather than tuned. For each penalty the lower bound is zero, corresponding to the absence of that regularization term, while the upper bounds differ across parameters and are chosen so that larger values leave the optimal weights essentially unchanged.

G.1 Hyperparameter Sensitivity Analysis

The regularized portfolio optimization problems include four regularization parameters, the asset penalties λ_1^a, λ_2^a and the currency penalties λ_1^c, λ_2^c. These are selected by grid search on a validation sample, which requires setting a bounded interval for each. The bounds are not arbitrary. As a penalty grows, the optimal weights approach a limiting allocation, and beyond some threshold a further increase leaves the solution essentially unchanged. A grid whose upper bound lies past that threshold is sufficient, in the sense that extending it further would only increase the computational cost of the grid search.

To locate this threshold we follow [30] and define, for each penalty $\lambda \in \{\lambda_1^a, \lambda_2^a, \lambda_1^c, \lambda_2^c\}$, an increasing sequence $(\lambda_n)_{n\in\mathbb{N}}$ with $\lambda_0 = 0$. Let $\boldsymbol{x}_t(\lambda_n)$ and $\boldsymbol{\psi}_t(\lambda_n)$ denote the optimal asset and currency overlay weights when only λ is active, that is, when the remaining three penalties are held at zero. The marginal effect of an incremental increase in regularization is measured in the supremum norm

$$d_n^a = \|\boldsymbol{x}_t(\lambda_n) - \boldsymbol{x}_t(\lambda_{n-1})\|_\infty \quad \text{for } \lambda \in \{\lambda_1^a, \lambda_2^a\},$$
$$d_n^c = \|\boldsymbol{\psi}_t(\lambda_n) - \boldsymbol{\psi}_t(\lambda_{n-1})\|_\infty \quad \text{for } \lambda \in \{\lambda_1^c, \lambda_2^c\}.$$

Tthe two penalty types approach their limits in different ways. Under an L_1 penalty the weights move to the corner of the feasible set, at which point the penalty term becomes

constant in the weights and doesnt influence the solution. An L_2 penalty instead shrinks the weights smoothly toward the origin and never reaches a corner, so its marginal effect decays gradually without ever vanishing. Such behavior is well documented in [**?**]. For this reason we report alongside d_n the marginal effect on portfolio exposure, $\|\boldsymbol{x}_t\|_1$ for the asset penalties and $\|\boldsymbol{\psi}_t\|_1$ for the currency penalties. Exposure is the economically interpretable quantity, and is the one that we use to determine the upper bound of the grid for the L_2 penalties.

The analysis is done on the first estimation window spanning July 1990 to June 2000. To account for sampling error, each reported quantity is averaged over 5 bootstraped replicates of the estimation window. The results are reported for the small investable universe.

Expected Shortfall Hyperparameter Sensitivity Analysis

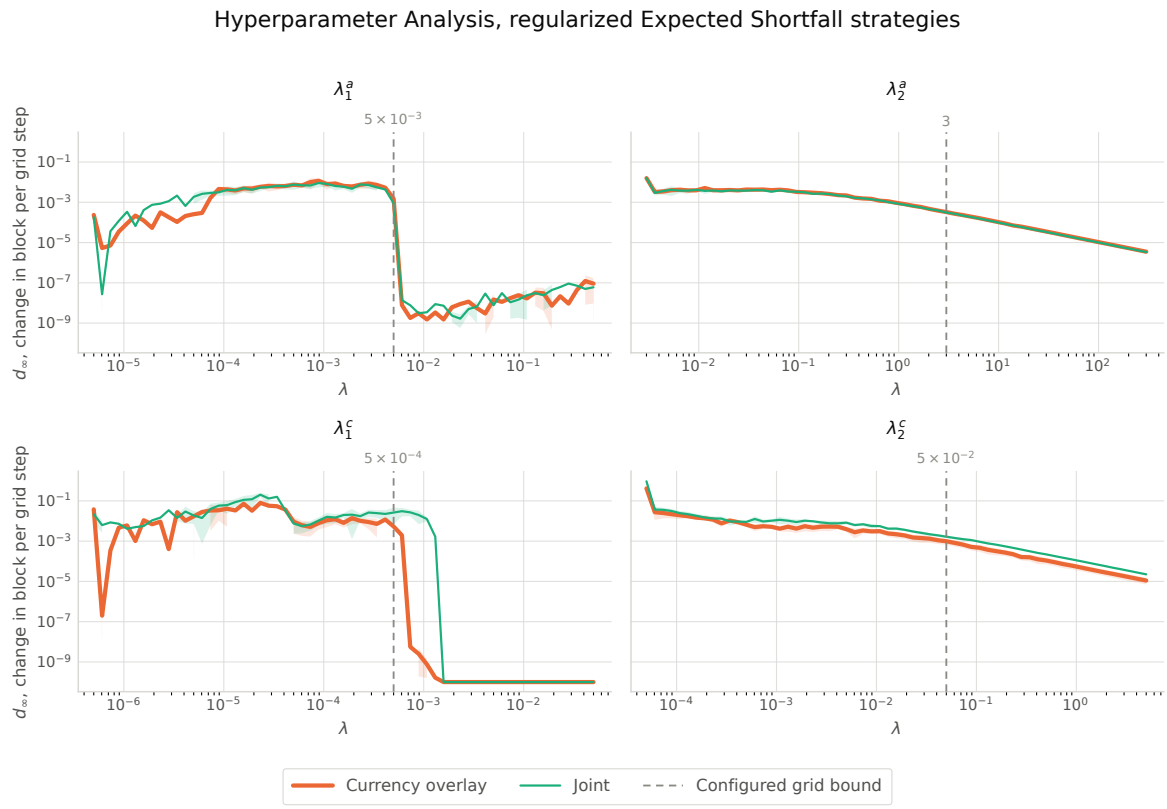


Figure G.1: Marginal effect of each regularization penalty, for the regularized Expected Shortfall currency overlay and joint optimizers. The dashed line marks the upper bound of the hyperparameter grid reported in Table G.1. The upper plots report d_n^a and the lower plots d_n^c. Estimates are averaged over five bootstrap replicates of the first estimation window on the small investment universe.

Figure G.1 and Figure G.2 report the marginal effect and the portfolio exposure. The difference between the two penalty types is clearest in the marginal effect plots. The L_1 penalties collapse abruptly near their bound, whereas the L_2 penalties decay smoothly throughout, which is why their bounds are determined by the exposure analysis instead.

For both asset penalties, exposure falls from the leverage limit in the unpenalized problem to 1. The budget constraint $\boldsymbol{x}_t^\top \mathbf{1}_N = 1$ implies $\|\boldsymbol{x}_t\|_1 \geq 1$, with equality precisely when the portfolio holds no short positions. The two penalties reach this floor for different reasons, and only the L_1 penalty ceases to act once it is there. Under λ_1^a the penalty term equals $\|\boldsymbol{x}_t\|_1 = 1$ everywhere on the long-only set and is therefore constant in $\boldsymbol{x}_t$, so the solution freezes exactly.

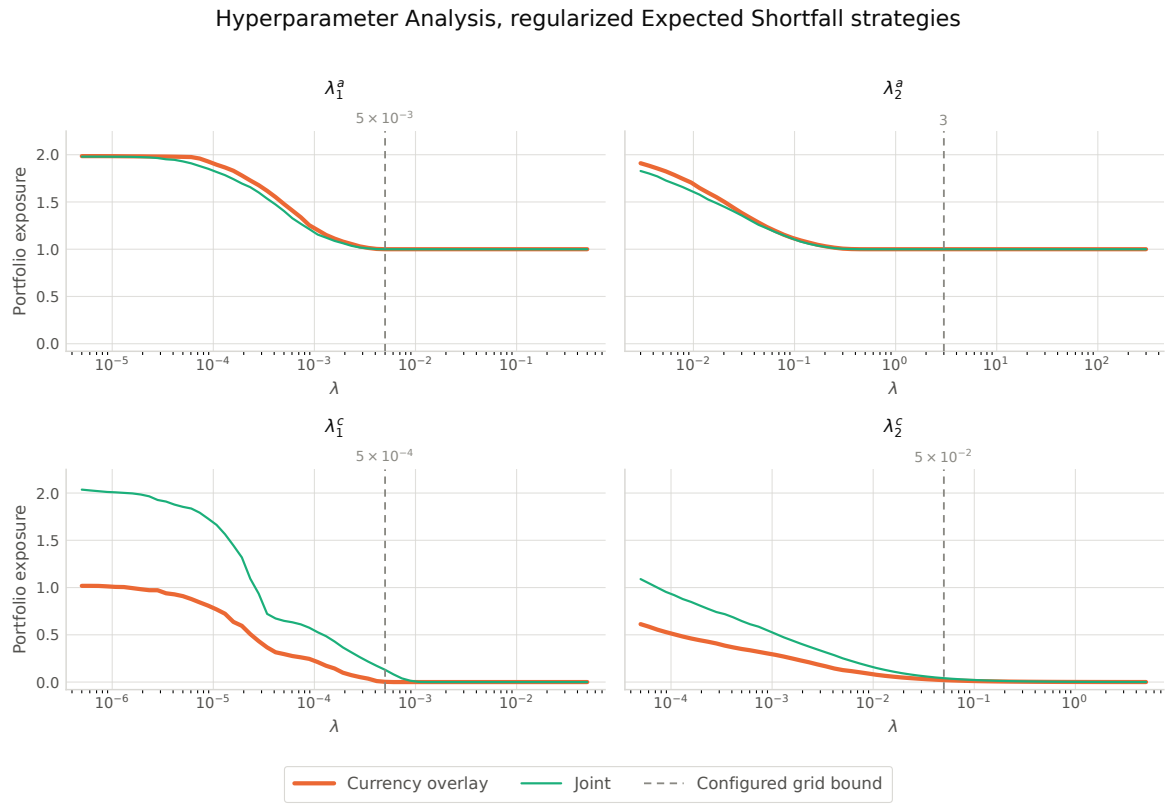


Figure G.2: Portfolio exposure as a function of each regularization penalty, for the regularized Expected Shortfall currency overlay and joint optimizers. The dashed line marks the upper bound of the validation grid reported in Table G.1. The upper panels report $\|\boldsymbol{x}_t\|_1$ and the lower panels $\|\boldsymbol{\psi}_t\|_1$. Estimates are averaged over five bootstrap replicates of the first estimation window on the small investment universe.

Under λ_2^a the penalty $\|\boldsymbol{x}_t\|_2^2$ is not constant on that set: exposure stays pinned at its floor while the weights keep spreading toward the equally weighted portfolio, which is the limiting allocation as $\lambda_2^a \to \infty$. For the asset L_2 penalty we therefore read the bound as the point at which the short positions have been eliminated and any further movement is a redistribution within a long-only book.

The overlay weights carry neither a budget nor a leverage constraint, so $\boldsymbol{\psi}_t$ shrinks to zero, exactly under λ_1^c and asymptotically under λ_2^c.

On the basis of this analysis we fix the ranges reported in Table G.1. Each range is discretised on a log-spaced grid that includes the zero endpoint, which corresponds to the unregularized optimization problem.

Hyperparameter	Description	Range
λ_1^a	Asset L_1 penalty	$[0,\ 5 \times 10^{-3}]$
λ_2^a	Asset L_2 penalty	$[0,\ 3]$
λ_1^c	Currency L_1 penalty	$[0,\ 5 \times 10^{-4}]$
λ_2^c	Currency L_2 penalty	$[0,\ 5 \times 10^{-2}]$

Table G.1: Hyperparameter ranges for the regularized Expected Shortfall optimizers. Each penalty is searched over a log-spaced grid spanning the reported range, where the lower endpoint corresponds to the unregularized problem and the upper endpoint is the bound established in Section G.1.

Mean-Variance Hyperparameter Sensitivity Analysis

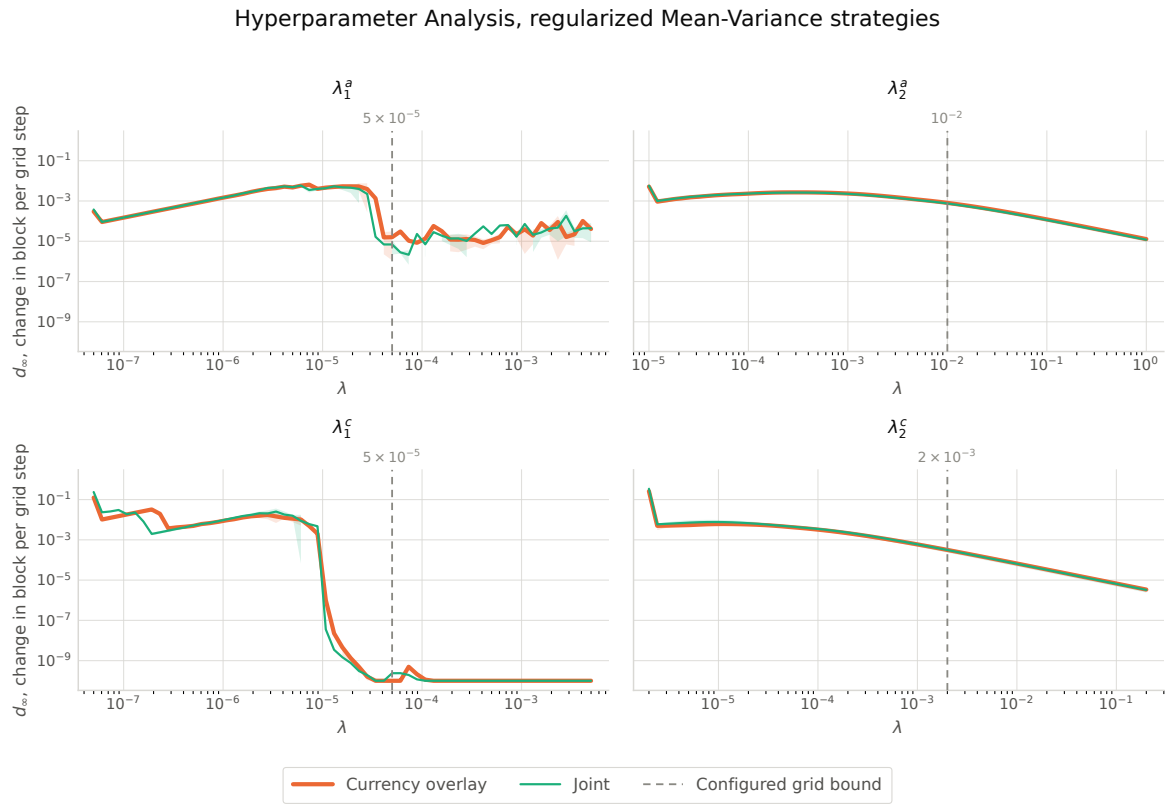


Figure G.3: Marginal effect of each regularization penalty, for the regularized mean-variance currency overlay and joint optimizers. The dashed line marks the upper bound of the hyperparameter grid reported in Table G.2. The upper plots report d_n^a and the lower plots d_n^c. Estimates are averaged over five bootstrap replicates of the first estimation window on the small investment universe.

Figure G.3 and Figure G.4 repeat the analysis for the regularized mean-variance optimizers. What differs is the scale. The penalties are added to a risk term measured in units of daily variance rather than daily expected shortfall, and the two differ by roughly two orders of magnitude. On the first estimation window the equally weighted portfolio has a daily variance on the order of 10^{-4} against a daily expected shortfall (at 80% level) of approximately $\times 10^{-2}$.

We fix the hyperparameter ranges reported in Table G.2 for the regularized mean-variance optimizers, which are discretized on a log-spaced grid the same way as for the Expected Shortfall optimizers.

Hyperparameter	Description	Range
λ_1^a	Asset L_1 penalty	$[0,\ 5 \times 10^{-5}]$
λ_2^a	Asset L_2 penalty	$[0,\ 10^{-2}]$
λ_1^c	Currency L_1 penalty	$[0,\ 5 \times 10^{-5}]$
λ_2^c	Currency L_2 penalty	$[0,\ 2 \times 10^{-3}]$

Table G.2: Hyperparameter ranges for the regularized mean-variance optimizers. Each penalty is searched over a log-spaced grid spanning the reported range, where the lower endpoint corresponds to the unregularized problem and the upper endpoint is the bound established in Section G.1.

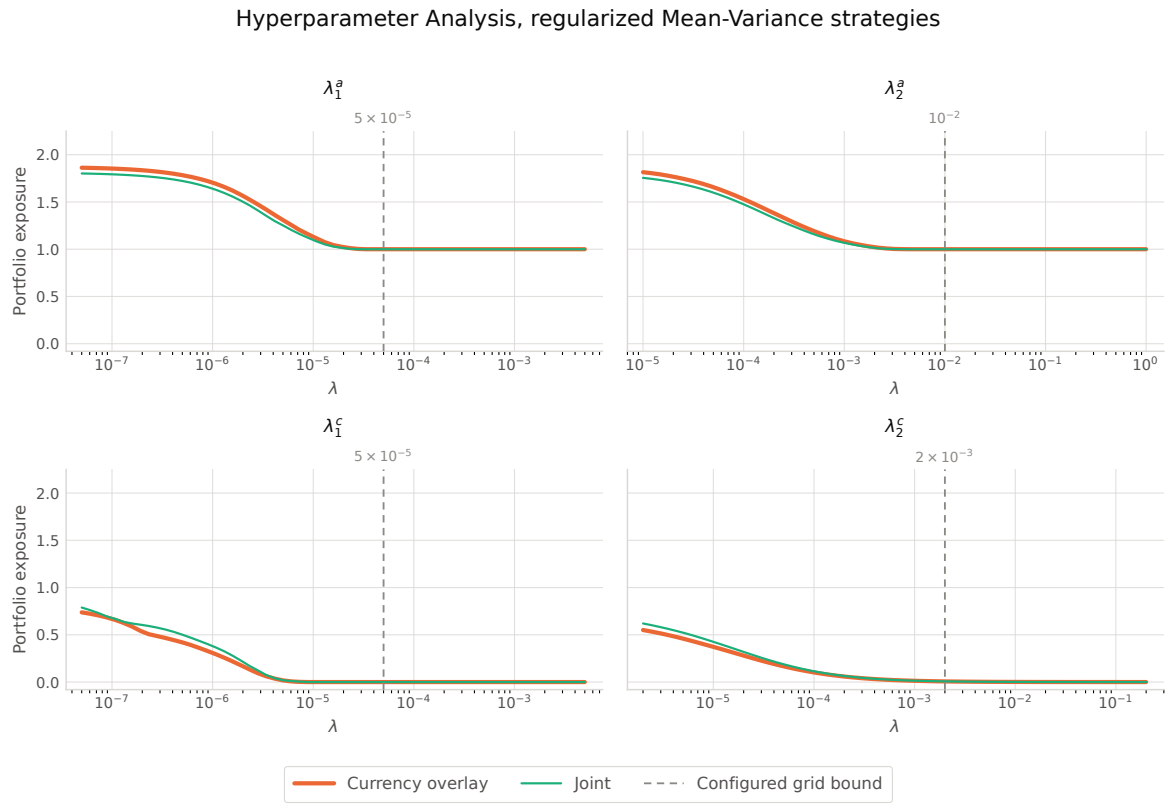


Figure G.4: Portfolio exposure as a function of each regularization penalty, for the regularized mean-variance currency overlay and joint optimizers. The dashed line marks the upper bound of the validation grid reported in Table G.2. The upper panels report $\|x_t\|_1$ and the lower panels $\|\psi_t\|_1$. Estimates are averaged over five bootstrap replicates of the first estimation window on the small investment universe.

Mean Absolute Deviation Hyperparameter Sensitivity Analysis

G.2 Portfolio Sparsity Analysis

Sparsity of the Regularized Expected Shortfall Optimizers

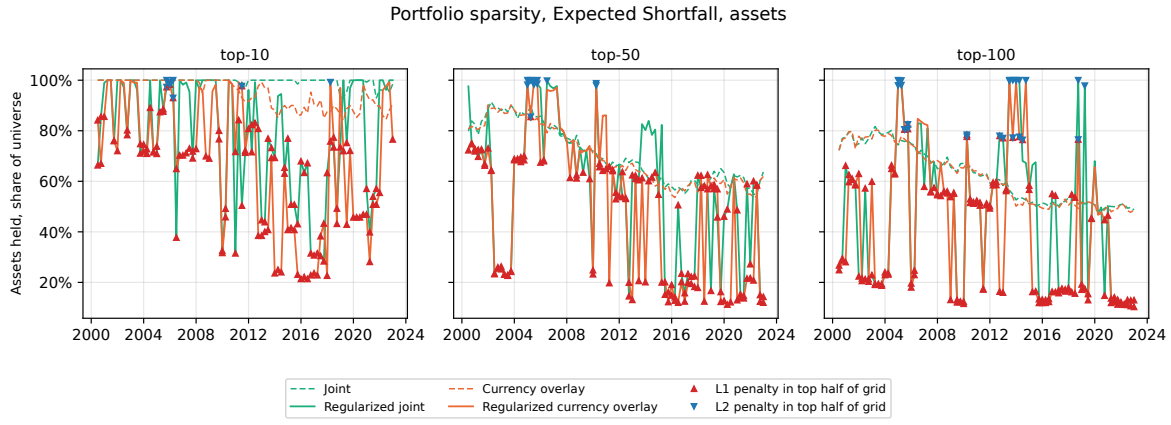


Figure G.5: Sparsity of the regularized Expected Shortfall currency overlay and joint optimizers.

Figure G.5 reports the share of the investable universe held at each rebalance date.

Regularization stabilizes the weight vector conditional on the penalty coefficients, but those coefficients are themselves re-selected at every rebalance date. The number of active positions is sensitive to this recalibration, an L1-dominant selection sets weights to zero and yields a concentrated portfolio, whereas an L2-dominant selection shrinks all weights toward the equally weighted portfolio. The oscillation therefore reflects the validation criterion alternating between these two regimes across quarters, not instability in the optimizer itself. The unregularized problem has no such degree of freedom. Its support is determined primarily by the leverage limit.

Sparsity of the Regularized Mean-Variance Optimizers

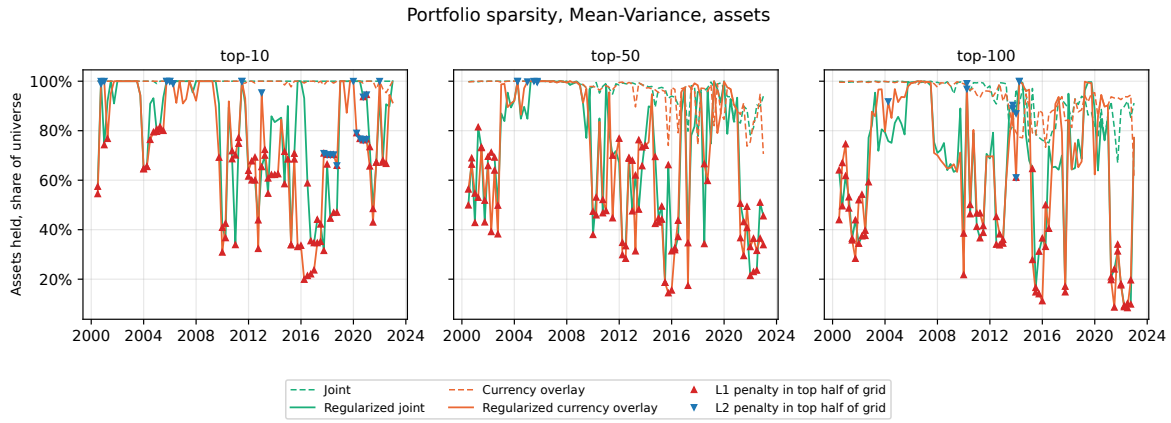


Figure G.6: Sparsity of the regularized Mean-Variance currency overlay and joint optimizers.

Sparsity of the Regularized Mean Absolute Deviation Optimizers

G.3 Risk Aversion Instability under Mean Variance Optimization

In Chapter 5 we report backtest results for the mean-variance optimizers under infinite risk aversion, which corresponds to the minimum variance problem. We do this because the mean-variance optimization problem is is unstable under finite risk aversion [**?**]. We showcase this by runnign the regularized overlay and joint optimizers under finite risk aversion on the top 10 universe and report the results in G.3.

We can see from table G.3 that the performance of the optimizers is highly unstable and results in poor out of sample performance of the portfolios. Two

G.4 Additional Backtest Performance Results

Mean-Variance Optimizers under Finite Risk Aversion		
Metric	**Overlay**	**Joint**
ES (%)	2.07	2.20
Vol (%)	26.50	28.54
MD (%)	70.15	69.31
SR	0.10	0.11
SOR	0.07	0.15
ESR	0.08	0.09
Turnover^a	1.82	1.84
Turnover^c	19.04	18.60

Table G.3: Performance metrics for the regularized mean-variance currency overlay and joint optimizers under finite risk aversion on the top 10 universe.

Panel A: Small Investment Universe (top 10 stocks per economy)							
Metric	**Equal Weight**	**Assets-Full**	**Assets-Zero**	**Overlay**	**Overlay-Reg**	**Joint**	**Joint-Reg**
ES (%)	1.15	0.71	0.98	0.69	0.67	0.67	0.66
Vol (%)	13.96	8.51	11.84	8.18	7.99	7.97	7.85
MD (%)	54.76	30.81	44.28	29.31	29.77	28.31	28.89
SR	0.12	0.10	0.21	0.08	0.13	0.08	0.17
SOR	0.17	0.13	0.28	0.10	0.18	0.11	0.23
ESR	0.09	0.08	0.15	0.06	0.10	0.06	0.13
Turnover^a	0.29	1.27	1.23	1.30	1.06	1.27	1.15
Turnover^c	0.92	0.93	0	2.39	1.67	2.98	1.77
Panel B: Medium Investment Universe (top 50 stocks per economy)							
Metric	**Equal Weight**	**Assets-Full**	**Assets-Zero**	**Overlay**	**Overlay-Reg**	**Joint**	**Joint-Reg**
ES (%)	1.04	0.43	0.80	0.44	0.42	0.44	0.42
Vol (%)	12.50	5.84	9.87	5.68	5.73	5.53	5.58
MD (%)	55.74	38.54	42.37	38.59	38.56	38.84	36.38
SR	0.32	0.56	0.45	0.56	0.55	0.54	0.55
SOR	0.44	0.77	0.64	0.75	0.77	0.73	0.76
ESR	0.24	0.48	0.36	0.45	0.48	0.43	0.46
Turnover^a	0.27	1.60	1.58	1.78	1.52	1.78	1.46
Turnover^c	0.92	0.95	0	0.97	1.26	1.67	1.29
Panel C: Large Investment Universe (top 100 stocks per economy)							
Metric	**Equal Weight**	**Assets-Full**	**Assets-Zero**	**Overlay**	**Overlay-Reg**	**Joint**	**Joint-Reg**
ES (%)	0.97	0.34	0.76	0.35	0.33	0.35	0.33
Vol (%)	11.67	4.46	9.21	4.57	4.45	4.55	4.40
MD (%)	56.26	31.60	35.12	31.32	30.61	30.45	30.11
SR	0.37	0.89	0.57	0.72	0.87	0.76	0.89
SOR	0.49	1.18	0.80	0.95	1.14	1.01	1.17
ESR	0.28	0.76	0.44	0.59	0.74	0.62	0.74
Turnover^a	0.27	1.61	1.61	1.89	1.42	1.91	1.62
Turnover^c	0.91	0.96	0	0.97	1.15	1.43	1.23

Table G.4: Performance metrics for Expected Shortfall optimization approaches on increasing universe sizes

Table G.4 reports out-of-sample performance metrics for the international portfolio optimization under expected shortfall, with portfolios rebalanced quarterly. The realized out-of-sample statistics shown are Expected Shortfall (ES), volatility (Vol), maximum drawdown (MD), the Sharpe Ratio (SR), the Sortino Ratio (SOR), the Expected Shortfall Ratio (ESR), and average asset and currency turnover, denoted Turnover^a and Turnover^c, respectively. Each panel corresponds to a different investment universe: Panel A covers the small universe, made up of the 10 largest companies by market capitalization in each economy; Panel B the medium universe, comprising the 50 largest companies per economy; and Panel C the large universe, consisting of the 100 largest companies per economy. The sample period runs from June 1990 to June 2023, and the first out-of-sample portfolio is formed in June 2000. Details on the optimization approaches and performance metrics are provided in section **??**.

Panel A: Small Investment Universe (top 10 stocks per economy)							
Metric	**Equal Weight**	**Assets-Full**	**Assets-Zero**	**Overlay**	**Overlay-Reg**	**Joint**	**Joint-Reg**
ES (%)	1.15	0.69	0.97	0.66	0.65	0.64	0.64
Vol (%)	13.96	8.39	11.80	7.97	7.82	7.73	7.75
MD (%)	54.76	28.00	42.18	28.23	27.80	27.92	27.79
SR	0.12	0.12	0.22	0.14	0.17	0.16	0.17
SOR	0.17	0.17	0.30	0.19	0.22	0.22	0.22
ESR	0.09	0.10	0.17	0.11	0.12	0.12	0.12
Turnover^a	0.29	1.02	1.02	0.87	0.92	0.82	0.98
Turnover^c	0.92	0.93	0	1.70	1.43	1.88	1.47
Panel B: Medium Investment Universe (top 50 stocks per economy)							
Metric	**Equal Weight**	**Assets-Full**	**Assets-Zero**	**Overlay**	**Overlay-Reg**	**Joint**	**Joint-Reg**
ES (%)	1.04	0.41	0.79	0.41	0.40	0.40	0.40
Vol (%)	12.50	5.31	9.58	5.17	5.26	5.09	5.19
MD (%)	55.74	35.05	42.05	35.85	35.03	34.89	34.26
SR	0.32	0.48	0.40	0.47	0.45	0.50	0.45
SOR	0.44	0.62	0.55	0.62	0.59	0.67	0.59
ESR	0.24	0.39	0.31	0.37	0.37	0.40	0.37
Turnover^a	0.27	1.07	1.06	1.04	1.12	1.05	1.13
Turnover^c	0.92	0.95	0	1.16	1.11	1.37	1.11
Panel C: Large Investment Universe (top 100 stocks per economy)							
Metric	**Equal Weight**	**Assets-Full**	**Assets-Zero**	**Overlay**	**Overlay-Reg**	**Joint**	**Joint-Reg**
ES (%)	0.97	0.32	0.74	0.32	0.31	0.31	0.31
Vol (%)	11.67	4.37	8.99	4.24	4.35	4.19	4.29
MD (%)	56.26	26.60	33.06	26.72	26.28	26.36	26.40
SR	0.37	0.72	0.49	0.65	0.77	0.67	0.76
SOR	0.49	0.93	0.68	0.85	0.99	0.89	1.00
ESR	0.28	0.63	0.38	0.54	0.67	0.55	0.65
Turnover^a	0.27	1.09	1.12	1.08	1.15	1.08	1.10
Turnover^c	0.91	0.96	0	1.06	1.06	1.24	1.03

Table G.5: Performance metrics for mean-variance optimization approaches on increasing universe sizes

Table G.5 reports out-of-sample performance metrics for the international portfolio optimization under mean-variance, with portfolios rebalanced quarterly. The realized out-of-sample statistics shown are Expected Shortfall (ES), volatility (Vol), maximum drawdown (MD), the Sharpe Ratio (SR), the Sortino Ratio (SOR), the Expected Shortfall Ratio (ESR), and average asset and currency turnover, denoted Turnover^a and Turnover^c, respectively. Each panel corresponds to a different investment universe: Panel A covers the small universe, made up of the 10 largest companies by market capitalization in each economy; Panel B the medium universe, comprising the 50 largest companies per economy; and Panel C the large universe, consisting of the 100 largest companies per economy. The sample period runs from June 1990 to June 2023, and the first out-of-sample portfolio is formed in June 2000. Details on the optimization approaches and performance metrics are provided in section **??**.

Panel A: Small Investment Universe (top 10 stocks per economy)							
Metric	**Equal Weight**	**Assets-Full**	**Assets-Zero**	**Overlay**	**Overlay-Reg**	**Joint**	**Joint-Reg**
ES (%)	1.15	0.69	0.96	0.65	0.65	0.64	0.64
Vol (%)	13.96	8.36	11.70	7.87	7.82	7.71	7.74
MAD (%)	54.76	5.67	8.06	5.51	5.47	5.38	5.35
MD (%)	54.76	28.59	40.88	27.67	27.83	27.31	26.01
SR	0.12	0.20	0.27	0.16	0.21	0.18	0.22
SOR	0.17	0.27	0.36	0.22	0.27	0.25	0.29
ESR	0.09	0.15	0.21	0.12	0.16	0.14	0.16
Turnover^a	0.29	1.08	1.08	1.10	1.04	1.06	1.06
Turnover^c	0.92	0.92	0	1.51	1.36	1.84	1.44
Panel B: Medium Investment Universe (top 50 stocks per economy)							
Metric	**Equal Weight**	**Assets-Full**	**Assets-Zero**	**Overlay**	**Overlay-Reg**	**Joint**	**Joint-Reg**
ES (%)	1.04	0.40	0.79	0.41	0.40	0.41	0.40
Vol (%)	12.50	6.19	9.97	5.72	6.16	5.55	6.08
MAD (%)	55.74	3.29	6.56	3.47	3.26	3.44	3.24
MD (%)	55.74	35.90	43.81	35.93	35.92	36.94	35.51
SR	0.32	0.48	0.39	0.48	0.46	0.48	0.49
SOR	0.44	0.77	0.57	0.69	0.75	0.69	0.79
ESR	0.24	0.46	0.31	0.42	0.44	0.41	0.47
Turnover^a	0.27	1.30	1.31	1.67	1.30	1.69	1.27
Turnover^c	0.92	0.96	0	1.98	1.09	1.25	1.05
Panel C: Large Investment Universe (top 100 stocks per economy)							
Metric	**Equal Weight**	**Assets-Full**	**Assets-Zero**	**Overlay**	**Overlay-Reg**	**Joint**	**Joint-Reg**
ES (%)	0.97	0.33	0.76	0.33	0.32	0.32	0.32
Vol (%)	11.67	4.94	9.36	4.63	4.97	4.57	4.85
MAD (%)	56.26	2.67	6.40	2.78	2.65	2.75	2.62
MD (%)	56.26	31.28	36.99	31.25	31.28	31.41	31.62
SR	0.37	0.57	0.44	0.53	0.55	0.52	0.56
SOR	0.49	0.77	0.62	0.70	0.73	0.69	0.76
ESR	0.28	0.54	0.34	0.47	0.52	0.46	0.53
Turnover^a	0.27	1.58	1.58	1.74	1.58	1.67	1.46
Turnover^c	0.91	0.97	0	0.97	1.01	1.16	1.03

Table G.6: Performance metrics for mean absolute deviation optimization approaches on increasing universe sizes

Table G.6 reports out-of-sample performance metrics for the international portfolio optimization under mean absolute deviation, with portfolios rebalanced quarterly. The realized out-of-sample statistics shown are Expected Shortfall (ES), volatility (Vol), maximum drawdown (MD), the Sharpe Ratio (SR), the Sortino Ratio (SOR), the Expected Shortfall Ratio (ESR), and average asset and currency turnover, denoted Turnover^a and Turnover^c, respectively. Each panel corresponds to a different investment universe: Panel A covers the small universe, made up of the 10 largest companies by market capitalization in each economy; Panel B the medium universe, comprising the 50 largest companies per economy; and Panel C the large universe, consisting of the 100 largest companies per economy. The sample period runs from June 1990 to June 2023, and the first out-of-sample portfolio is formed in June 2000. Details on the optimization approaches and performance metrics are provided in section **??**.

Bibliography

[1] Carlo Acerbi and Dirk Tasche. On the coherence of expected shortfall, 2002.

[2] Stephen Boyd and Lieven Vandenberghe. *Convex Optimization*. Cambridge University Press, Cambridge, 2004.

[3] JOHN Y. CAMPBELL, KARINE SERFATY-DE MEDEIROS, and LUIS M. VICEIRA. Global currency hedging. *The Journal of Finance*, 65(1):87–121, 2010.

[4] Yu chin Chen and Kenneth Rogoff. Commodity currencies. *Journal of International Economics*, 60(1):133–160, 2003. Emperical Exchange Rate Models.

[5] Stefano Ciliberti, Imre Kondor, and Marc Mézard. On the feasibility of portfolio optimization under expected shortfall. *Quantitative Finance*, 7(4):389–396, 2007.

[6] Noureddine El Karoui. High-dimensionality effects in the Markowitz problem and other quadratic programs with linear constraints: Risk underestimation. *The Annals of Statistics*, 38(6):3487–3566, 2010.

[7] Andrea Frazzini, Ronen Israel, and Tobias J. Moskowitz. Trading costs. 2018. Working paper, AQR Capital Management / Yale University.

[8] Lisa R. Goldberg, Michael Y. Hayes, and Ola Mahmoud. Minimizing shortfall, 2013.

[9] Trevor Hastie, Robert Tibshirani, and Jerome Friedman. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer Series in Statistics. Springer, New York, 2 edition, 2009.

[10] Charles M. Jones. A century of stock market liquidity and trading costs. 2002. Working paper, Columbia University.

[11] Philippe Jorion. Bayes-stein estimation for portfolio analysis. *The Journal of Financial and Quantitative Analysis*, 21(3):279–292, 1986.

[12] Imre Kondor, Szilárd Pafka, and Gábor Nagy. Noise sensitivity of portfolio selection under various risk measures. *Journal of Banking & Finance*, 31(5):1545–1573, 2007.

[13] Pavlo A. Krokhmal. Higher moment coherent risk measures. *Quantitative Finance*, 7(4):373–387, 2007.

[14] Olivier Ledoit and Michael Wolf. Honey, i shrunk the sample covariance matrix. Economics Working Papers 691, Department of Economics and Business, Universitat Pompeu Fabra, Jun 2003.

[15] Olivier Ledoit and Michael Wolf. A well-conditioned estimator for large-dimensional covariance matrices. *Journal of Multivariate Analysis*, 88(2):365–411, 2004.

[16] Jiahan Li. Sparse and stable portfolio selection with parameter uncertainty. *Journal of Business & Economic Statistics*, 33(3):381–392, 2015.

[17] Hanno Lustig, Nikolai Roussanov, and Adrien Verdelhan. Common risk factors in currency markets. *The Review of Financial Studies*, 24(11):3731–3777, 11 2011.

[18] Harry Markowitz. Portfolio selection. *The Journal of Finance*, 7(1):77–91, 1952.

[19] Alexander J. McNeil and Rudiger Frey. Estimation of tail-related risk measures for heteroscedastic financial time series: an extreme value approach. *Journal of Empirical Finance*, 7(3-4):271–300, November 2000.

[20] LUKAS MENKHOFF, Lucio Sarno, MAIK SCHMELING, and Andreas Schrimpf. Carry trades and global foreign exchange volatility. *Journal of Finance*, 67:681–718, 02 2011.

[21] Lukas Menkhoff, Lucio Sarno, Maik Schmeling, and Andreas Schrimpf. Carry trades and global foreign exchange volatility. *The Journal of Finance*, 67(2):681–718, 2012.

[22] Richard Michaud. The markowitz optimization enigma: Is 'optimized' optimal? *Financial Analysts Journal - FINANC ANAL J*, 45:31–42, 01 1989.

[23] Christopher J. Neely, Paul A. Weller, and Joshua M. Ulrich. The adaptive markets hypothesis: Evidence from the foreign exchange market. *Journal of Financial and Quantitative Analysis*, 44(2):467–488, 2009.

[24] Johanna G. NeÅ¡LehovÃ¡. Quantitative risk management: Concepts, techniques and tools , by alexander j. mcneil , rÃ¼diger frey and paul embrechts . revised edition. published by princeton university press , 2015 . total number of pages: 720. isbn: 978-0-691-16627-8 (hardback). *Journal of Time Series Analysis*, 37(3):431–432, May 2016.

[25] Angelo Ranaldo and Paul Söderlind. Safe haven currencies. *Review of Finance*, 14(3):385–407, 07 2010.

[26] R. Tyrrell Rockafellar, Stan Uryasev, and Michael Zabarankin. Generalized deviations in risk analysis. *Finance and Stochastics*, 10(1):51–74, 2006.

[27] R. Tyrrell Rockafellar and Stanislav Uryasev. Conditional value-at-risk for general loss distributions. *Journal of Banking & Finance*, 26(7):1443–1471, 2002.

[28] R. Tyrrell Rockafellar and Stanislav Uryasev. Optimization of conditional value-at-risk. *Journal of Risk*, None(None):None, None None.

[29] Thierry Roncalli. *Introduction to Risk Parity and Budgeting*. Chapman & Hall/CRC Financial Mathematics Series. Chapman & Hall/CRC, Boca Raton, FL, 2013.

[30] Patrick Lucescu Urban Ulrych. Regularized multi-currency expected shortfall portfolios, 2026.

[31] Raphael Burkhardt Urban Ulrych. Sparse and stable international portfolio optimization and currency risk management, 2023.

Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Declaration of originality

The signed declaration of originality is a component of every written paper or thesis authored during the course of studies. **In consultation with the supervisor**, one of the following two options must be selected:

☐ I hereby declare that I authored the work in question independently, i.e. that no one helped me to author it. Suggestions from the supervisor regarding language and content are excepted. I used no generative artificial intelligence technologies[1].

☐ I hereby declare that I authored the work in question independently. In doing so I only used the authorised aids, which included suggestions from the supervisor regarding language and content and generative artificial intelligence technologies. The use of the latter and the respective source declarations proceeded in consultation with the supervisor.

Title of paper or thesis:

Authored by:
If the work was compiled in a group, the names of all authors are required.

Last name(s):	**First name(s):**

With my signature I confirm the following:
- I have adhered to the rules set out in the Citation Guidelines.
- I have documented all methods, data and processes truthfully and fully.
- I have mentioned all persons who were significant facilitators of the work.

I am aware that the work may be screened electronically for originality.

Place, date	**Signature(s)**

If the work was compiled in a group, the names of all authors are required. Through their signatures they vouch jointly for the entire content of the written work.

[1] For further information please consult the ETH Zurich websites, e.g. https://ethz.ch/en/the-eth-zurich/education/ai-in-education.html and https://library.ethz.ch/en/researching-and-publishing/scientific-writing-at-eth-zurich.html (subject to change).